

Optical Flow and Object Tracking

簡韶逸 Shao-Yi Chien

Department of Electrical Engineering

National Taiwan University

Slide Credits

- Lucas-Kanade algorithm from Prof. Yung-Yu Chuang's VFX course
- Feature tracker and optical flow from Prof. Jia-Bin Huang's Computer Vision
- Visual Tracking from Prof. Alexandre Alahi, Stanford Vision Lab
- Slides from members of Media IC and System Lab

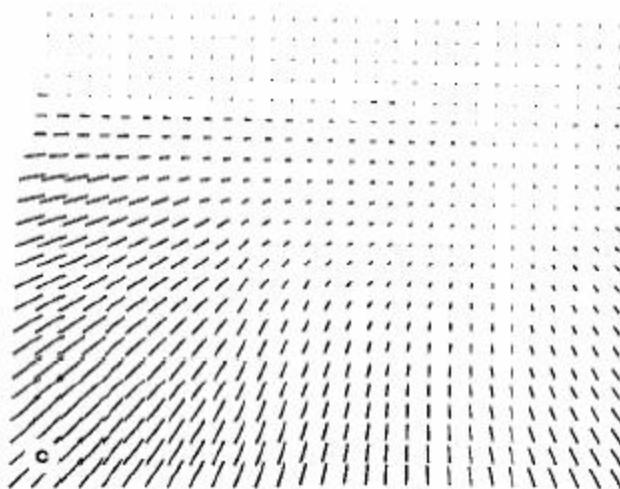
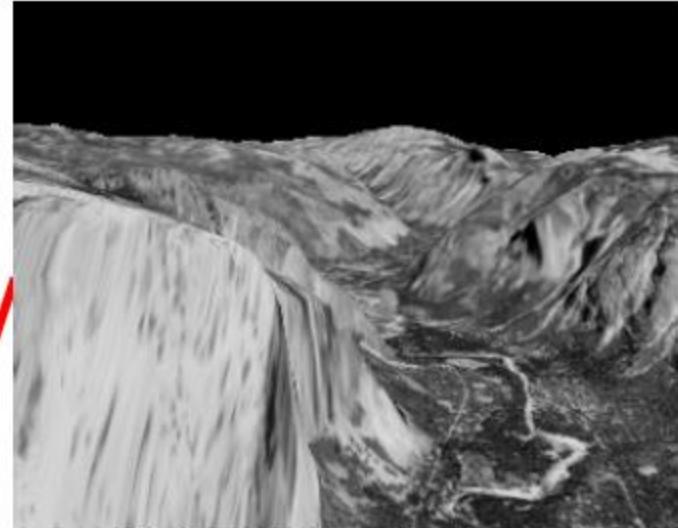
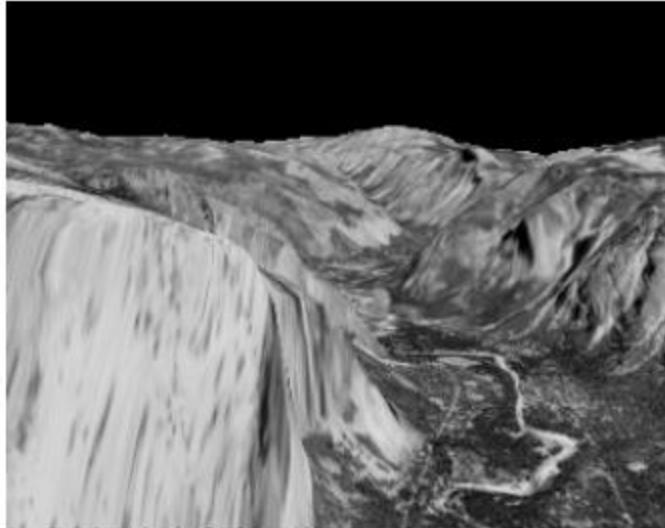
Outline

- Optical flow
 - Introduction
 - Lucas-Kanade algorithm
 - State-of-the-art optical flow
 - Epicflow
 - Flownet
 - PWC-Net

Introduction to Optical Flow

- Motion vector for each pixel
- Developed for 3-D motion estimation in the computer vision community
- Optical flow is caused by movement of intensity patterns in an image plane
- Optical flow is what we can only get from video frames

Introduction to Optical Flow



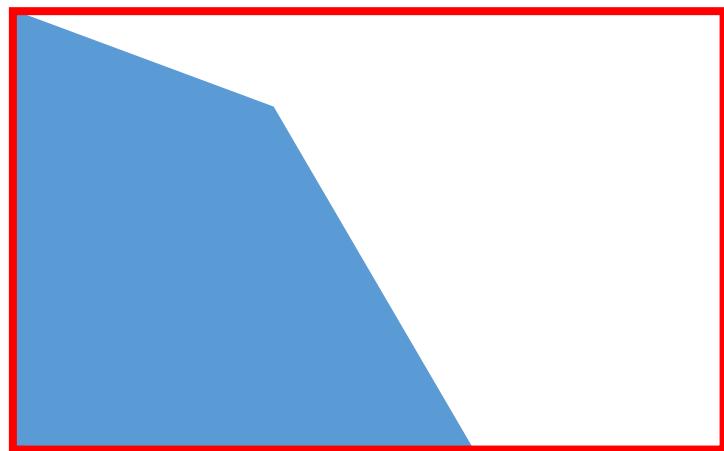
Introduction to Optical Flow

- Aperture problem
- Ill-posed inverse problem
 - Not well posed
 - Existance: No. Ex. rotating uniform sphere
 - Uniqueness: No. Aperture problem
 - Continuity: No. Sensitive to noise

Aperture Problem



Aperture Problem



Aperture Problem



Basic Assumptions for Optical Flow

- Brightness consistency
- Spatial coherence
- Temporal persistence

Horn-Schunck Algorithm

- B. K. P. Horn and B. G. Schunck, “Determining optical flow,” Artificial Intelligence, vol. 17, pp. 185—203, 1981.

For image function $f(x, y, t)$

Invariance of brightness $\frac{df}{dt} = 0$

$$\frac{\partial f}{\partial x} u + \frac{\partial f}{\partial y} v + \frac{\partial f}{\partial t} = 0$$

$$u = \frac{\partial x}{\partial t}, v = \frac{\partial y}{\partial t}$$

Smoothness constraint : to minimize $\left(\frac{\partial u}{\partial x}\right)^2 + \left(\frac{\partial u}{\partial y}\right)^2 + \left(\frac{\partial v}{\partial x}\right)^2 + \left(\frac{\partial v}{\partial y}\right)^2$

Horn-Schunck Algorithm

For image function $f(x, y, t)$

Invariance of brightness $\frac{df}{dt} = 0$

$$\varepsilon_b^2 = \left(\frac{\partial f}{\partial x} u + \frac{\partial f}{\partial y} v + \frac{\partial f}{\partial t} \right)^2$$

$$\varepsilon_s^2 = \left(\frac{\partial u}{\partial x} \right)^2 + \left(\frac{\partial u}{\partial y} \right)^2 + \left(\frac{\partial v}{\partial x} \right)^2 + \left(\frac{\partial v}{\partial y} \right)^2$$

To minimize $\varepsilon = \sum_x \sum_y \varepsilon_b^2 + \alpha^2 \varepsilon_s^2$

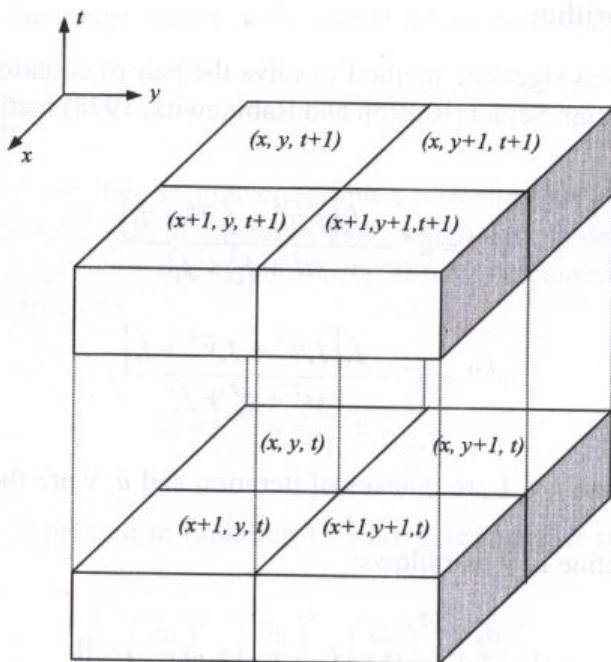
Iteration solution :

$$u^{k+1} = u^k - \frac{f_x [f_x u^k + f_y v^k + f_t]}{\alpha^2 + f_x^2 + f_y^2}$$

$$v^{k+1} = v^k - \frac{f_y [f_x u^k + f_y v^k + f_t]}{\alpha^2 + f_x^2 + f_y^2}$$

$$f_x = \frac{\partial f}{\partial x}, f_y = \frac{\partial f}{\partial y}, f_t = \frac{\partial f}{\partial t}$$

Horn-Schunck Algorithm



$$f_x = \frac{1}{4} \left\{ [f(x+1, y, t) - f(x, y, t)] + [f(x+1, y+1, t) - f(x, y+1, t)] \right.$$

$$\left. + [f(x+1, y+1, t) - f(x, y, t)] + [f(x+1, y+1, t+1) - f(x, y+1, t+1)] \right\}$$

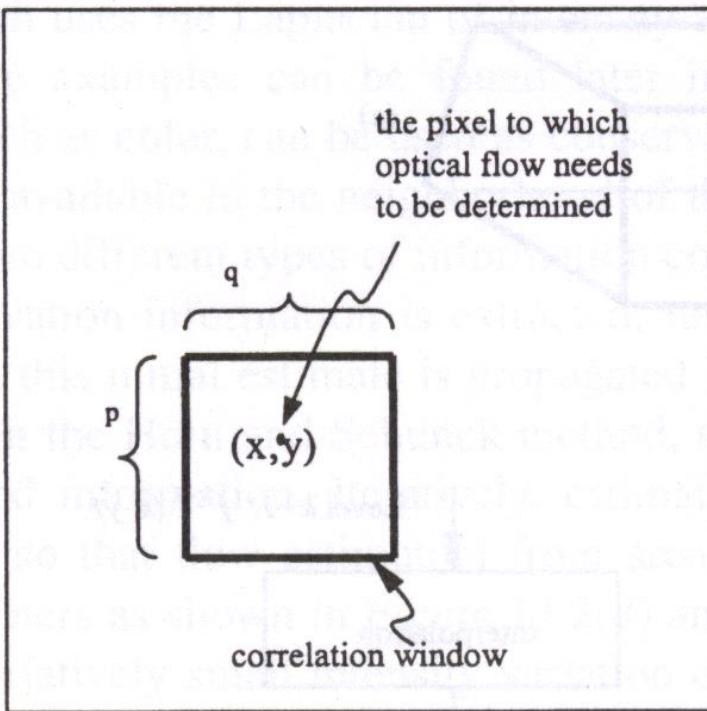
$$f_y = \frac{1}{4} \left\{ [f(x, y+1, t) - f(x, y, t)] + [f(x+1, y+1, t) - f(x+1, y, t)] \right.$$

$$\left. + [f(x, y+1, t+1) - f(x, y, t+1)] + [f(x+1, y+1, t+1) - f(x+1, y, t+1)] \right\}$$

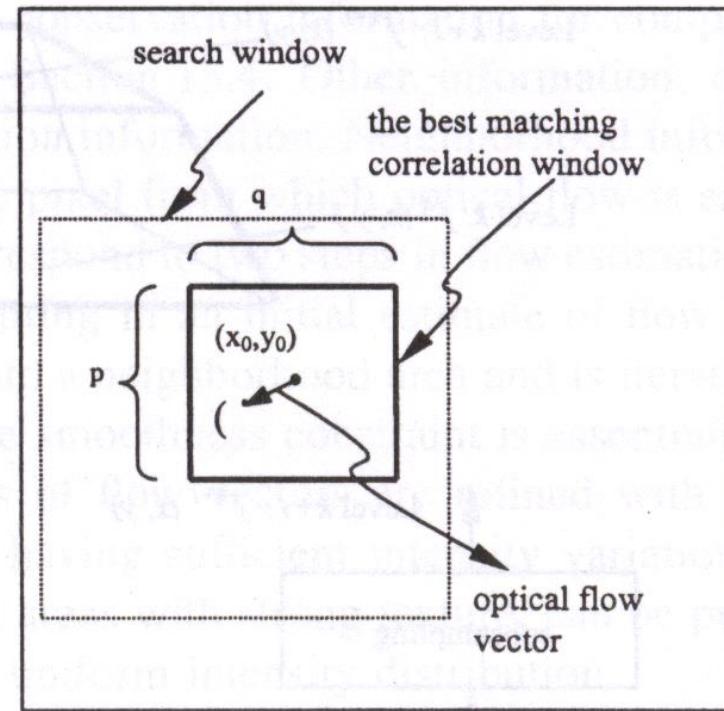
$$f_x = \frac{1}{4} \left\{ [f(x, y, t+1) - f(x, y, t)] + [f(x+1, y, t+1) - f(x+1, y, t)] \right.$$

$$\left. + [f(x, y+1, t+1) - f(x, y+1, t)] + [f(x+1, y+1, t+1) - f(x+1, y+1, t)] \right\}$$

Correlation-Based Approach

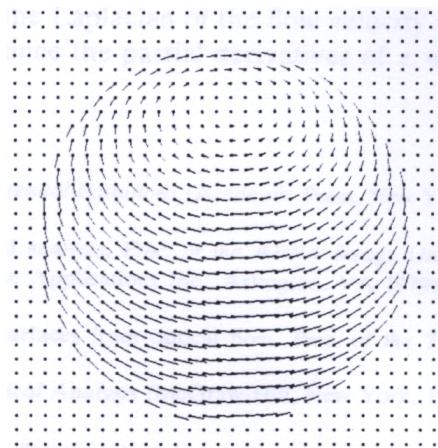
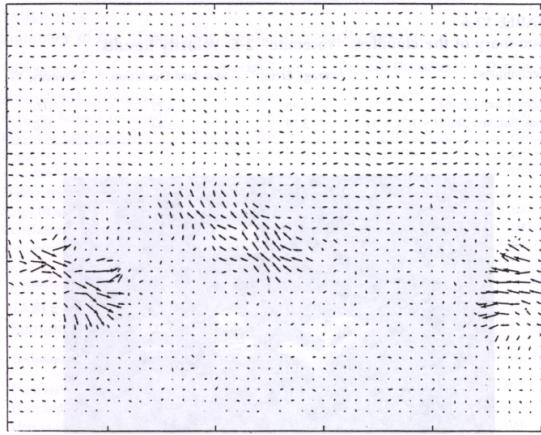
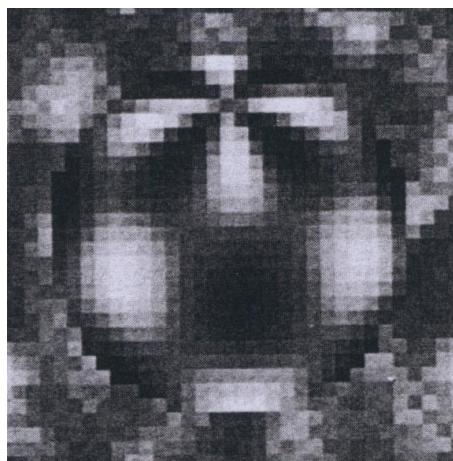


$$f(x, y, t)$$



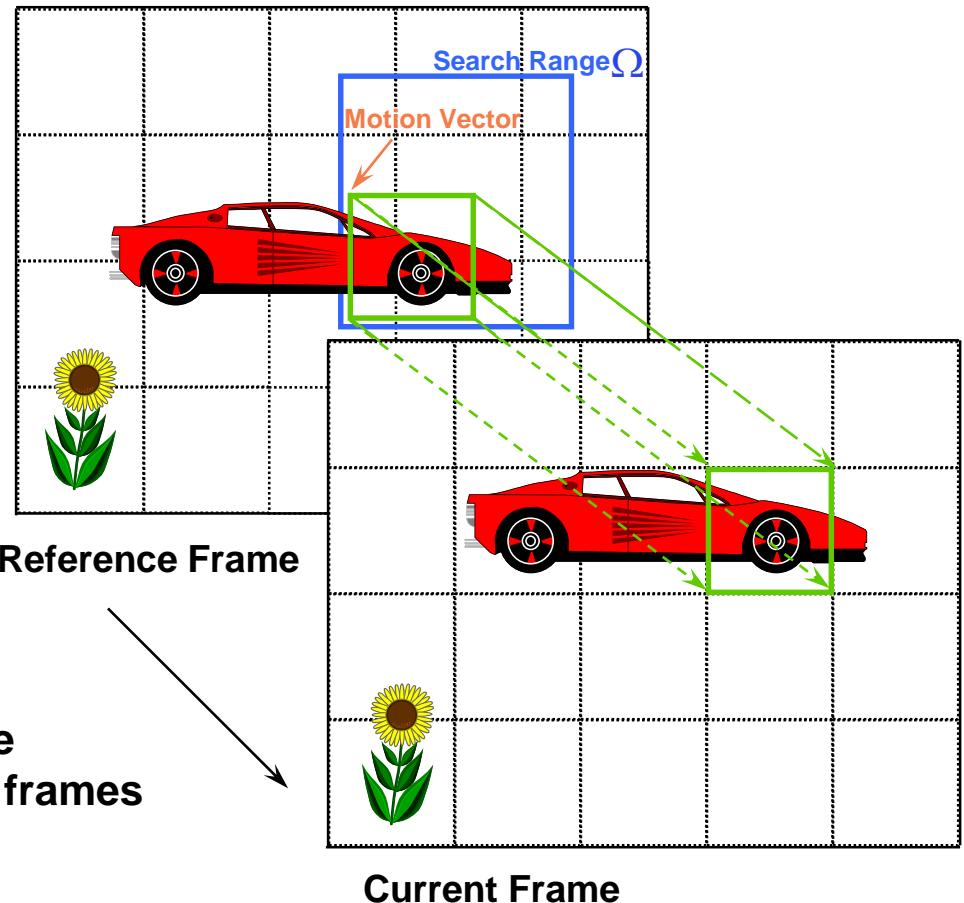
$$f(x, y, t - 1)$$

Some Example



Needle diagram

Block-Matching Motion Estimation



Motion Vector

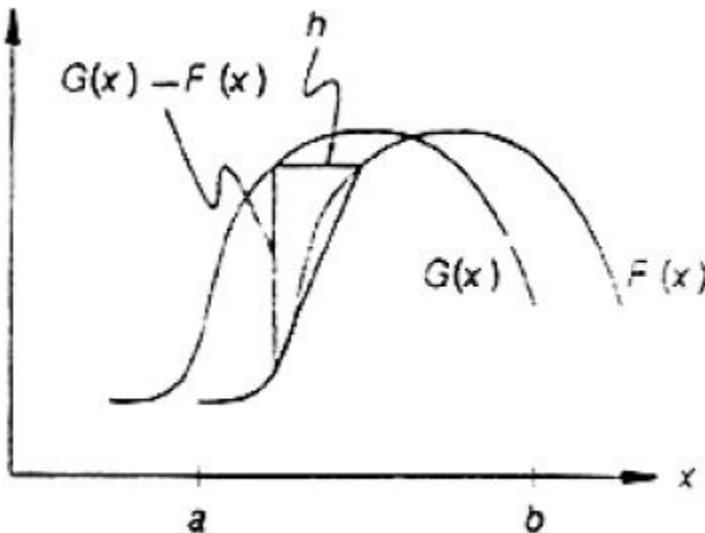
$$V_t(p,q) = (Vec_i, Vec_j)$$

the location in the search range Ω
that has the maximum correlation value
between blocks in temporally adjacent frames

Lucas-Kanade Optical Flow

To minimize

$$E = \sum_x [F(x + h) - G(x)]^2.$$



$$F(x+h) \approx F(x) + hF'(x),$$

$$\begin{aligned} 0 &= \frac{\partial E}{\partial h} \\ &\approx \frac{\partial}{\partial h} \sum_x [F(x) + hF'(x) - G(x)]^2 \\ &= \sum_x 2F'(x)[F(x) + hF'(x) - G(x)] \end{aligned}$$

from which

$$h \approx \frac{\sum_x F'(x)[G(x) - F(x)]}{\sum_x F'(x)^2}.$$

Ref: Lucas and Kanade, "An iterative image registration technique with an application to stereo vision," IJCAI, 1981.

Lucas-Kanade Optical Flow

- Iterative solution

$$h_0 = 0,$$

$$h_{k+1} = h_k + \frac{\sum_x w(x) F(x + h_k) [G(x) - F(x + h_k)]}{\sum_x w(x) F(x + h_k)^2},$$

Add weights

$$w(x) = \frac{1}{|G(x) - F(x)|}.$$

- More dimensions

$$h = \left[\sum_x \left(\frac{\partial F}{\partial x} \right)^T [G(x) - F(x)] \right] \left[\sum_x \left(\frac{\partial F}{\partial x} \right)^T \left(\frac{\partial F}{\partial x} \right) \right]^{-1}$$

Lucas-Kanade Algorithm

$$E(u, v) = \sum_{x, y} (I(x + u, y + v) - T(x, y))^2$$

$$I(x + u, y + v) \approx I(x, y) + uI_x + vI_y$$

$$= \sum_{x, y} (I(x, y) - T(x, y) + uI_x + vI_y)^2$$

$$0 = \frac{\partial E}{\partial u} = \sum_{x, y} 2I_x (I(x, y) - T(x, y) + uI_x + vI_y)$$

$$0 = \frac{\partial E}{\partial v} = \sum_{x, y} 2I_y (I(x, y) - T(x, y) + uI_x + vI_y)$$

Lucas Kanade Algorithm

$$0 = \frac{\partial E}{\partial u} = \sum_{x,y} 2I_x (I(x, y) - T(x, y) + uI_x + vI_y)$$

$$0 = \frac{\partial E}{\partial v} = \sum_{x,y} 2I_y (I(x, y) - T(x, y) + uI_x + vI_y)$$

$$\begin{aligned} \rightarrow & \begin{cases} \sum_{x,y} I_x^2 u + \sum_{x,y} I_x I_y v = \sum_{x,y} I_x (T(x, y) - I(x, y)) \\ \sum_{x,y} I_x I_y u + \sum_{x,y} I_y^2 v = \sum_{x,y} I_y (T(x, y) - I(x, y)) \end{cases} \end{aligned}$$

$$\begin{aligned} \rightarrow & \begin{bmatrix} \sum_{x,y} I_x^2 & \sum_{x,y} I_x I_y \\ \sum_{x,y} I_x I_y & \sum_{x,y} I_y^2 \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = \begin{bmatrix} \sum_{x,y} I_x (T(x, y) - I(x, y)) \\ \sum_{x,y} I_y (T(x, y) - I(x, y)) \end{bmatrix} \end{aligned}$$

Lucas Kanade Algorithm

iterate

shift $I(x,y)$ with (u,v)

compute gradient image I_x, I_y

compute error image $T(x,y) - I(x,y)$

compute Hessian matrix

solve the linear system

$$(u,v) = (u,v) + (\Delta u, \Delta v)$$

until converge

$$\begin{bmatrix} \sum_{x,y} I_x^2 & \sum_{x,y} I_x I_y \\ \sum_{x,y} I_x I_y & \sum_{x,y} I_y^2 \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = \begin{bmatrix} \sum_{x,y} I_x (T(x,y) - I(x,y)) \\ \sum_{x,y} I_y (T(x,y) - I(x,y)) \end{bmatrix}$$

Lucas Kanade Algorithm: Parametric

$$E(u, v) = \sum_{x, y} (I(x+u, y+v) - T(x, y))^2$$

→ $E(\mathbf{p}) = \sum_{\mathbf{x}} (I(\mathbf{W}(\mathbf{x}; \mathbf{p})) - T(\mathbf{x}))^2$ ← Our goal is to find \mathbf{p} to minimize $E(\mathbf{p})$
for all \mathbf{x} in T 's domain

translation $\mathbf{W}(\mathbf{x}; \mathbf{p}) = \begin{pmatrix} x + d_x \\ y + d_y \end{pmatrix}, p = (d_x, d_y)^T$

affine $\mathbf{W}(\mathbf{x}; \mathbf{p}) = \mathbf{A}\mathbf{x} + \mathbf{d} = \begin{pmatrix} 1 + d_{xx} & d_{xy} & d_x \\ d_{yx} & 1 + d_{yy} & d_y \end{pmatrix} \begin{pmatrix} x \\ y \\ 1 \end{pmatrix},$

$$p = (d_{xx}, d_{xy}, d_{yx}, d_{yy}, d_x, d_y)^T$$

Lucas Kanade Algorithm: Parametric

$$\text{minimize} \quad \sum_{\mathbf{x}} (I(\mathbf{W}(\mathbf{x}; \mathbf{p} + \Delta\mathbf{p})) - T(\mathbf{x}))^2$$

with respect to $\Delta\mathbf{p}$

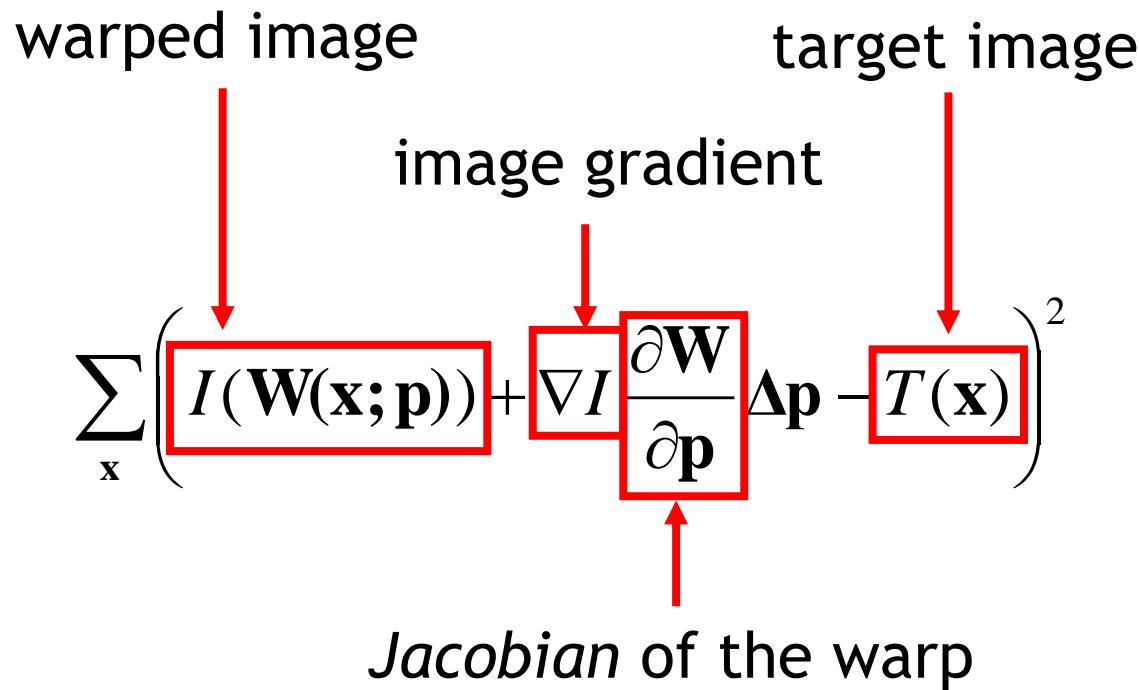
$$\mathbf{W}(\mathbf{x}; \mathbf{p} + \Delta\mathbf{p}) \approx \mathbf{W}(\mathbf{x}; \mathbf{p}) + \frac{\partial \mathbf{W}}{\partial \mathbf{p}} \Delta\mathbf{p}$$

$$I(\mathbf{W}(\mathbf{x}; \mathbf{p} + \Delta\mathbf{p})) \approx I(\mathbf{W}(\mathbf{x}; \mathbf{p}) + \frac{\partial \mathbf{W}}{\partial \mathbf{p}} \Delta\mathbf{p})$$

$$\approx I(\mathbf{W}(\mathbf{x}; \mathbf{p})) + \frac{\partial I}{\partial \mathbf{x}} \frac{\partial \mathbf{W}}{\partial \mathbf{p}} \Delta\mathbf{p}$$

$$\rightarrow \text{minimize} \sum_{\mathbf{x}} \left(I(\mathbf{W}(\mathbf{x}; \mathbf{p})) + \nabla I \frac{\partial \mathbf{W}}{\partial \mathbf{p}} \Delta\mathbf{p} - T(\mathbf{x}) \right)^2$$

Lucas Kanade Algorithm: Parametric



$$\frac{\partial \mathbf{W}}{\partial \mathbf{p}} = \begin{pmatrix} \frac{\partial \mathbf{W}_x}{\partial \mathbf{p}} \\ \frac{\partial \mathbf{W}_y}{\partial \mathbf{p}} \end{pmatrix} = \begin{pmatrix} \frac{\partial \mathbf{W}_x}{\partial \mathbf{p}_1} & \frac{\partial \mathbf{W}_x}{\partial \mathbf{p}_2} & \dots & \frac{\partial \mathbf{W}_x}{\partial \mathbf{p}_n} \\ \frac{\partial \mathbf{W}_y}{\partial \mathbf{p}_1} & \frac{\partial \mathbf{W}_y}{\partial \mathbf{p}_2} & \dots & \frac{\partial \mathbf{W}_y}{\partial \mathbf{p}_n} \end{pmatrix}$$

Lucas Kanade Algorithm: Parametric

$$\frac{\partial \mathbf{W}}{\partial \mathbf{p}} = \begin{pmatrix} \frac{\partial \mathbf{W}_x}{\partial \mathbf{p}} \\ \frac{\partial \mathbf{W}_y}{\partial \mathbf{p}} \end{pmatrix} = \begin{pmatrix} \frac{\partial \mathbf{W}_x}{\partial \mathbf{p}_1} & \frac{\partial \mathbf{W}_x}{\partial \mathbf{p}_2} & \dots & \frac{\partial \mathbf{W}_x}{\partial \mathbf{p}_n} \\ \frac{\partial \mathbf{W}_y}{\partial \mathbf{p}_1} & \frac{\partial \mathbf{W}_y}{\partial \mathbf{p}_2} & \dots & \frac{\partial \mathbf{W}_y}{\partial \mathbf{p}_n} \end{pmatrix}$$

For example, for affine

$$\mathbf{W}(\mathbf{x}; \mathbf{p}) = \begin{pmatrix} 1 + d_{xx} & d_{xy} & d_x \\ d_{yx} & 1 + d_{yy} & d_y \end{pmatrix} \begin{pmatrix} x \\ y \\ 1 \end{pmatrix} = \begin{pmatrix} (1 + d_{xx})x + d_{xy}y + d_x \\ d_{yx}x + (1 + d_{yy})y + d_y \end{pmatrix}$$

→ $\frac{\partial \mathbf{W}}{\partial \mathbf{p}} = \begin{pmatrix} x & 0 & y & 0 & 1 & 0 \\ 0 & x & 0 & y & 0 & 1 \\ d_{xx} & d_{yx} & d_{xy} & d_{yy} & d_x & d_y \end{pmatrix}$

Lucas Kanade Algorithm: Parametric

$$\arg \min_{\Delta \mathbf{p}} \sum_{\mathbf{x}} \left(I(\mathbf{W}(\mathbf{x}; \mathbf{p})) + \nabla I \frac{\partial \mathbf{W}}{\partial \mathbf{p}} \Delta \mathbf{p} - T(\mathbf{x}) \right)^2$$

 $0 = \sum_{\mathbf{x}} \left[\nabla I \frac{\partial \mathbf{W}}{\partial \mathbf{p}} \right]^T \left[I(\mathbf{W}(\mathbf{x}; \mathbf{p})) + \nabla I \frac{\partial \mathbf{W}}{\partial \mathbf{p}} \Delta \mathbf{p} - T(\mathbf{x}) \right]$

$$\Delta \mathbf{p} = \mathbf{H}^{-1} \sum_{\mathbf{x}} \left[\nabla I \frac{\partial \mathbf{W}}{\partial \mathbf{p}} \right]^T [T(\mathbf{x}) - I(\mathbf{W}(\mathbf{x}; \mathbf{p}))]$$

(Approximated) Hessian $\mathbf{H} = \sum_{\mathbf{x}} \left[\nabla I \frac{\partial \mathbf{W}}{\partial \mathbf{p}} \right]^T \left[\nabla I \frac{\partial \mathbf{W}}{\partial \mathbf{p}} \right]$

Lucas Kanade Algorithm

iterate

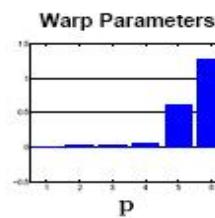
- 1) warp I with $W(x; p)$
- 2) compute error image $T(x, y) - I(W(x, p))$
- 3) compute gradient image ∇I with $W(x, p)$
- 4) evaluate Jacobian $\frac{\partial W}{\partial p}$ at $(x; p)$
- 5) compute $\nabla I \frac{\partial W}{\partial p}$
- 6) compute Hessian
- 7) compute $\sum_x \left[\nabla I \frac{\partial W}{\partial p} \right]^T [T(x) - I(W(x; p))]$
- 8) solve Δp
- 9) update p by $p + \Delta p$

until converge

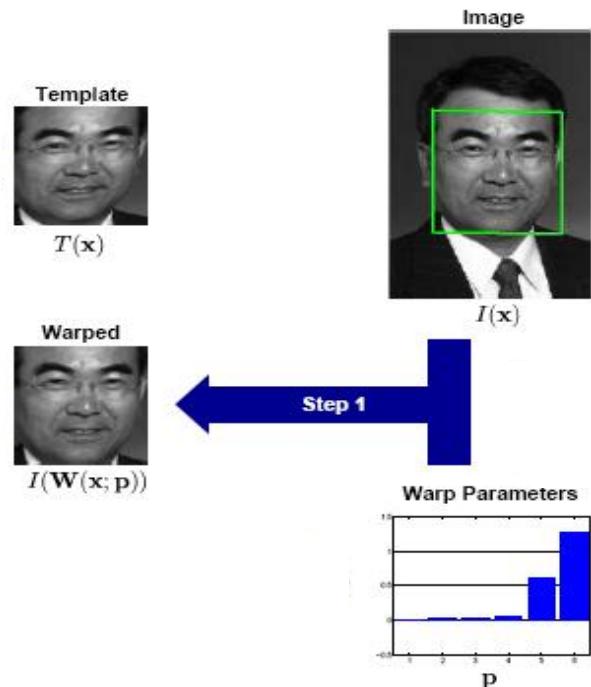
$$\Delta p = H^{-1} \sum_x \left[\nabla I \frac{\partial W}{\partial p} \right]^T [T(x) - I(W(x; p))]$$



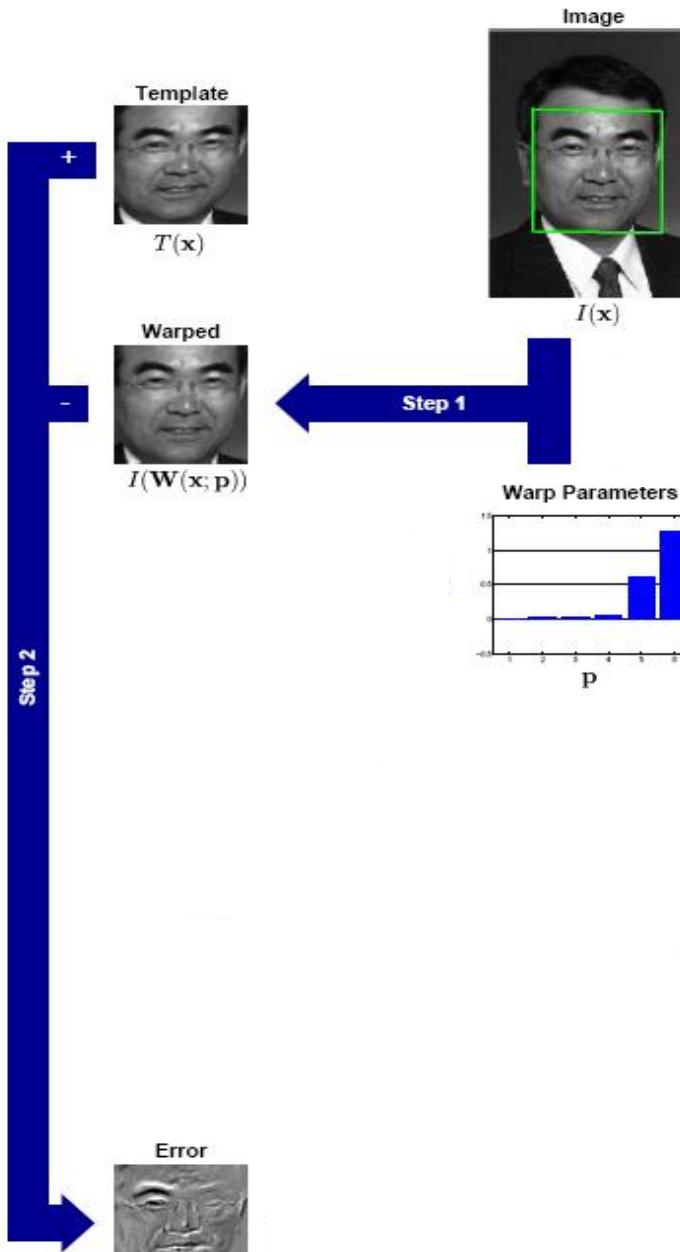
$$\Delta \mathbf{p} = \mathbf{H}^{-1} \sum_{\mathbf{x}} \left[\nabla I \frac{\partial \mathbf{W}}{\partial \mathbf{p}} \right]^T [T(\mathbf{x}) - I(\mathbf{W}(\mathbf{x}; \mathbf{p}))]$$



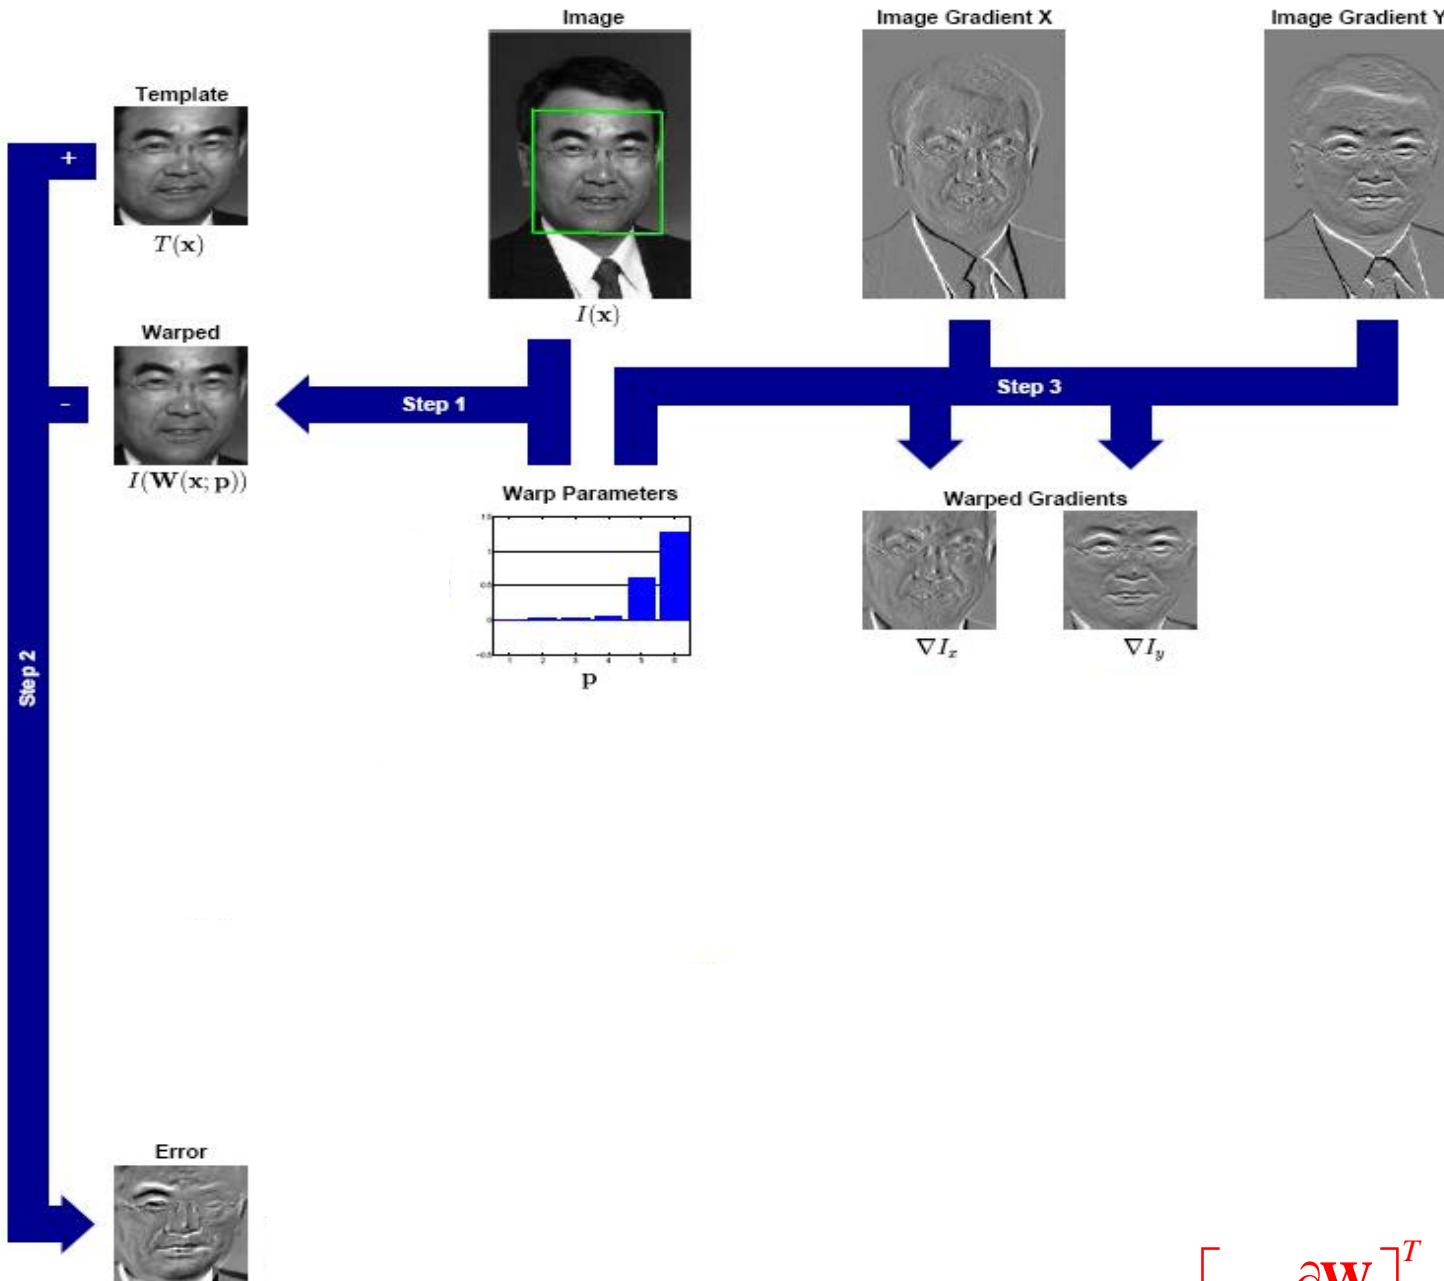
$$\Delta \mathbf{p} = \mathbf{H}^{-1} \sum_{\mathbf{x}} \left[\nabla I \frac{\partial \mathbf{W}}{\partial \mathbf{p}} \right]^T [T(\mathbf{x}) - I(\mathbf{W}(\mathbf{x}; \mathbf{p}))]$$



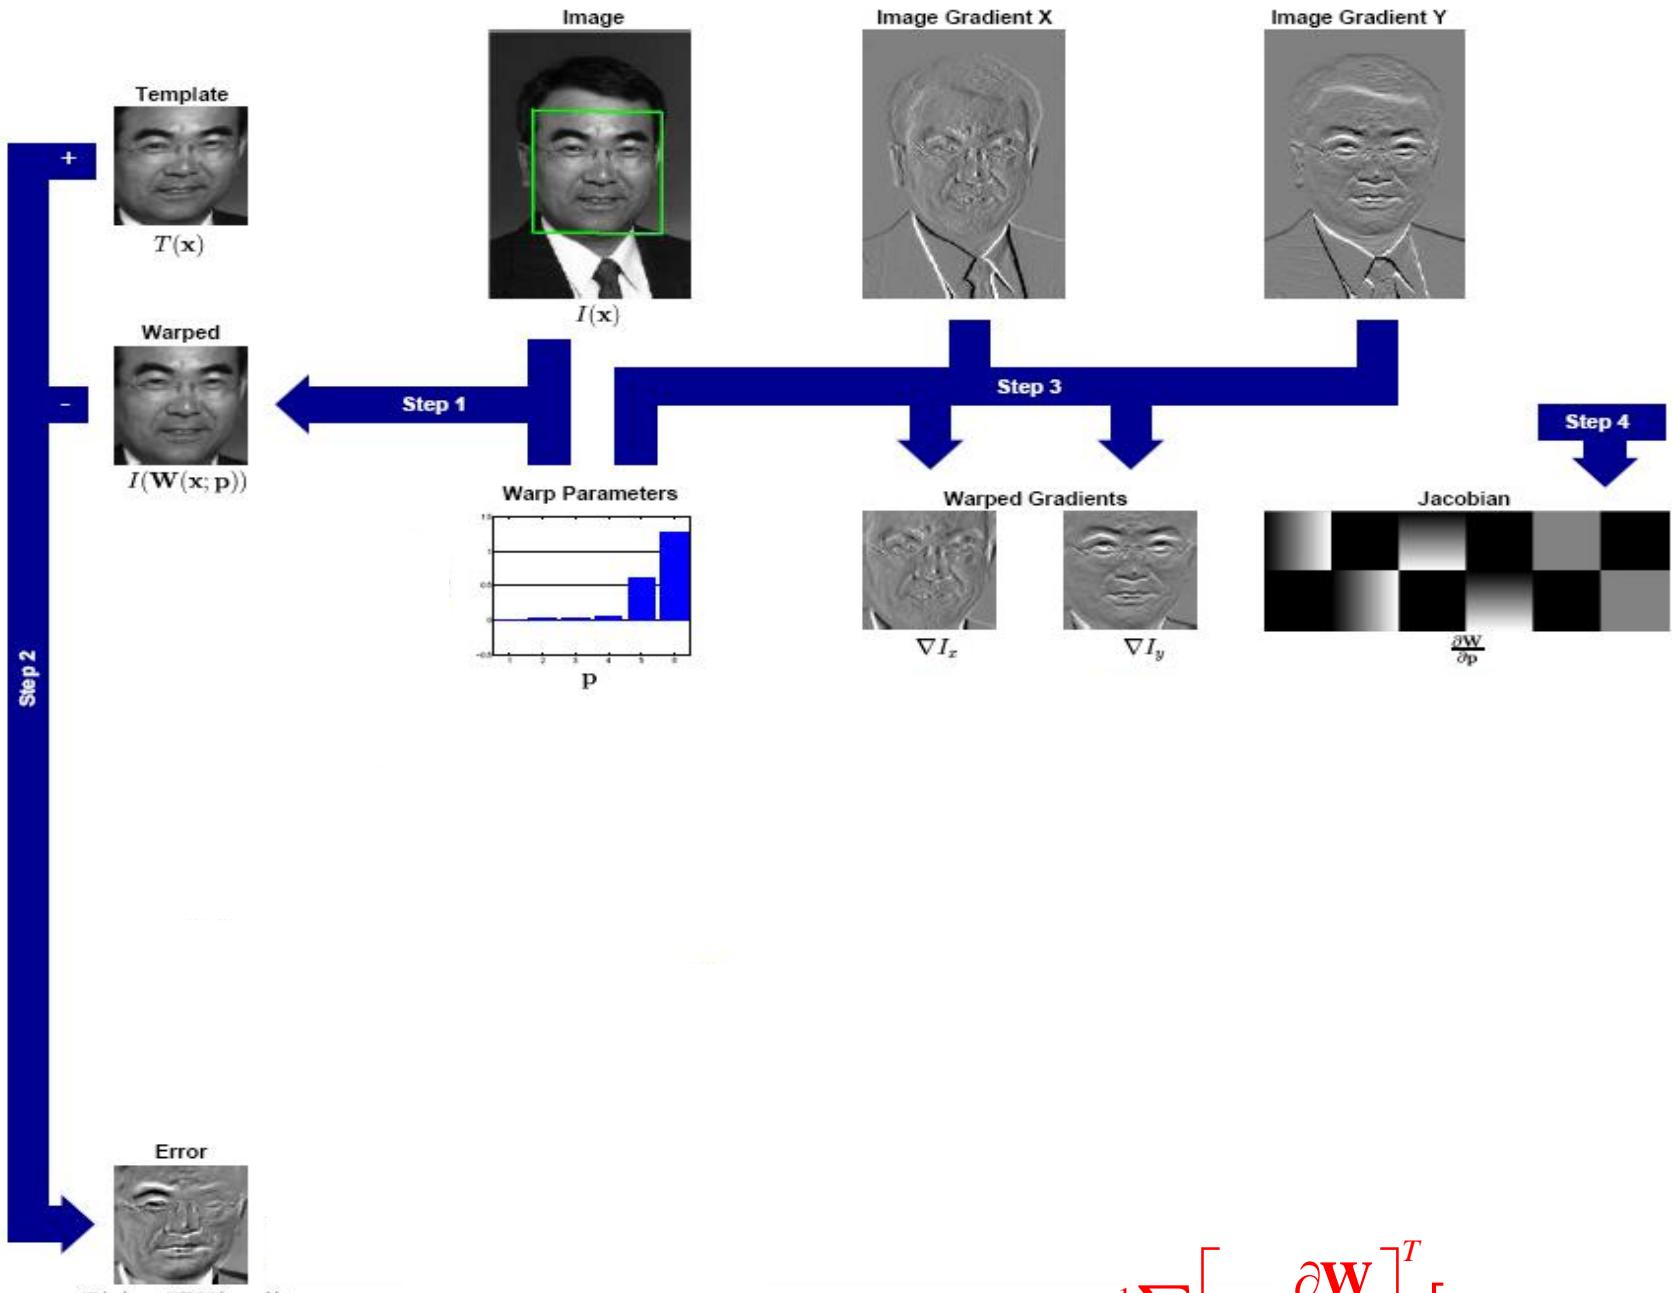
$$\Delta \mathbf{p} = \mathbf{H}^{-1} \sum_{\mathbf{x}} \left[\nabla I \frac{\partial \mathbf{W}}{\partial \mathbf{p}} \right]^T [T(\mathbf{x}) - I(\mathbf{W}(\mathbf{x}; \mathbf{p}))]$$



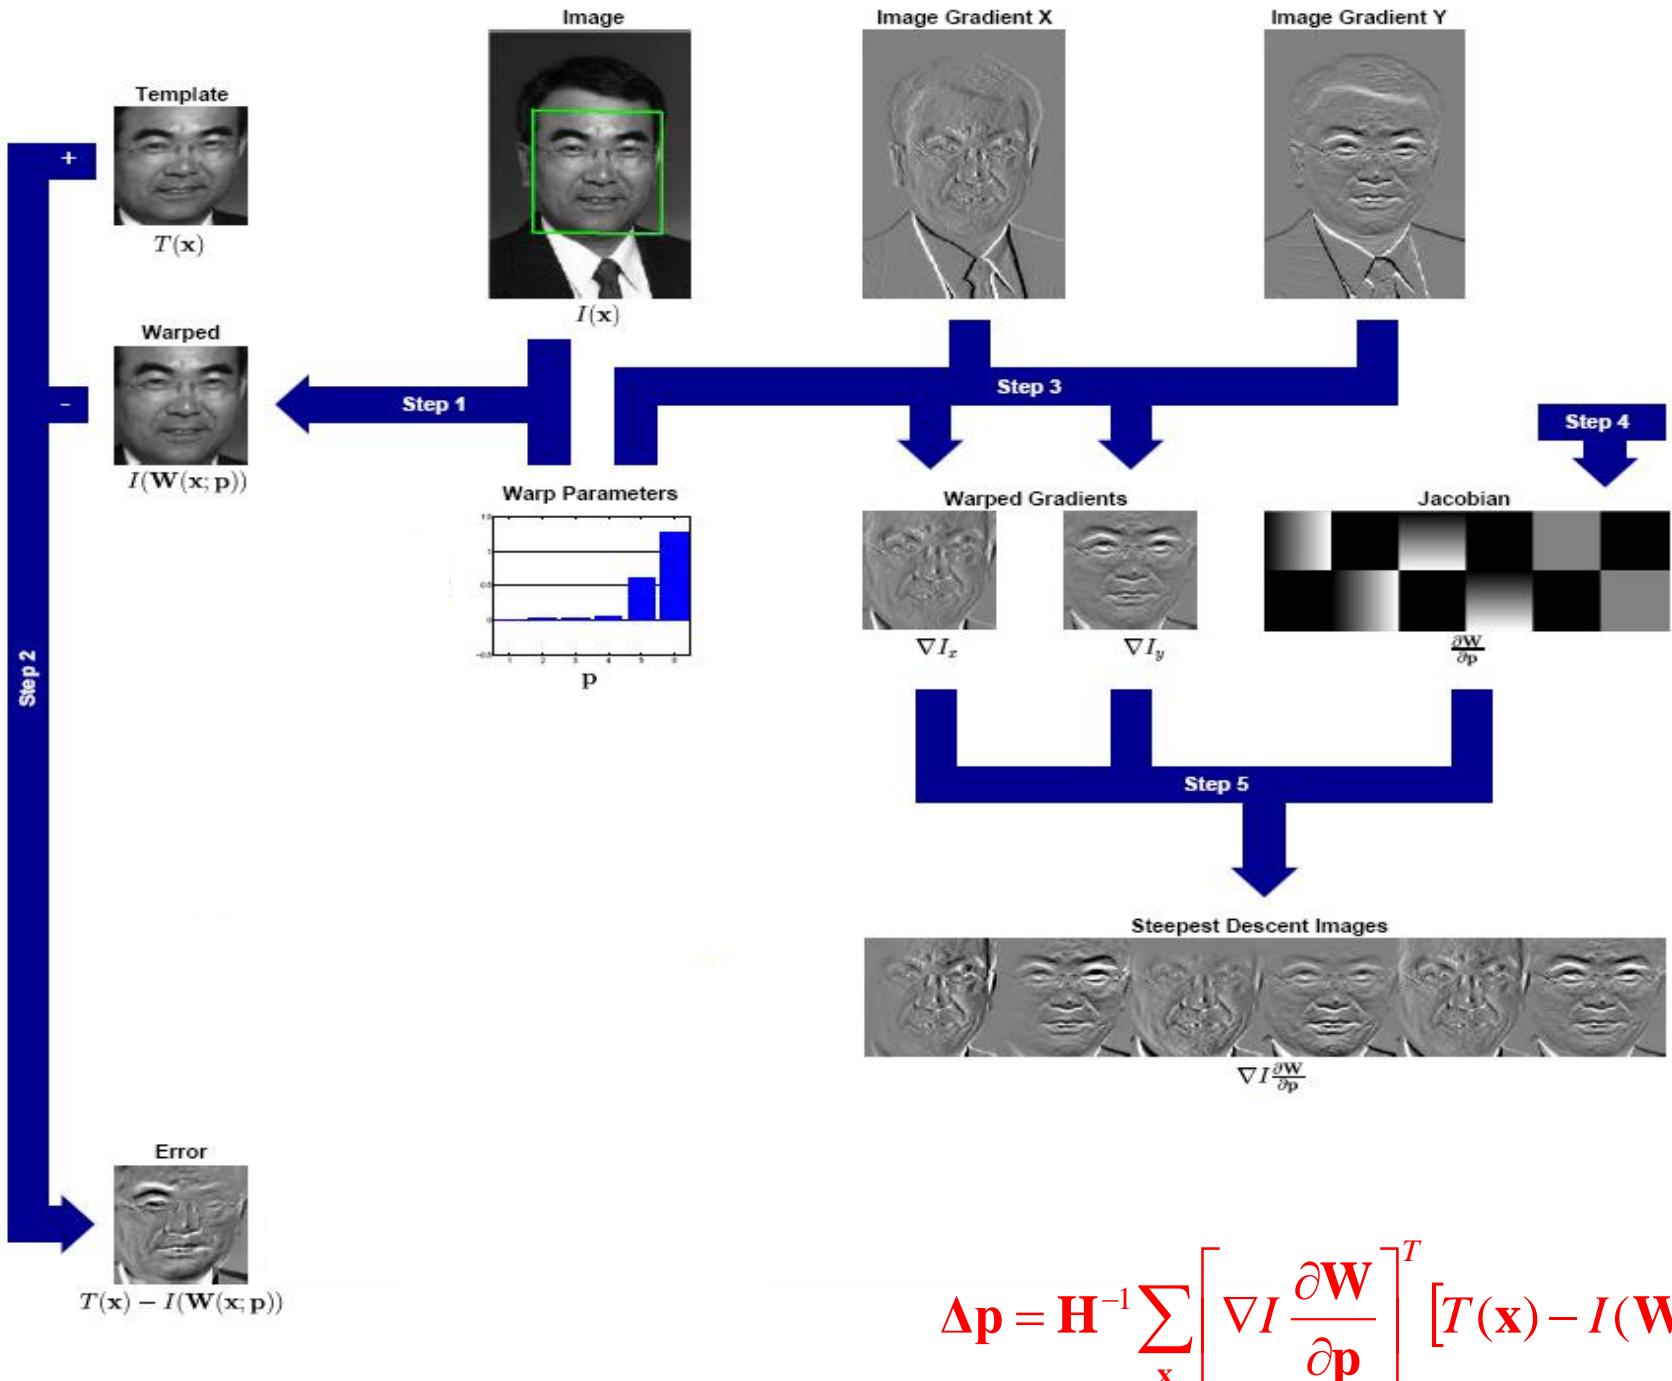
$$\Delta \mathbf{p} = \mathbf{H}^{-1} \sum_{\mathbf{x}} \left[\nabla I \frac{\partial \mathbf{W}}{\partial \mathbf{p}} \right]^T [T(\mathbf{x}) - I(\mathbf{W}(\mathbf{x}; \mathbf{p}))]$$

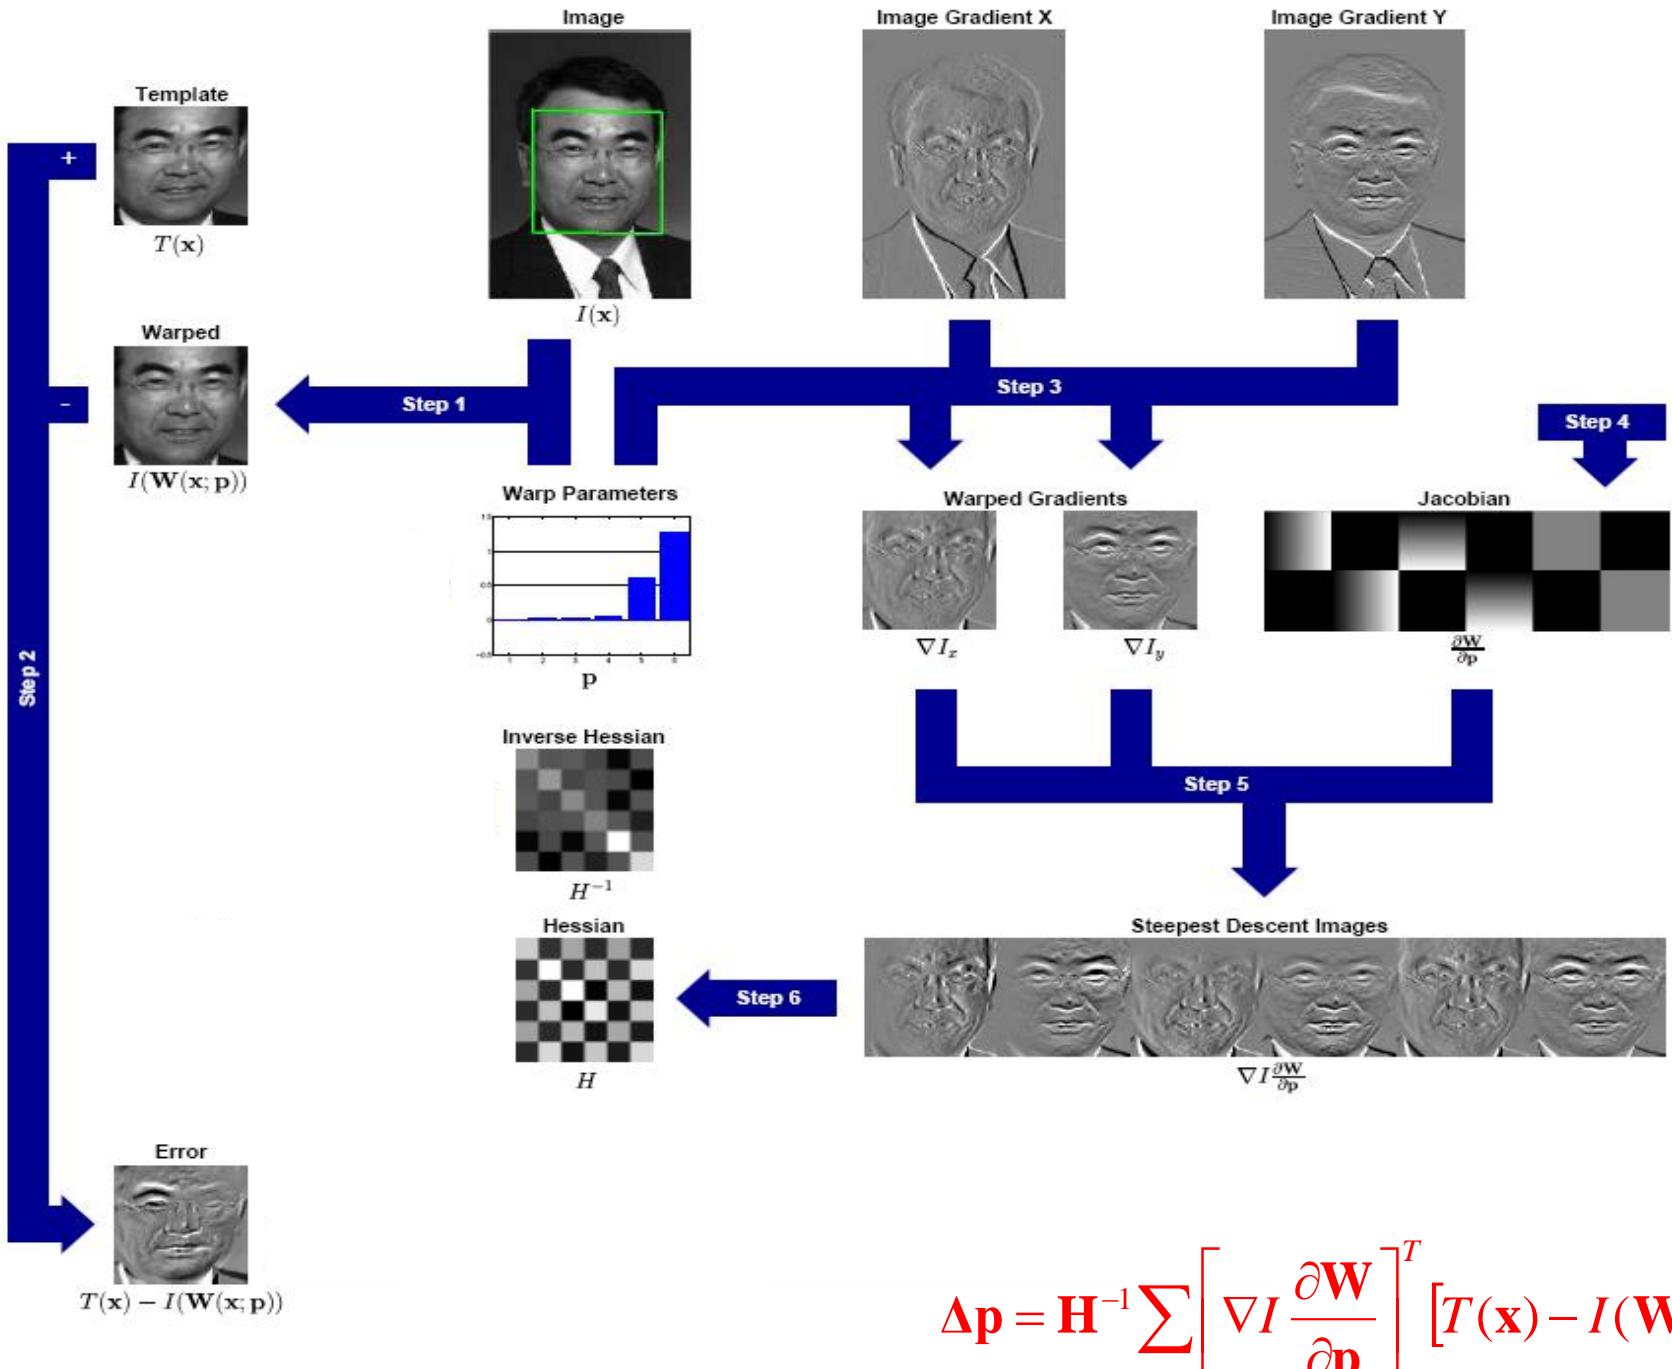


$$\Delta \mathbf{p} = \mathbf{H}^{-1} \sum_{\mathbf{x}} \left[\nabla I \frac{\partial \mathbf{W}}{\partial \mathbf{p}} \right]^T [T(\mathbf{x}) - I(\mathbf{W}(\mathbf{x}; \mathbf{p}))]$$

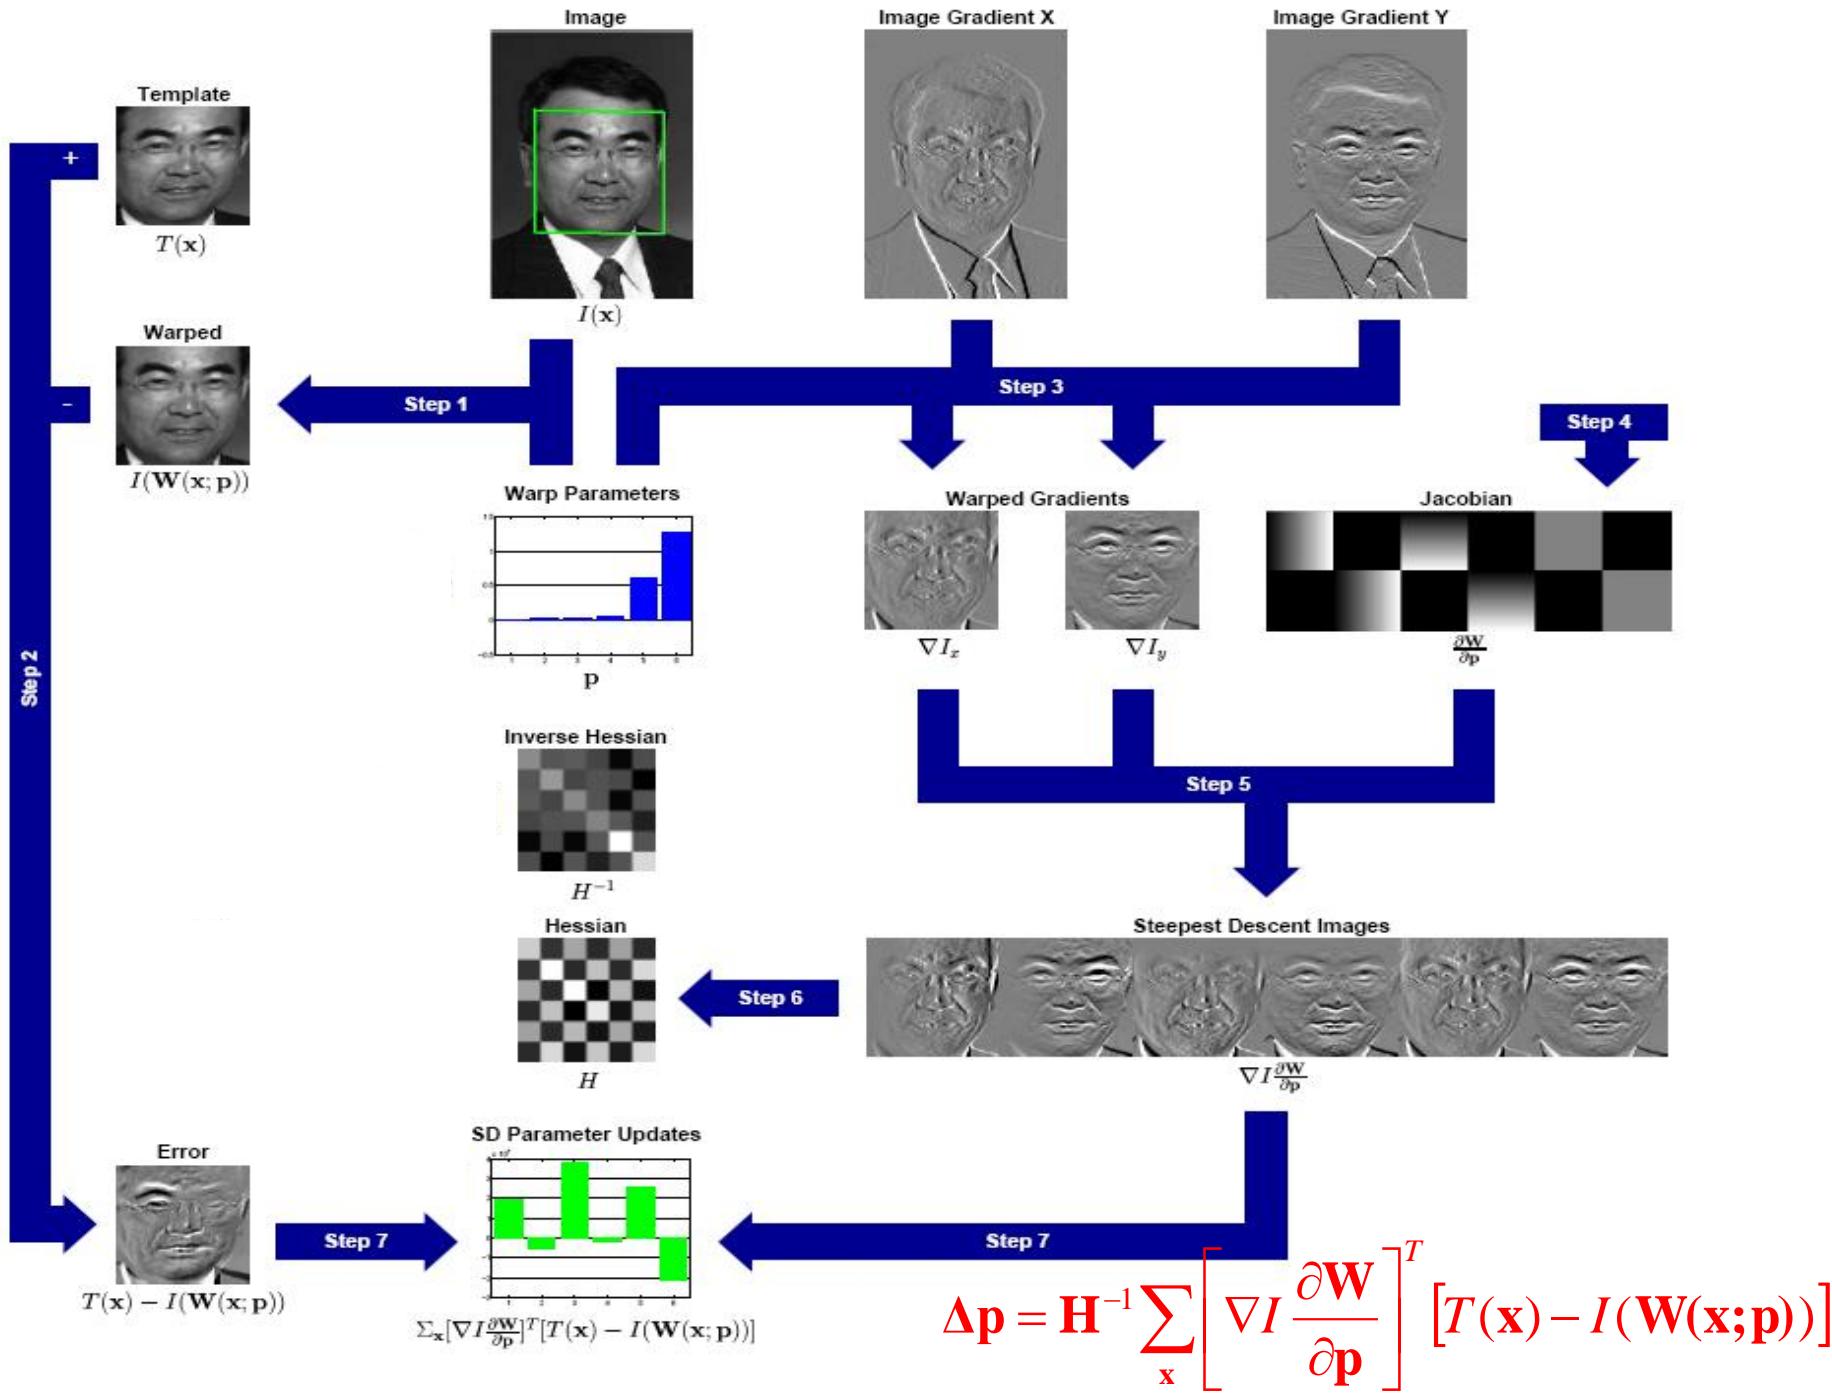


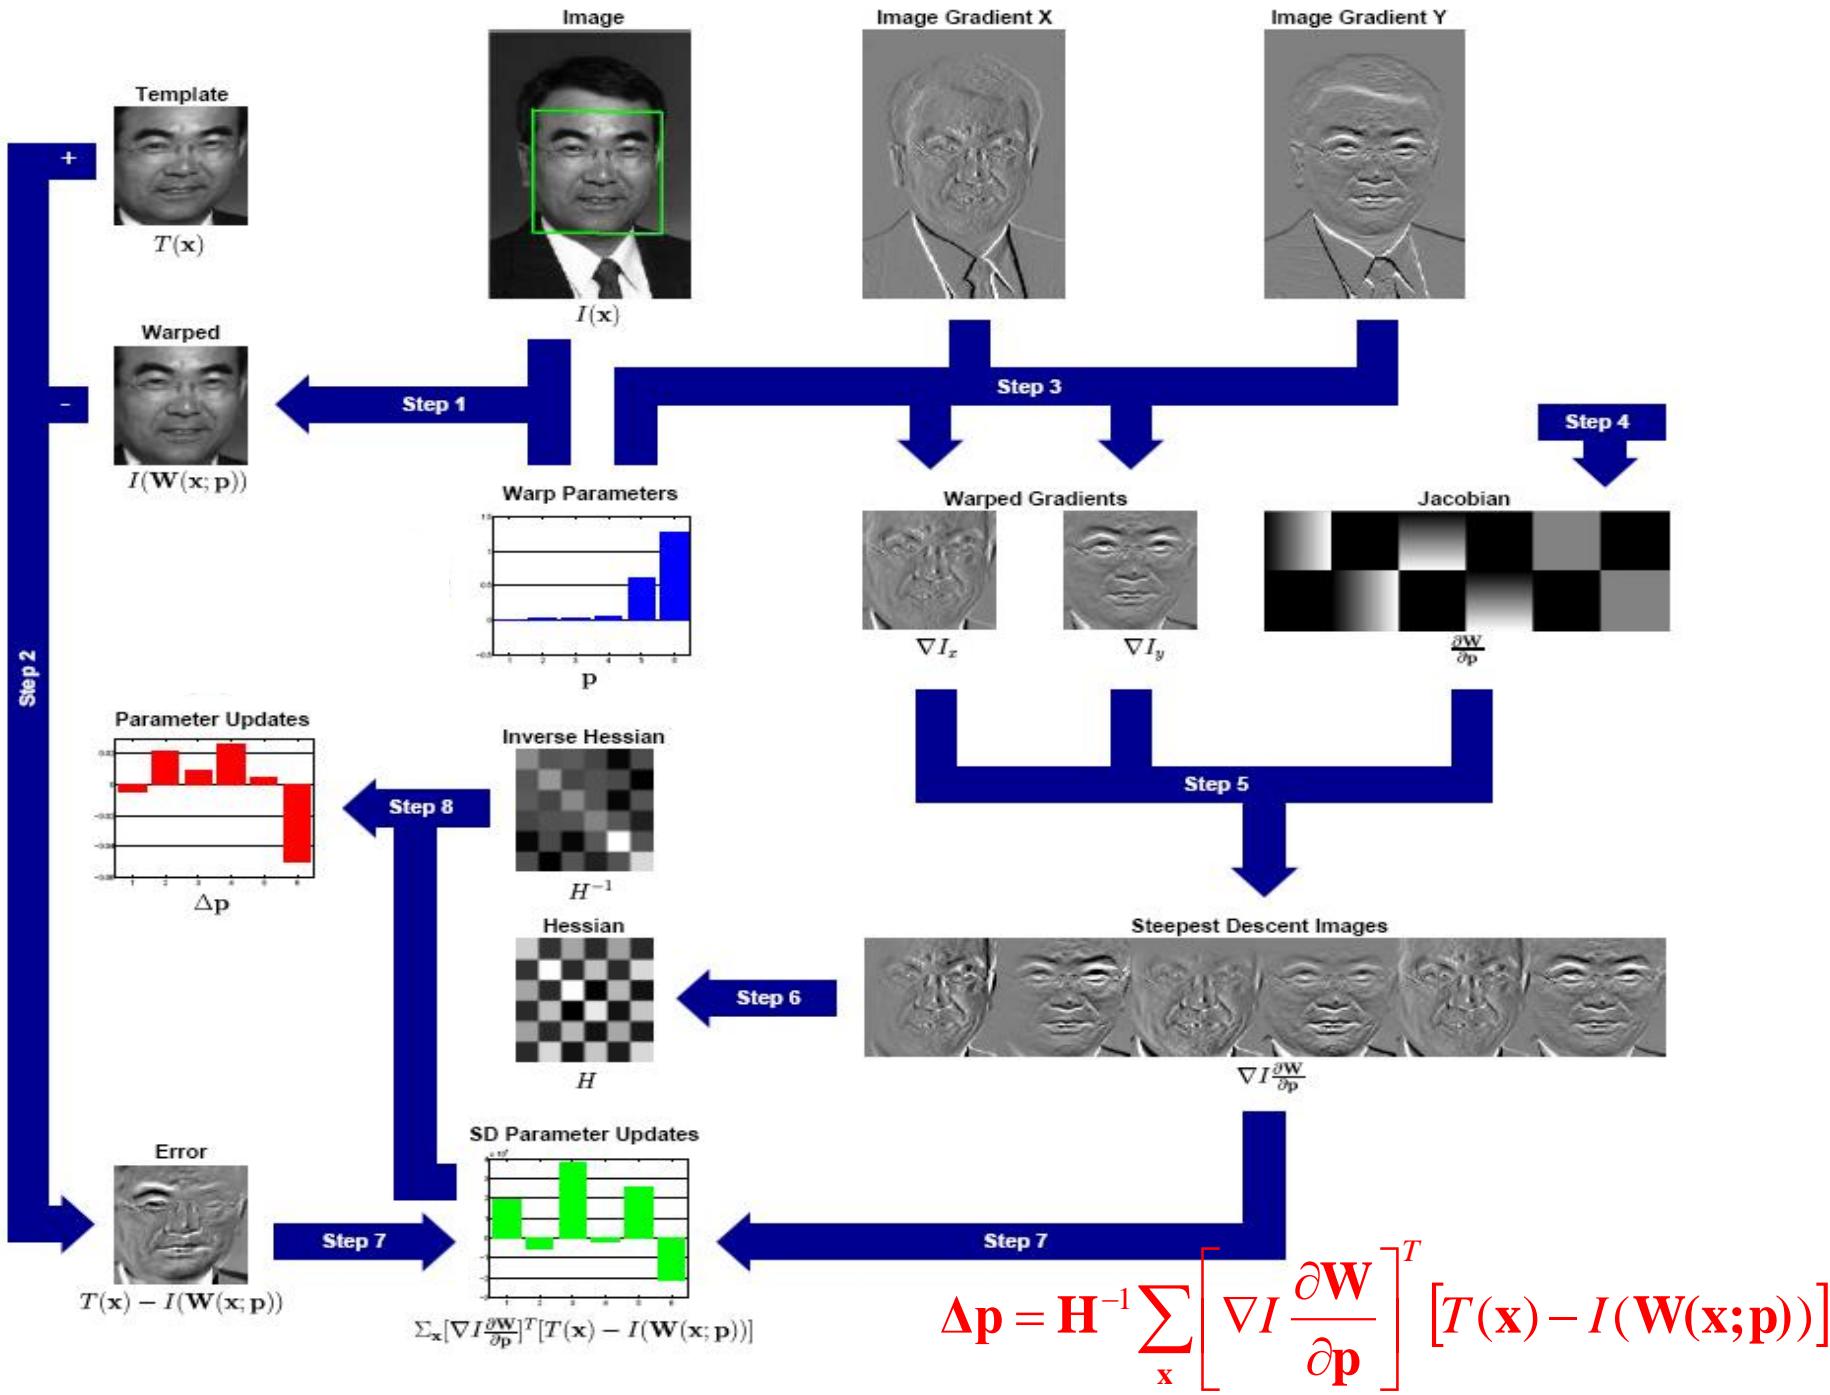
$$\Delta \mathbf{p} = \mathbf{H}^{-1} \sum_{\mathbf{x}} \left[\nabla I \frac{\partial \mathbf{W}}{\partial \mathbf{p}} \right]^T [T(\mathbf{x}) - I(\mathbf{W}(\mathbf{x}; \mathbf{p}))]$$

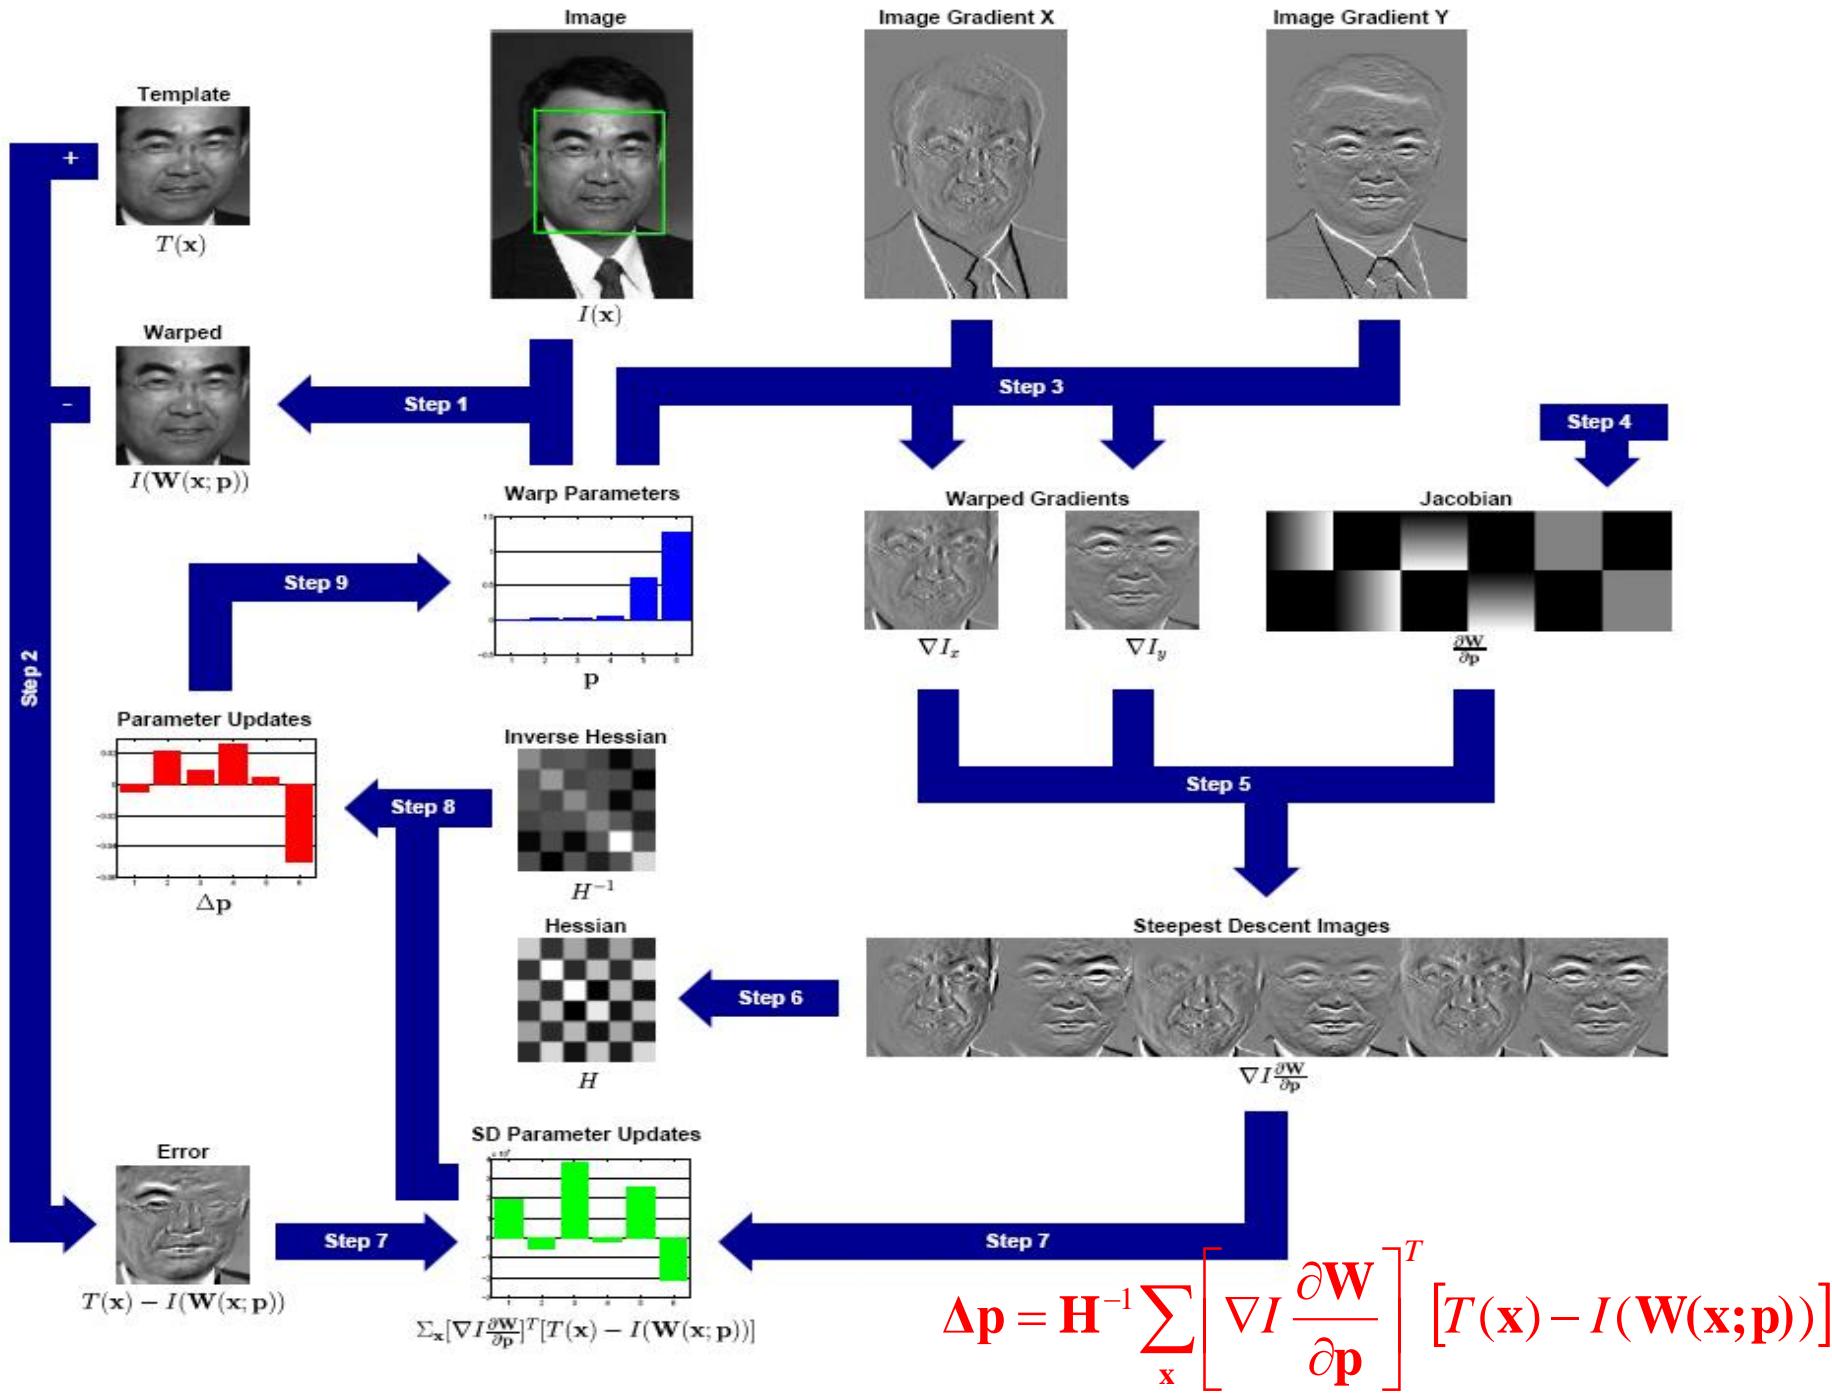




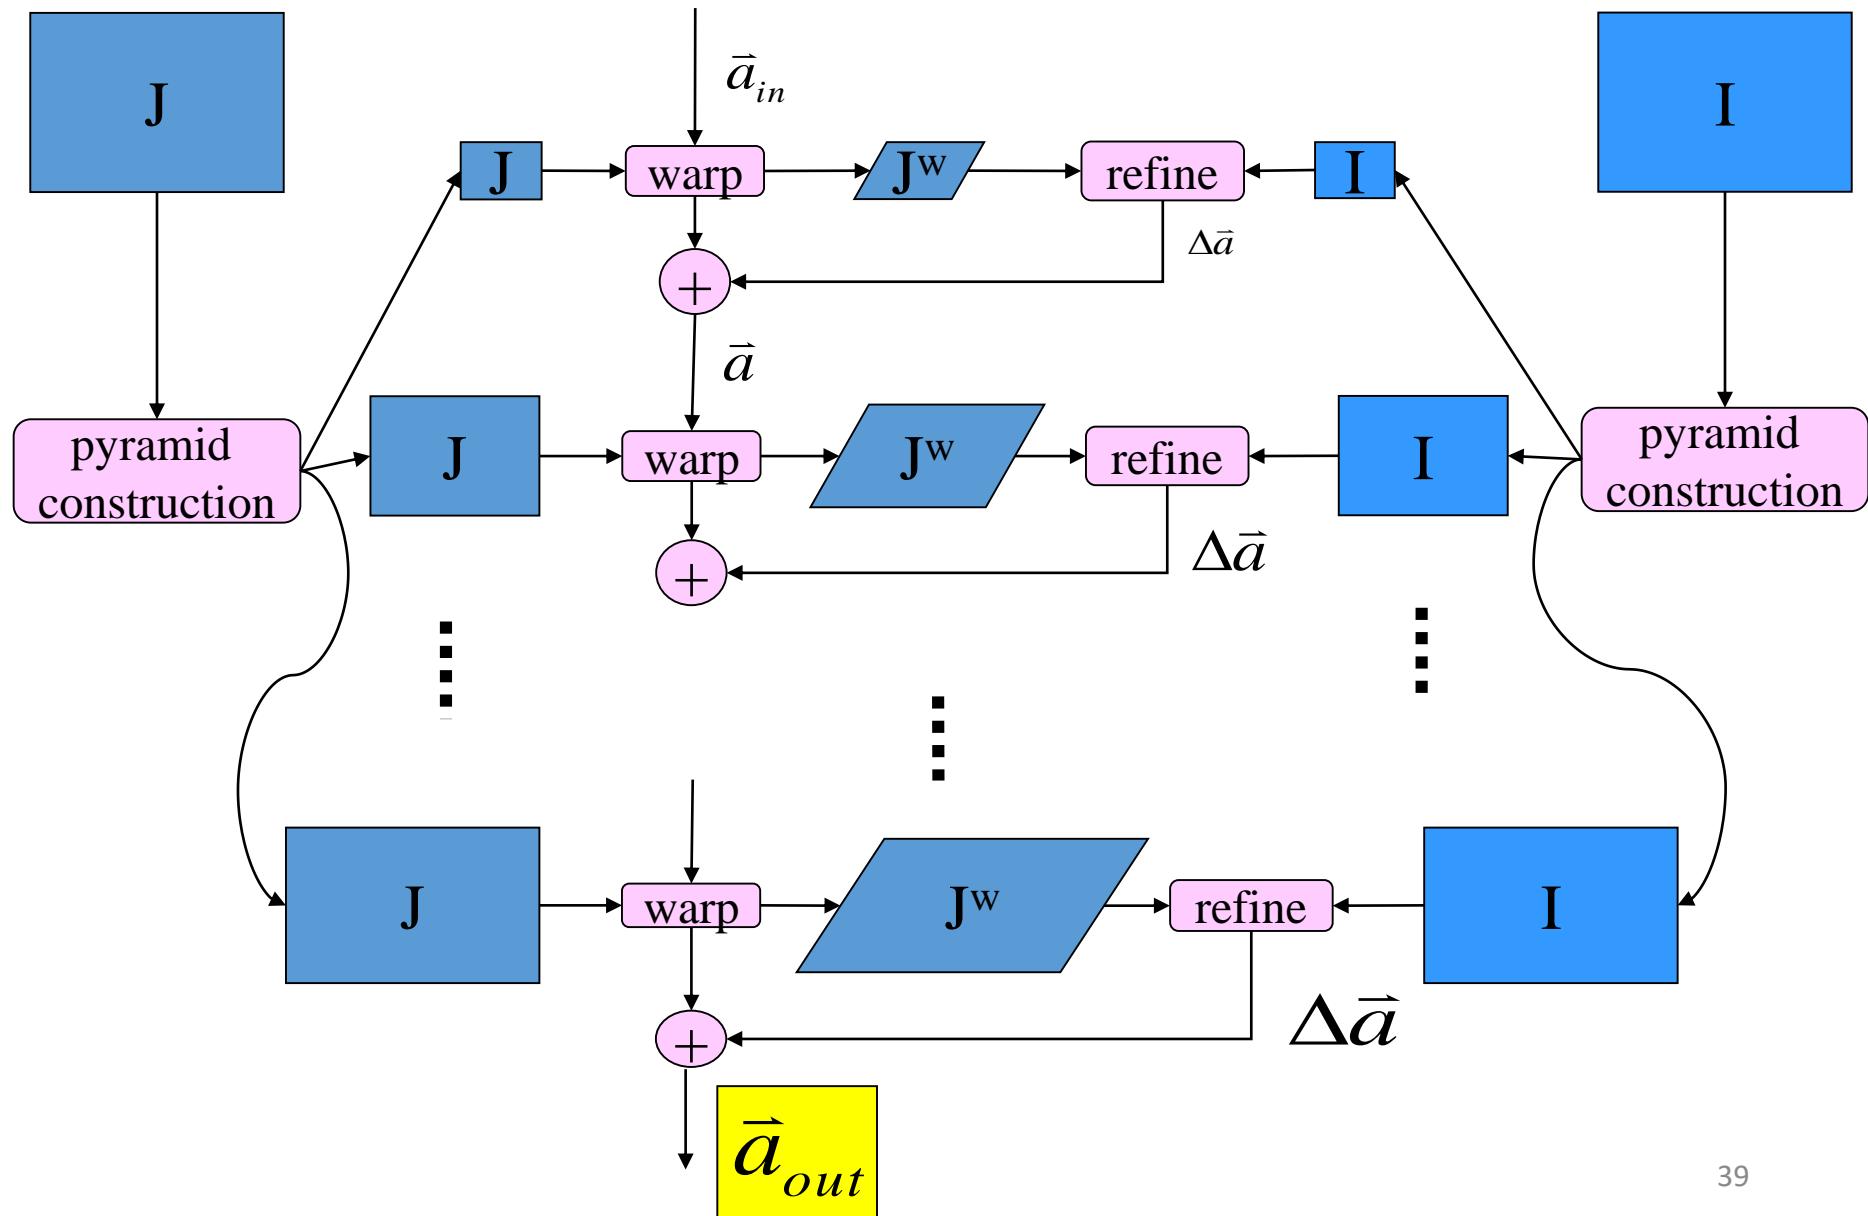
$$\Delta \mathbf{p} = \mathbf{H}^{-1} \sum_{\mathbf{x}} \left[\nabla I \frac{\partial \mathbf{W}}{\partial \mathbf{p}} \right]^T [T(\mathbf{x}) - I(\mathbf{W}(\mathbf{x}; \mathbf{p}))]$$







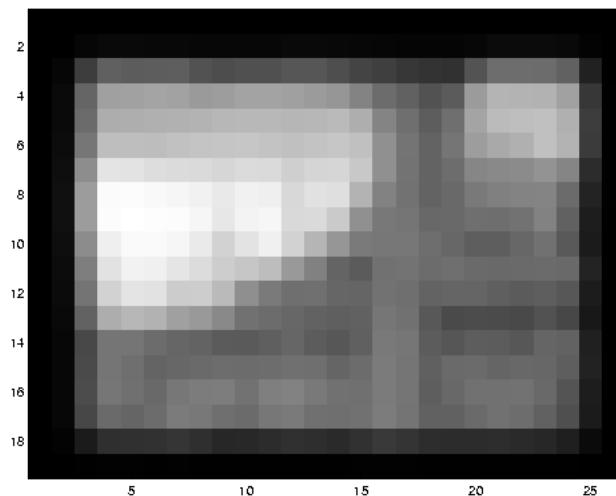
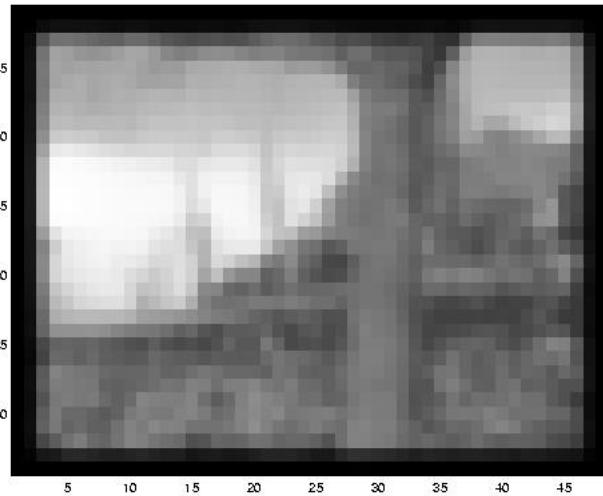
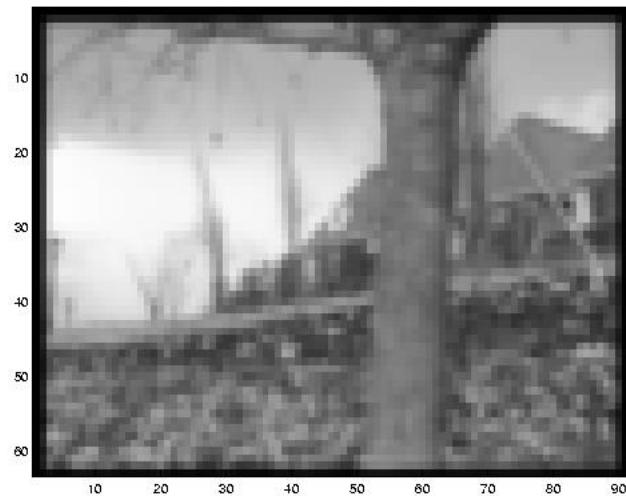
Coarse-to-Fine Strategy



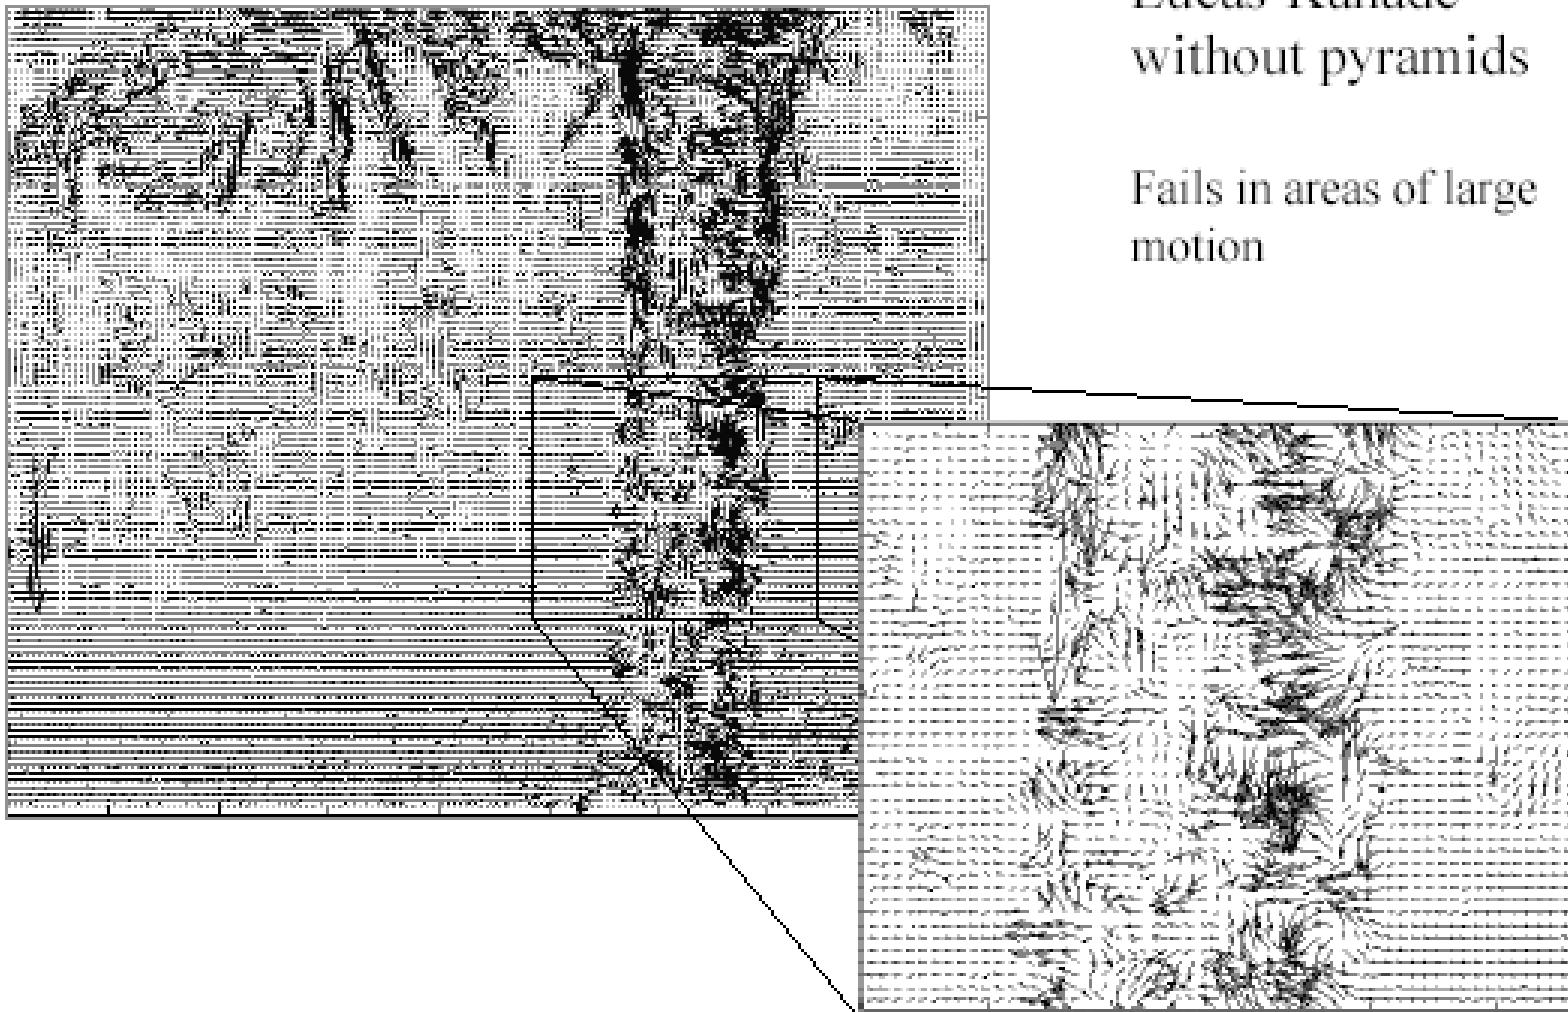
Example



Multi-resolution registration



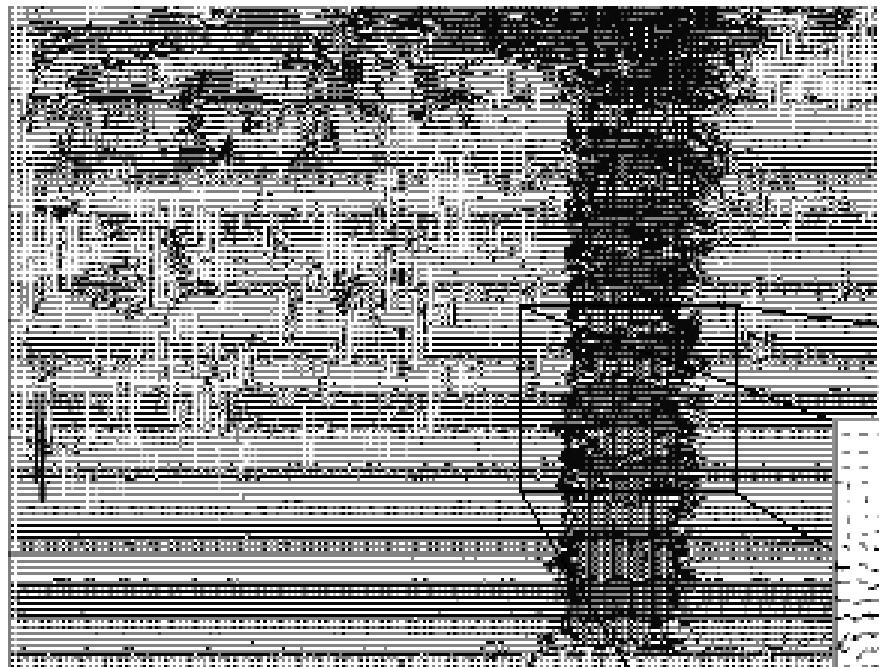
Optical Flow Results



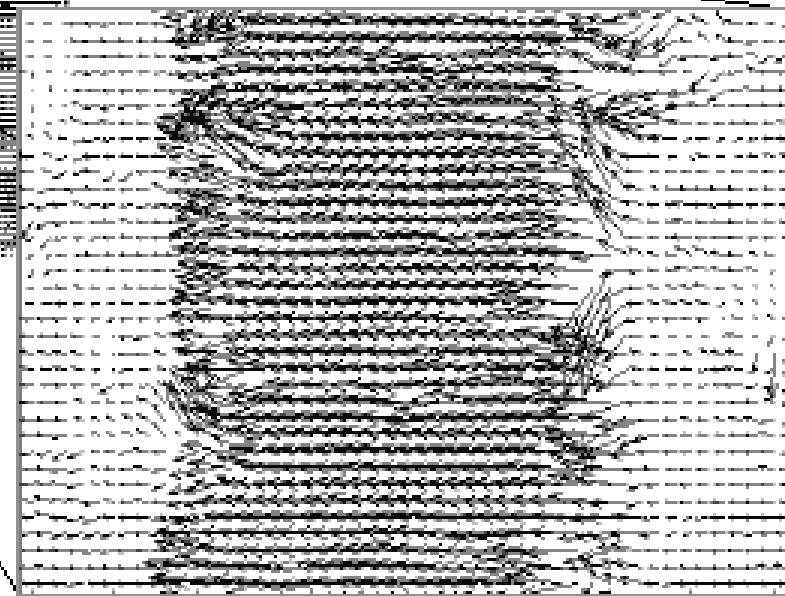
Lucas-Kanade
without pyramids

Fails in areas of large
motion

Optical Flow Results

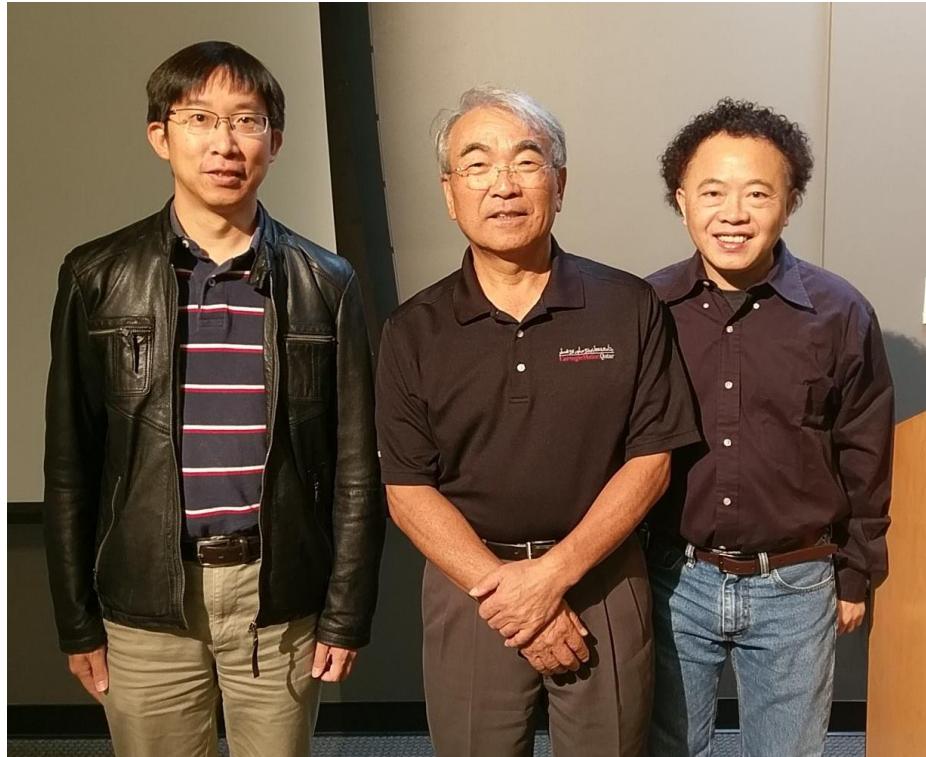


Lucas-Kanade with Pyramids



Lucas-Kanade Optical Flow

- <https://youtu.be/D7r3-fHXvRU?t=1h40m54s>



Errors of Lucas-Kanade

- The motion is large
 - Possible Fix: Keypoint matching
- A point does not move like its neighbors
 - Possible Fix: Region-based matching
- Brightness constancy does not hold
 - Possible Fix: Gradient constancy

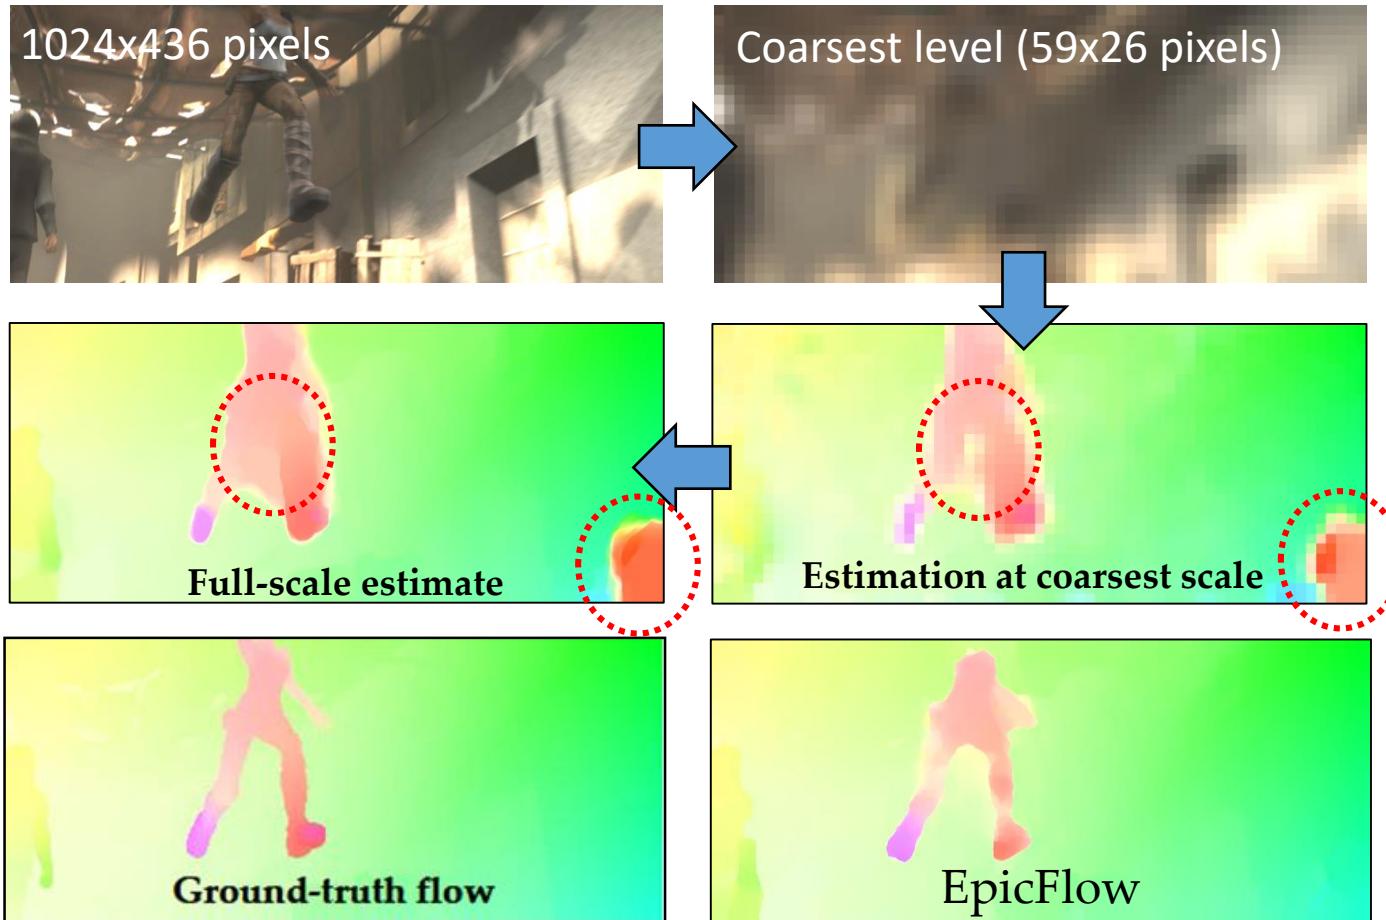
Epicflow

- Main remaining problems:
 - large displacements
 - occlusions
 - motion discontinuities
- **EpicFlow**
Epic: Edge-Preserving Interpolation of Correspondences
- leverages state-of-the-art matching algorithm
 - invariant to large displacements
- incorporate an edge-aware distance:
 - handles occlusions and motion discontinuities
- state-of-the-art results

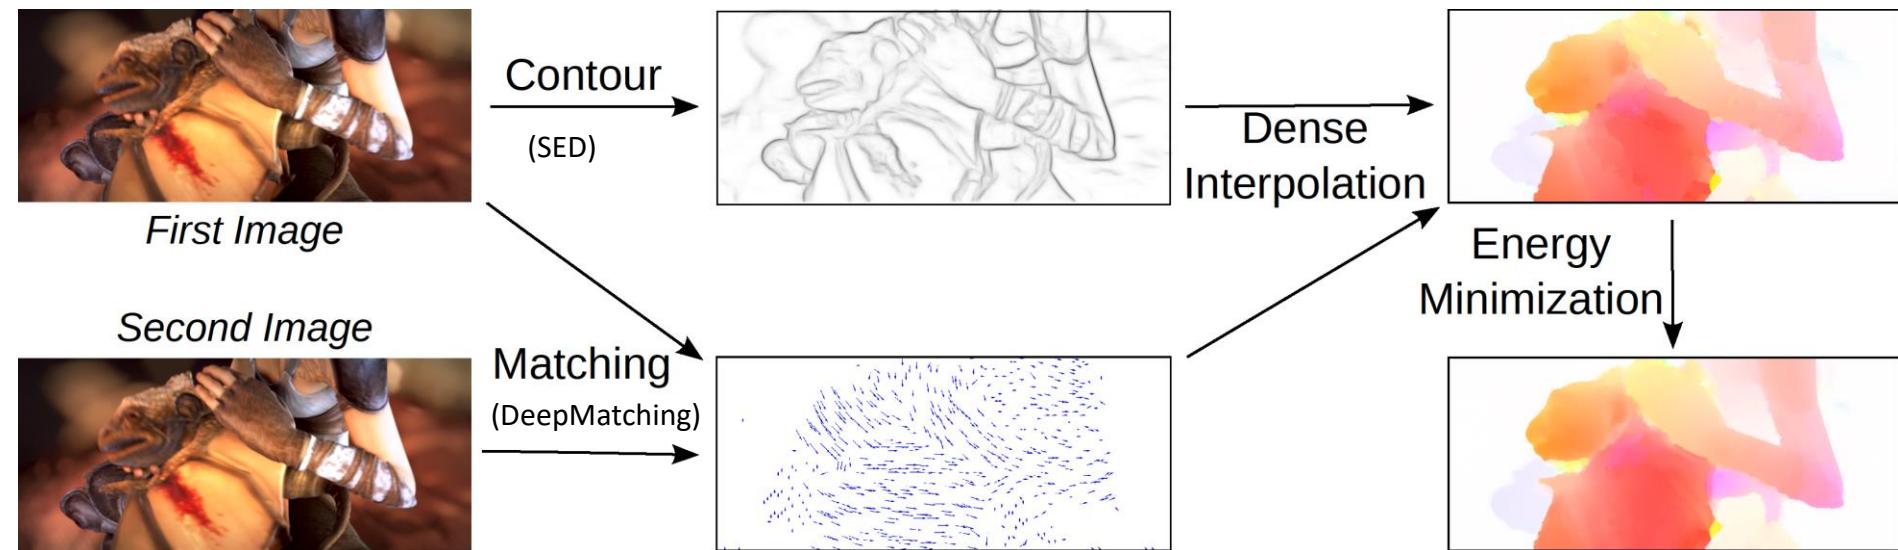
Ref: J. Revaud, P. Weinzaepfel, Z. Harchaoui and C. Schmid, “EpicFlow: Edge-Preserving Interpolation of Correspondences for Optical Flow,” *CVPR 2015*.

Epicflow

- Problems with coarse-to-fine:
 - flow discontinuities overlap at coarsest scales
 - errors are propagated across scales
 - no theoretical guarantees or proof of convergence!



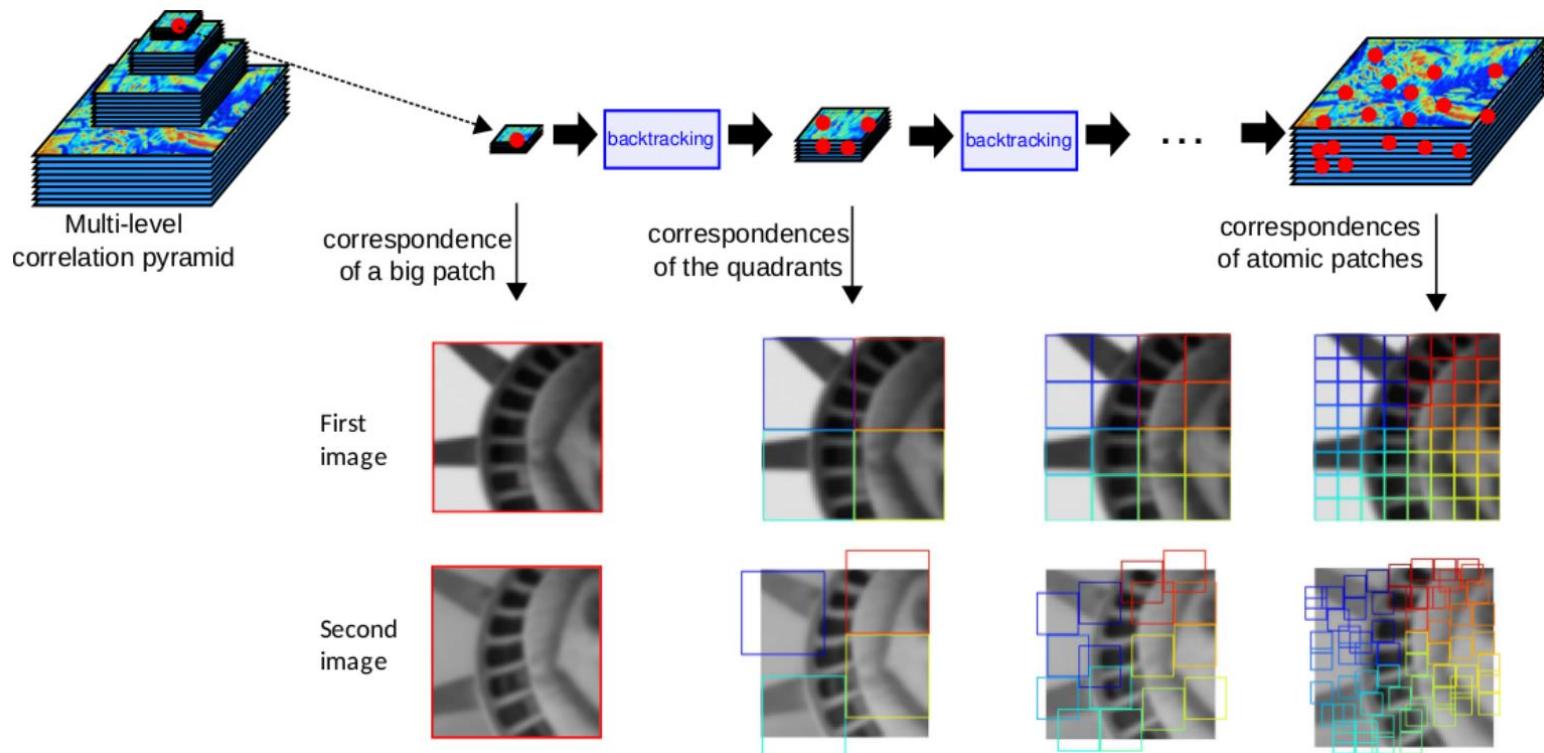
Epicflow: Overview



- Avoid coarse-to-fine scheme

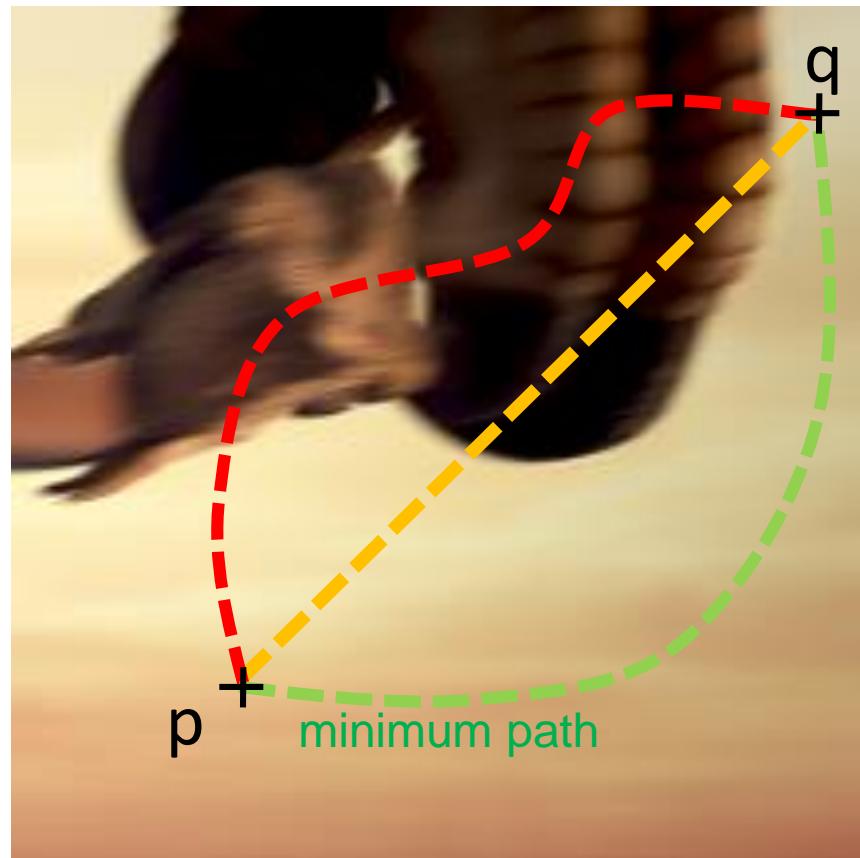
Epicflow: DeepMatching

- Based on correlation of SIFT features of a patch

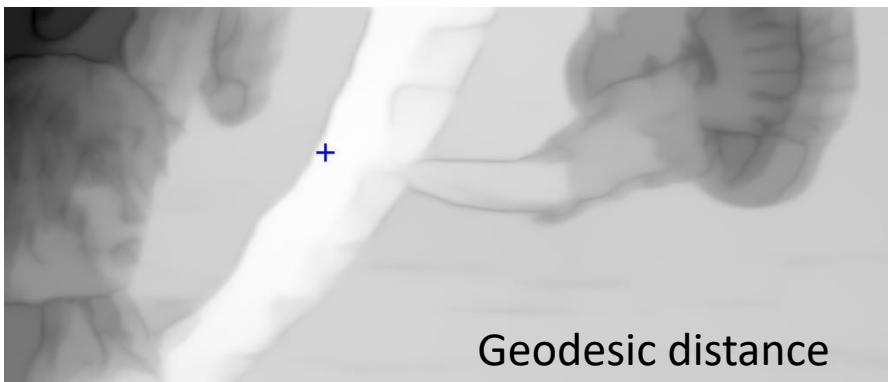
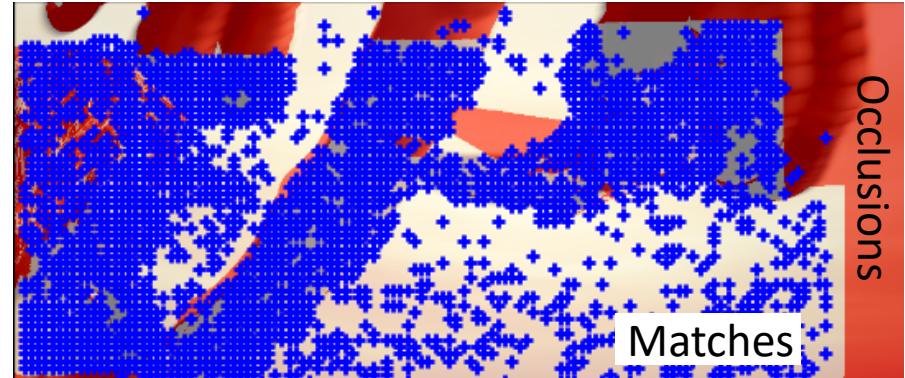


Epicflow: Dense Interpolation

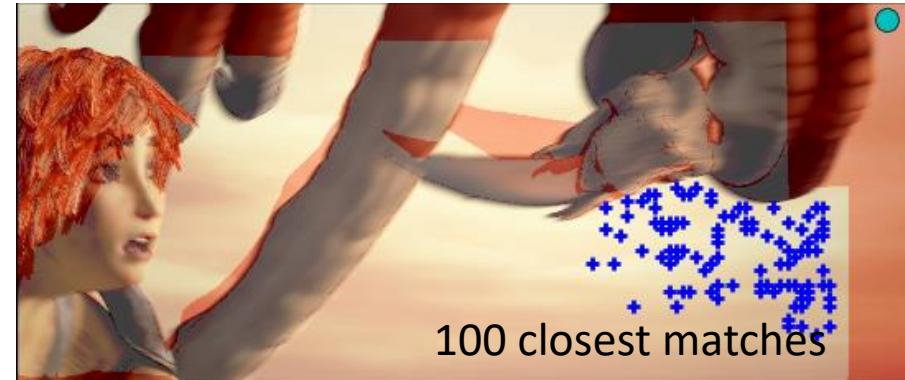
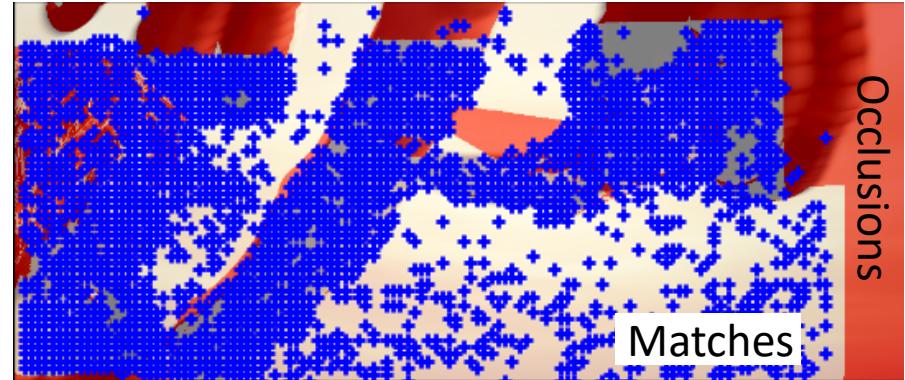
- Replace Euclidean distance with edge-aware distance to find NNs
- Geodesic distance:
 - Anisotropic cost map = **image edges**
 - ▶ Cost of a path: sum of the cost of all traversed pixels
 - ▶ Geodesic distance: minimum cost among all possible paths



Geodesic Distance



Geodesic Distance

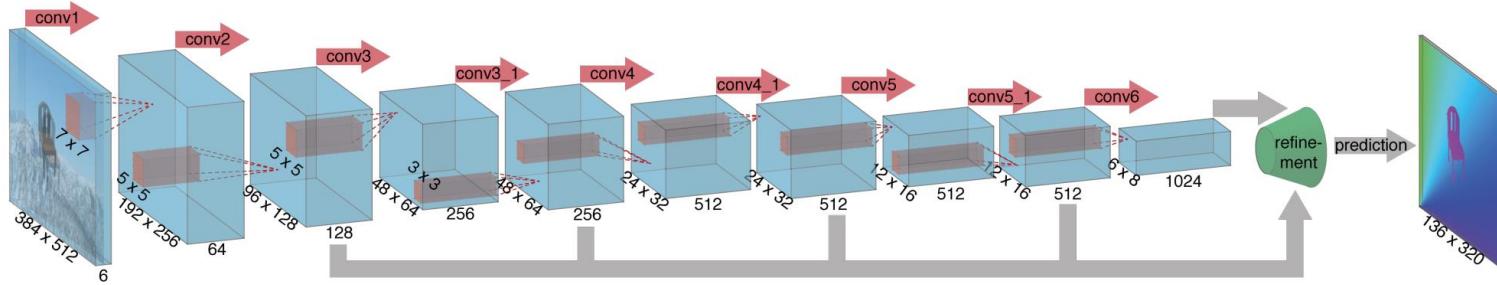


Experimental Results

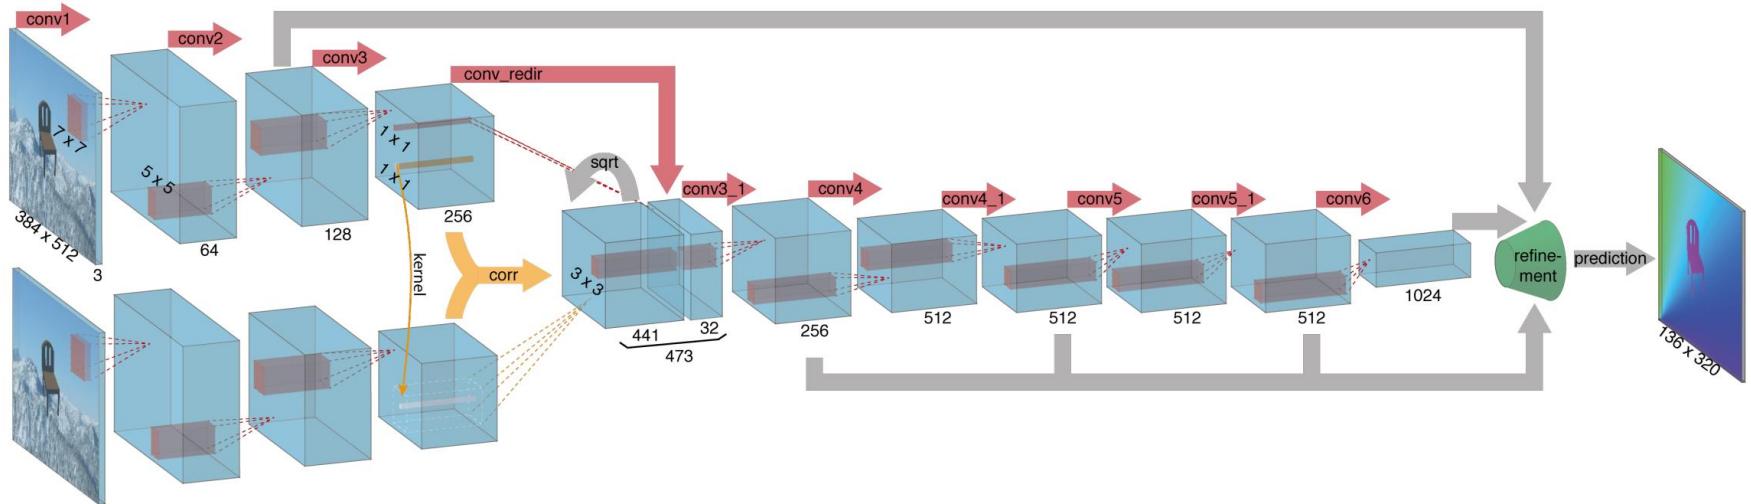


FlowNet

FlowNetSimple

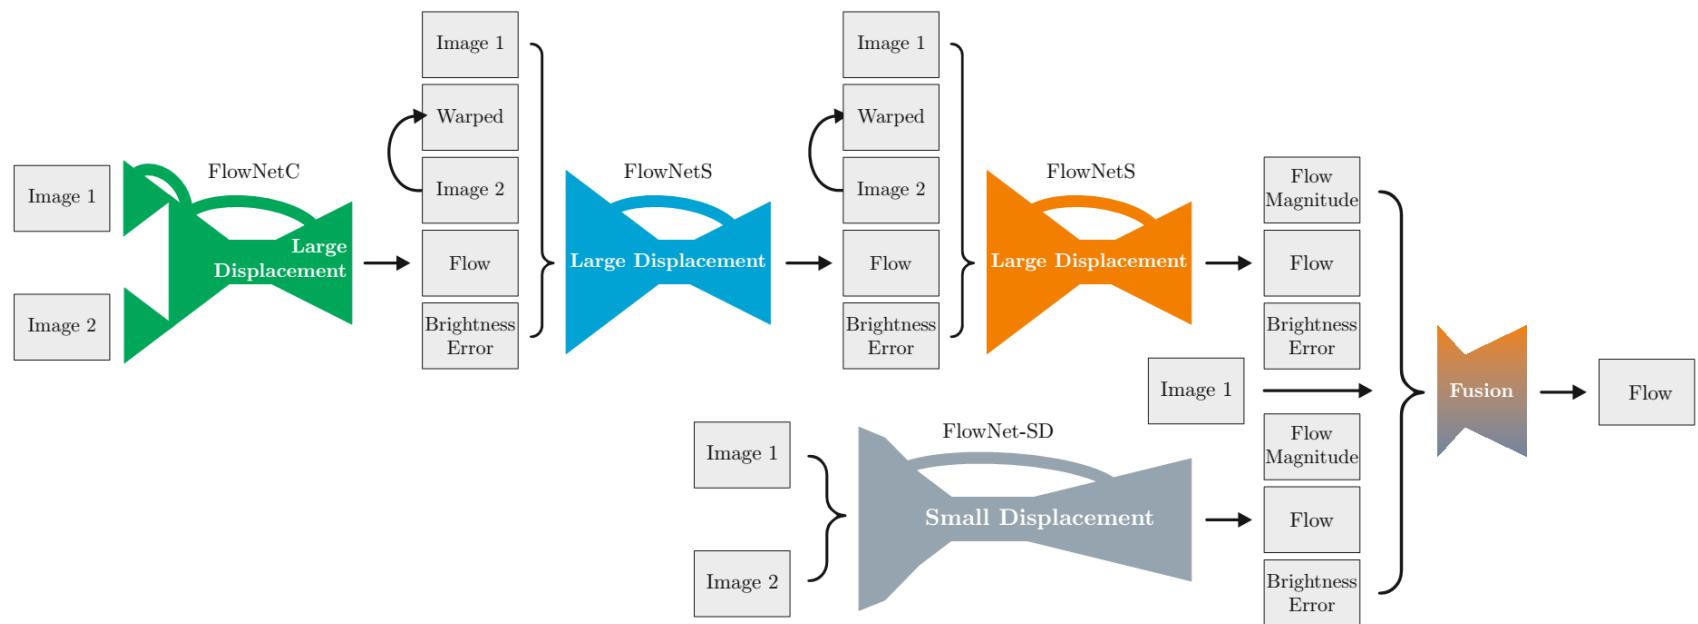


FlowNetCorr



Ref: A. Dosovitskiy, P. Fischer, E. Ilg, P. Hausser, C. Hazirbas, V. Golkov, P. van der Smagt, D. Cremers, T. Brox, "FlowNet: Learning Optical Flow With Convolutional Networks," CVPR2015.

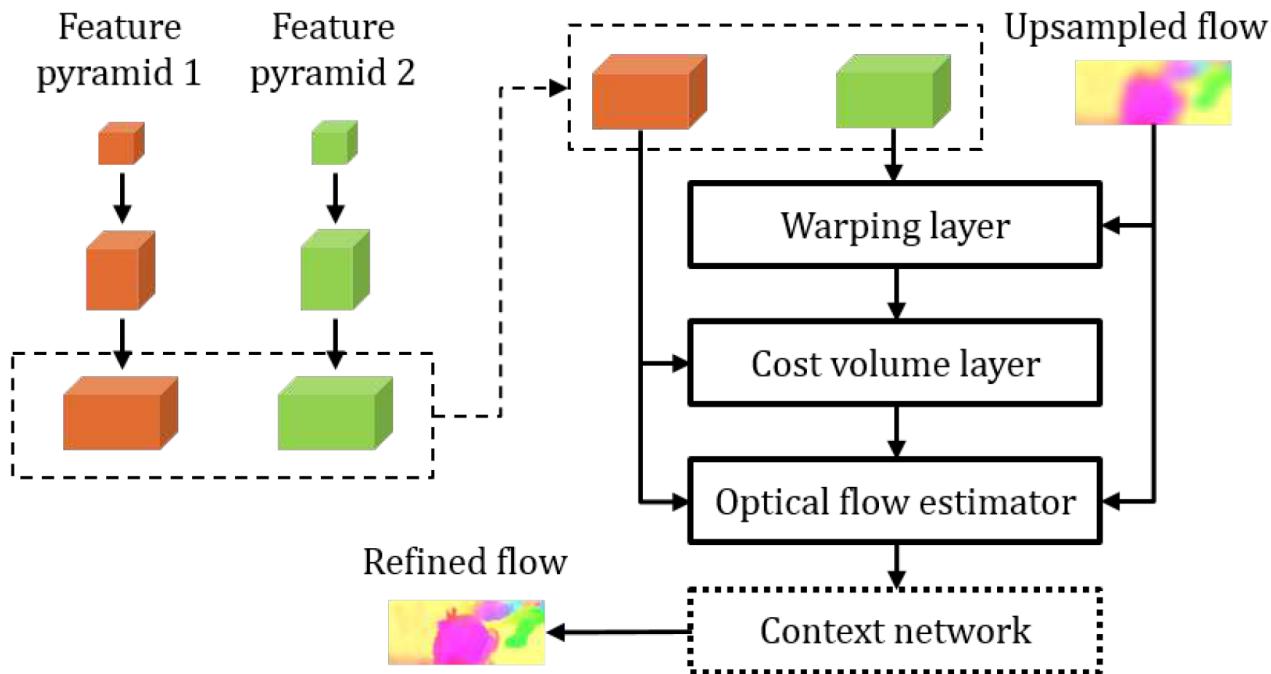
FlowNet 2.0



Ref: E. Ilg, N. Mayer, T. Saikia, M. Keuper, A. Dosovitskiy and T. Brox, "FlowNet 2.0: Evolution of Optical Flow Estimation with Deep Networks," *CVPR2017*.

PWC-Net

- Pyramid of features
- Warping features
- Cost volume with small search range

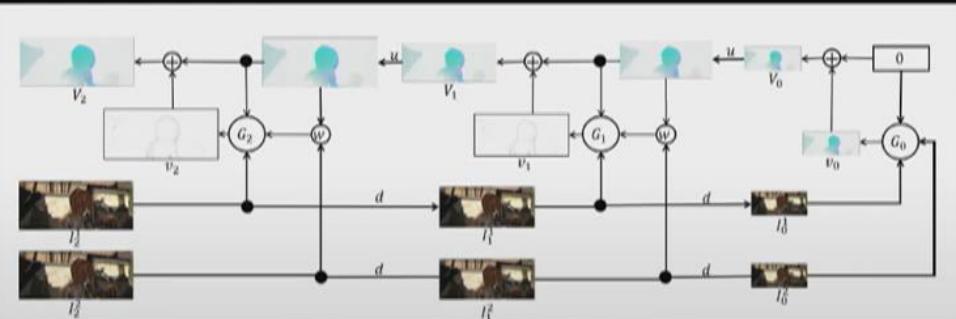


Ref: D. Sun, X. Yang, M.-Y. Liu, and J. Kautz, “PWC-Net: CNNs for optical flow using pyramid, warping, and cost volume,” *CVPR2018*.

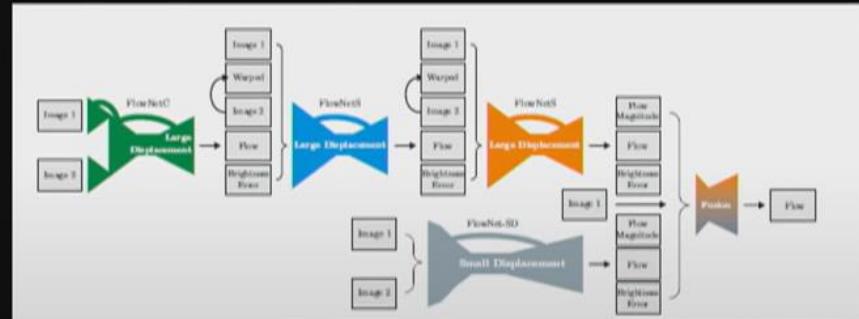
Comparison with Prior Work

Pyramid, Warping, and Cost volume

	SpyNet	FlowNet2	PWC-Net
Pyramid	Image	Feature (3-level)	Feature (6-level)
Warping	Image	Image	Feature
Cost volume	-	Once search range: 41x41	Every level search range: 9x9

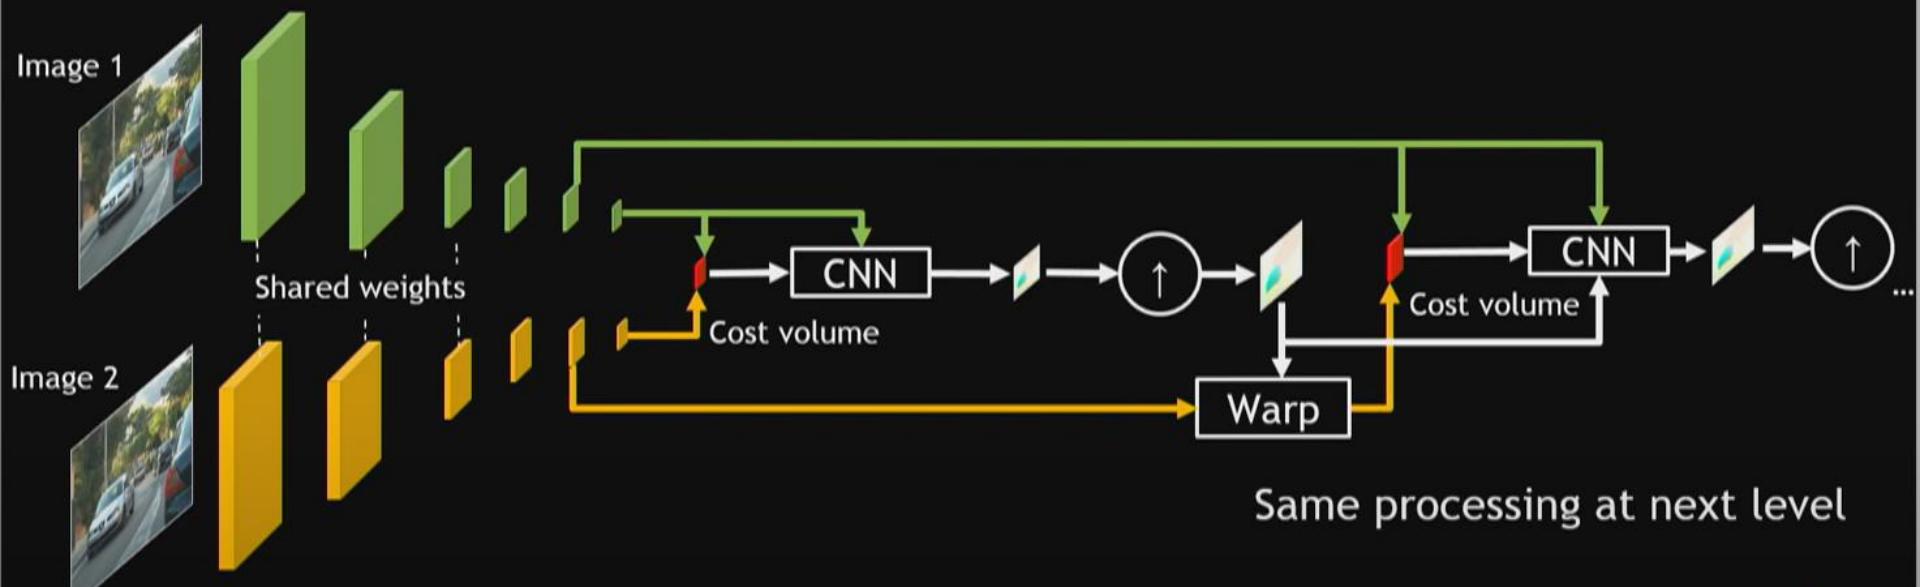


SpyNet: Ranjan & Black CVPR'17

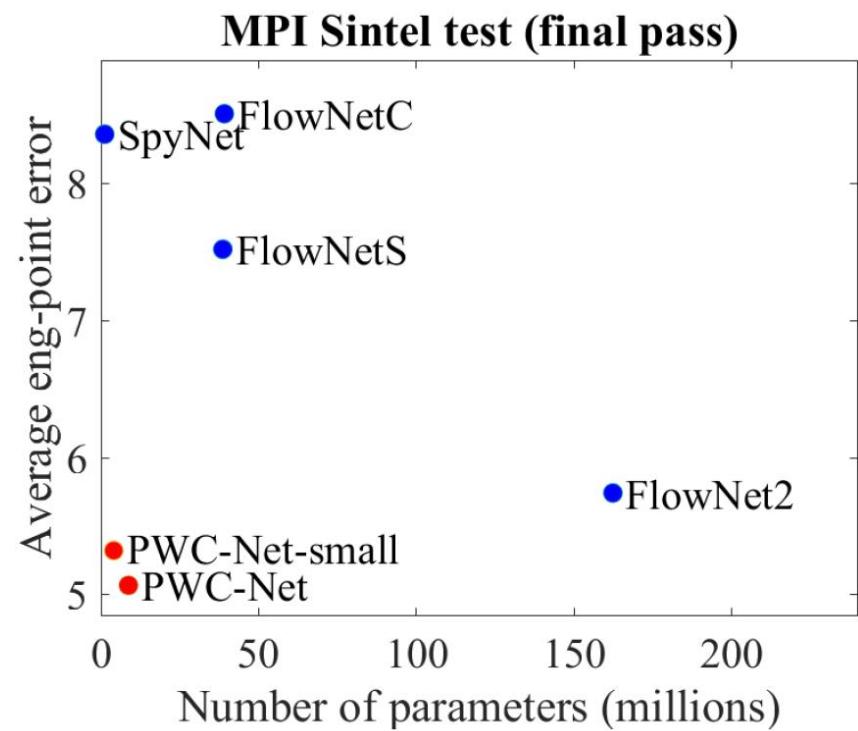
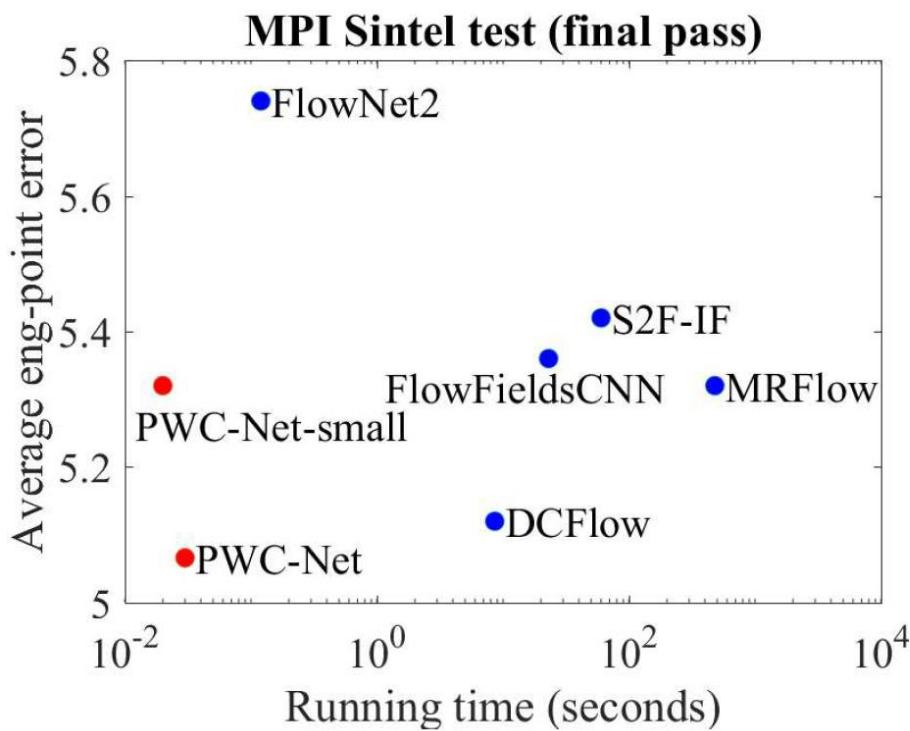


FlowNet2: Ilg *et al.* CVPR'17

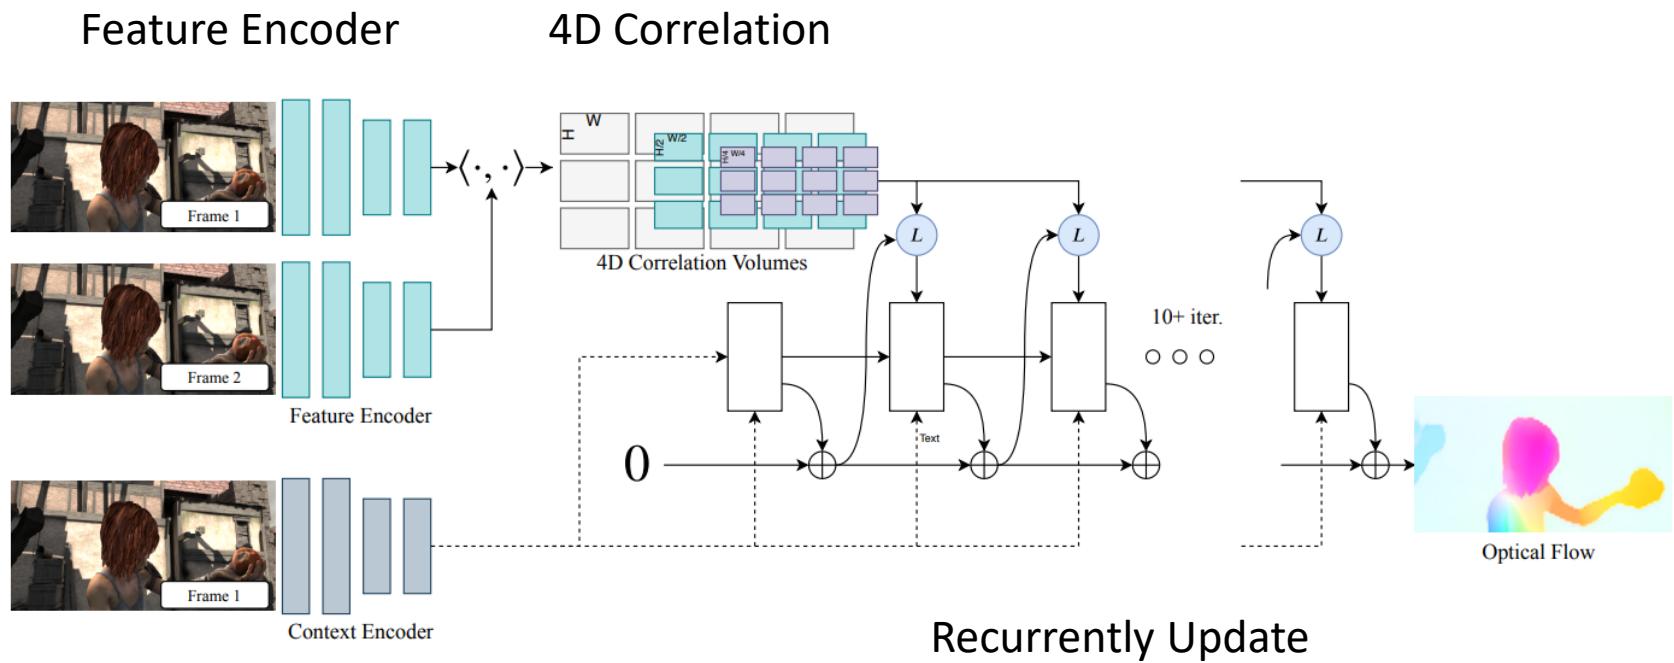
Warping Features



PWC-Net

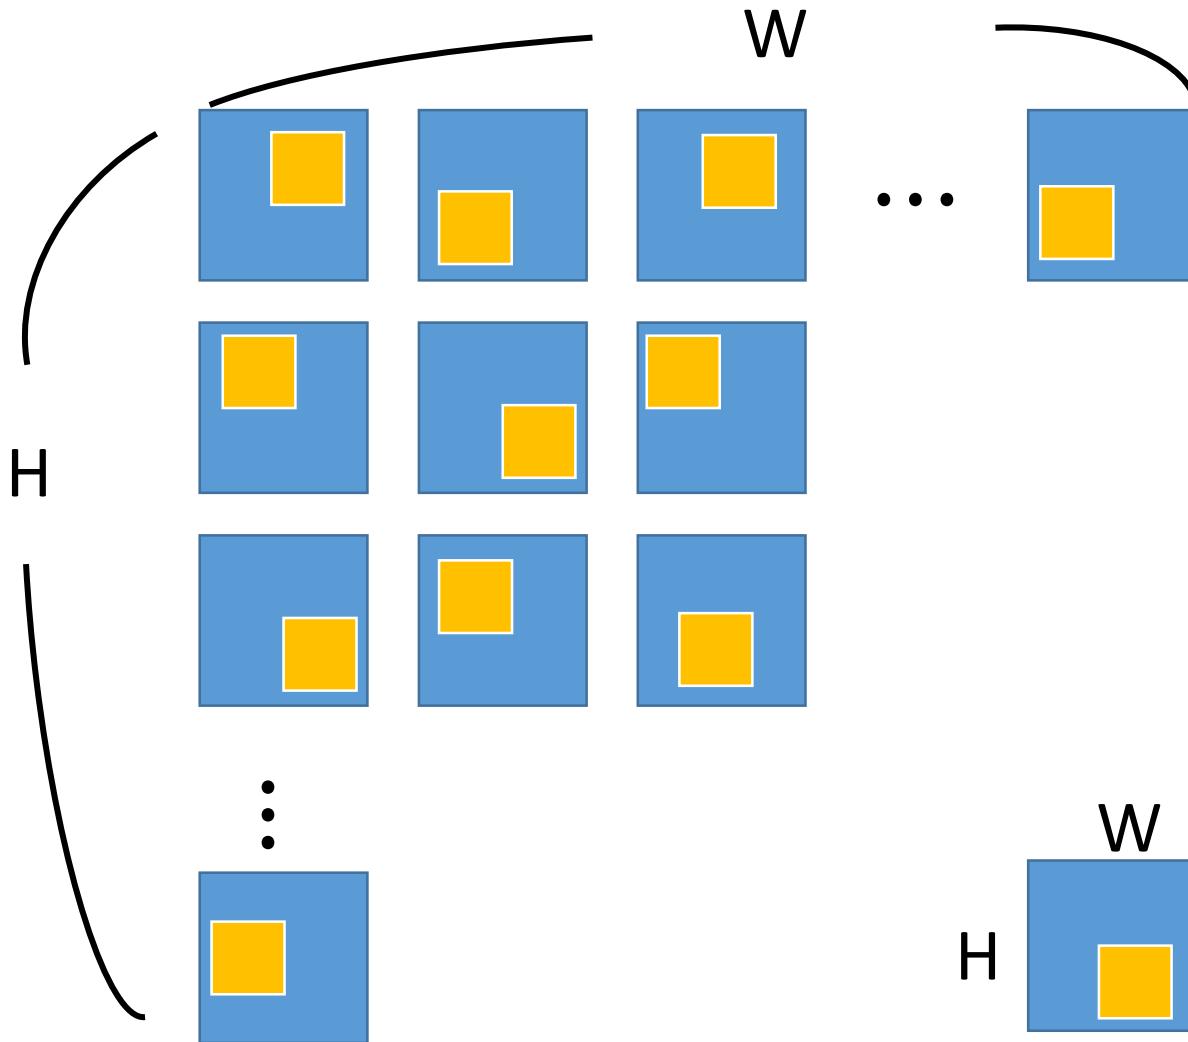


RAFT

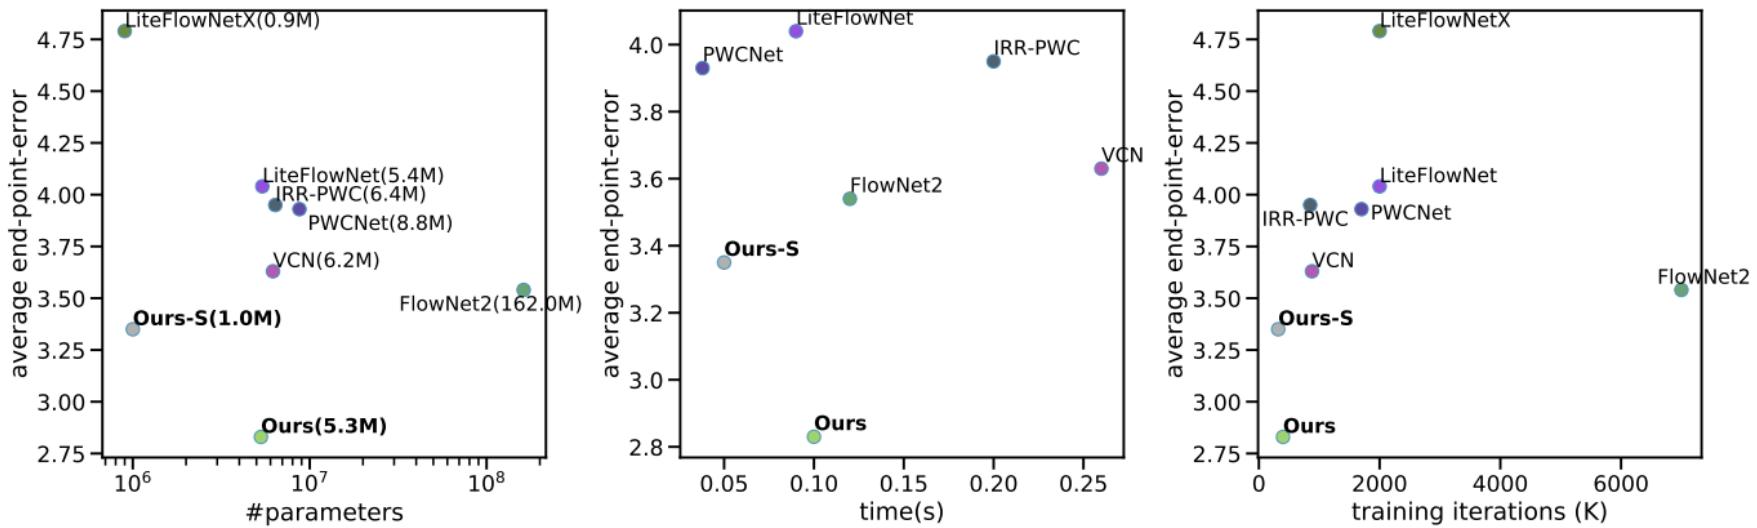


Ref: Z. Teed and J. Deng, "Recurrent All-Pairs Field Transforms for Optical Flow," in Proc. ECCV 2020.

4D Correlation



Experimental Results



Ref: Z. Teed and J. Deng, "Recurrent All-Pairs Field Transforms for Optical Flow," in Proc. ECCV 2020.

	Method	Setting	Code	Fl-bg	Fl-fg	Fl-all	Density	Runtime	Environment	Compare	
1	CamLiRAFT		code	2.08 %	7.37 %	2.96 %	100.00 %	1 s	GPU @ 2.5 Ghz (Python + C/C++)	<input type="checkbox"/>	
H. Liu, T. Lu, Y. Xu, J. Liu and L. Wang: Learning Optical Flow and Scene Flow with Bidirectional Camera-LiDAR Fusion . arXiv preprint arXiv:2303.12017 2023.											
2	CamLiFlow		code	2.31 %	7.04 %	3.10 %	100.00 %	1.2 s	GPU @ 2.5 Ghz (Python + C/C++)	<input type="checkbox"/>	
H. Liu, T. Lu, Y. Xu, J. Liu, W. Li and L. Chen: CamLiFlow: Bidirectional Camera-LiDAR Fusion for Joint Optical Flow and Scene Flow Estimation . CVPR 2022.											
3	TBD			2.90 %	5.05 %	3.26 %	100.00 %	TBD s	1 core @ 2.5 Ghz (C/C++)	<input type="checkbox"/>	
4	CamLiRAFT-NR		code	2.76 %	6.78 %	3.43 %	100.00 %	1 s	GPU @ 2.5 Ghz (Python + C/C++)	<input type="checkbox"/>	
H. Liu, T. Lu, Y. Xu, J. Liu and L. Wang: Learning Optical Flow and Scene Flow with Bidirectional Camera-LiDAR Fusion . arXiv preprint arXiv:2303.12017 2023.											
5	M-FUSE			code	2.66 %	7.47 %	3.46 %	100.00 %	1.3 s	GPU	<input type="checkbox"/>
L. Mehl, A. Jahedi, J. Schmalfuss and A. Bruhn: M-FUSE: Multi-frame Fusion for Scene Flow Estimation . Proc. Winter Conference on Applications of Computer Vision (WACV) 2023.											
6	RigidMask+ISF		code	2.63 %	7.85 %	3.50 %	100.00 %	3.3 s	GPU @ 2.5 Ghz (Python)	<input type="checkbox"/>	
G. Yang and D. Ramanan: Learning to Segment Rigid Motions from Two Frames . CVPR 2021.											
7	ScaleRAFT3D			2.37 %	9.26 %	3.51 %	100.00 %	1 s	1 core @ 2.5 Ghz (C/C++)	<input type="checkbox"/>	
8	TPCV+RAFT3D			2.48 %	10.19 %	3.76 %	100.00 %	0.2 s	1 core @ 2.5 Ghz (C/C++)	<input type="checkbox"/>	
9	RAFT3D+mscv			2.48 %	10.23 %	3.77 %	100.00 %	0.2 s	1 core @ 2.5 Ghz (C/C++)	<input type="checkbox"/>	
10	CMAX+			3.39 %	6.21 %	3.86 %	100.00 %	1.5 s	GPU @ 2.5 Ghz (Python)	<input type="checkbox"/>	
11	RAFT-it+ RVC		code	3.62 %	5.33 %	3.90 %	100.00 %	0.14 s	1 core @ 2.5 Ghz (Python)	<input type="checkbox"/>	
D. Sun, C. Herrmann, F. Reda, M. Rubinstein, D. Fleet and W. Freeman: Disentangling Architecture and Training for Optical Flow . ECCV 2022.											
12	RAFT-OCTC			3.72 %	5.39 %	4.00 %	100.00 %	0.2 s	GPU @ 2.5 Ghz (Python)	<input type="checkbox"/>	
J. Jeong, J. Lin, F. Porikli and N. Kwak: Imposing Consistency for Optical Flow Estimation (Qualcomm AI Research) . CVPR 2022.											
13	RCA-Flow			3.67 %	6.25 %	4.10 %	100.00 %	0.16 s	1 core @ 2.5 Ghz (Python)	<input type="checkbox"/>	
14	SF2SE3		code	3.17 %	8.79 %	4.11 %	100.00 %	2.7 s	GPU @ >3.5 Ghz (Python)	<input type="checkbox"/>	
L. Sommer, P. Schröppel and T. Brox: SF2SE3: Clustering Scene Flow into SE (3)-Motions via Proposal and Selection . DAGM German Conference on Pattern Recognition 2022.											
15	RAFT-CF-CE-PL3			3.80 %	5.65 %	4.11 %	100.00 %	0.05 s	GPU @ 2.5 Ghz (Python)	<input type="checkbox"/>	
16	RAFT-S-AF		code	3.86 %	5.38 %	4.12 %	100.00 %	1 s	1 core @ 2.5 Ghz (C/C++)	<input type="checkbox"/>	
17	MS RAFT+ corr RVC		code	3.83 %	5.71 %	4.15 %	100.00 %	0.65 s	GPU @ 2.5 Ghz (Python + C/C++)	<input type="checkbox"/>	
A. Jahedi, M. Luz, L. Mehl, M. Rivinius and A. Bruhn: High Resolution Multi-Scale RAFT . Robust Vision Challenge 2022, arXiv preprint arXiv:2210.16900 2022.											
A. Jahedi, L. Mehl, M. Rivinius and A. Bruhn: Multi-Scale Raft: Combining Hierarchical Concepts for Learning-Based Optical Flow Estimation . IEEE International Conference on Image Processing (ICIP) 2022.											
18	MS RAFT+ RVC			3.89 %	5.67 %	4.19 %	100.00 %	0.65 s	GPU @ 2.5 Ghz (Python + C/C++)	<input type="checkbox"/>	
19	DIP		code	3.86 %	5.96 %	4.21 %	100.00 %	0.15 s	1 core @ 2.5 Ghz (Python)	<input type="checkbox"/>	
Z. Zheng, N. Nie, Z. Ling, P. Xiong, J. Liu, H. Wang and J. Li: DIP: Deep Inverse Patchmatch for High- Resolution Optical Flow . Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition 2022.											
20	APCAFlow			3.86 %	6.14 %	4.24 %	100.00 %	0.2 s	1 core @ 2.5 Ghz (Python)	<input type="checkbox"/>	

Outline

- Tracking
 - Introduction
 - Single camera tracking
 - Feature tracker, KLT tracker
 - Region area tracker
 - Mean-shift
 - Tracking with adaptive filters
 - On-line learning
 - Multi-camera tracking

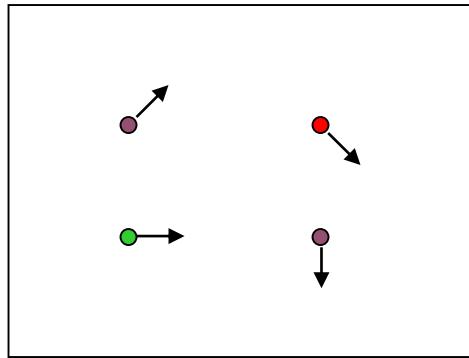
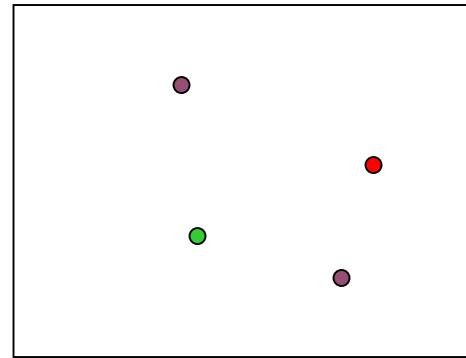
Introduction to Tracking



Introduction to Tracking

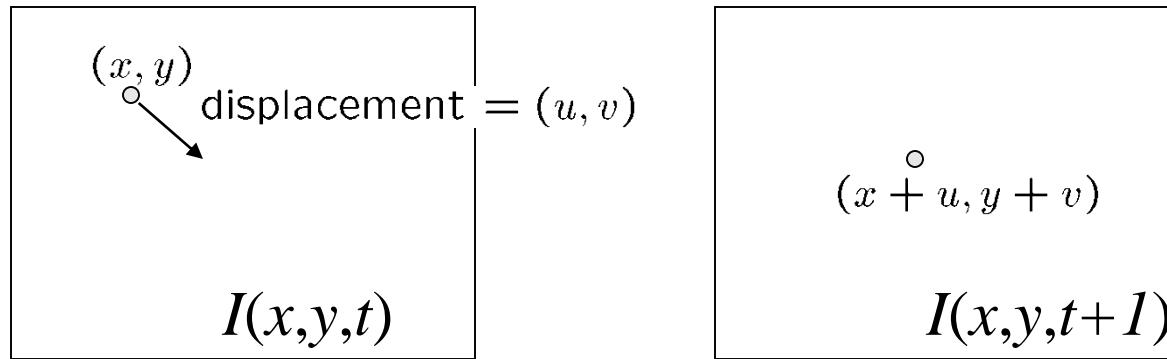
- Input: target
- Objective: Estimate target state over time
- State:
 - Position
 - Appearance
 - Shape
 - Velocity
 - ...
- It is about data association, similarity measurement, and correlation

Feature Tracking

 $I(x,y,t)$  $I(x,y,t+1)$

- Given two subsequent frames, estimate the point translation
- Key assumptions of Lucas-Kanade Tracker
 - **Brightness constancy:** projection of the same point looks the same in every frame
 - **Small motion:** points do not move very far
 - **Spatial coherence:** points move like their neighbors

The brightness constancy constraint



- Brightness Constancy Equation:

$$I(x, y, t) = I(x + u, y + v, t + 1)$$

Take Taylor expansion of $I(x + u, y + v, t + 1)$ at (x, y, t) to linearize the right side:

$$I(x + u, y + v, t + 1) \approx I(x, y, t) + \boxed{I_x} \cdot u + I_y \cdot v + \boxed{I_t}$$

Image derivative along x Difference over frames

$$I(x + u, y + v, t + 1) - I(x, y, t) = I_x \cdot u + I_y \cdot v + I_t$$

So: $I_x \cdot u + I_y \cdot v + I_t \approx 0 \rightarrow \nabla I \cdot [u \ v]^T + I_t = 0$

The Brightness Constancy Constraint

Can we use this equation to recover image motion (u, v) at each pixel?

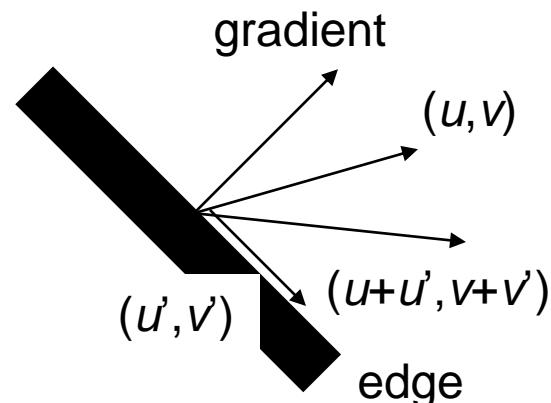
$$\nabla I \cdot [u \ v]^T + I_t = 0$$

- How many equations and unknowns per pixel?
 - One equation (this is a scalar equation!), two unknowns (u, v)

The component of the motion perpendicular to the gradient (i.e., parallel to the edge) cannot be measured

If (u, v) satisfies the equation,
so does $(u+u', v+v')$ if

$$\nabla I \cdot [u' \ v']^T = 0$$



Solving the ambiguity...

B. Lucas and T. Kanade. An iterative image registration technique with an application to stereo vision.
In *Proceedings of the International Joint Conference on Artificial Intelligence*, pp. 674–679, 1981.

- How to get more equations for a pixel?
- **Spatial coherence constraint**
 - Assume the pixel's neighbors have the same (u,v)
 - If we use a 5x5 window, that gives us 25 equations per pixel

$$0 = I_t(\mathbf{p}_i) + \nabla I(\mathbf{p}_i) \cdot [u \ v]$$

$$\begin{bmatrix} I_x(\mathbf{p}_1) & I_y(\mathbf{p}_1) \\ I_x(\mathbf{p}_2) & I_y(\mathbf{p}_2) \\ \vdots & \vdots \\ I_x(\mathbf{p}_{25}) & I_y(\mathbf{p}_{25}) \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = - \begin{bmatrix} I_t(\mathbf{p}_1) \\ I_t(\mathbf{p}_2) \\ \vdots \\ I_t(\mathbf{p}_{25}) \end{bmatrix}$$

Solving the ambiguity...

- Least squares problem:

$$\begin{bmatrix} I_x(p_1) & I_y(p_1) \\ I_x(p_2) & I_y(p_2) \\ \vdots & \vdots \\ I_x(p_{25}) & I_y(p_{25}) \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = - \begin{bmatrix} I_t(p_1) \\ I_t(p_2) \\ \vdots \\ I_t(p_{25}) \end{bmatrix}$$

$A \quad d = b$
 $25 \times 2 \quad 2 \times 1 \quad 25 \times 1$

Matching patches across images

- Overconstrained linear system

$$\begin{bmatrix} I_x(p_1) & I_y(p_1) \\ I_x(p_2) & I_y(p_2) \\ \vdots & \vdots \\ I_x(p_{25}) & I_y(p_{25}) \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = - \begin{bmatrix} I_t(p_1) \\ I_t(p_2) \\ \vdots \\ I_t(p_{25}) \end{bmatrix}$$

$A \quad d = b$
 $25 \times 2 \quad 2 \times 1 \quad 25 \times 1$

Least squares solution for d given by $(A^T A)^{-1} A^T b$

$$\begin{bmatrix} \sum I_x I_x & \sum I_x I_y \\ \sum I_x I_y & \sum I_y I_y \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = - \begin{bmatrix} \sum I_x I_t \\ \sum I_y I_t \end{bmatrix}$$

$A^T A \qquad \qquad \qquad A^T b$

The summations are over all pixels in the $K \times K$ window

Conditions for solvability

Optimal (u, v) satisfies Lucas-Kanade equation

$$\begin{bmatrix} \sum I_x I_x & \sum I_x I_y \\ \sum I_x I_y & \sum I_y I_y \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = - \begin{bmatrix} \sum I_x I_t \\ \sum I_y I_t \end{bmatrix}$$

$A^T A$ $A^T b$

When is this solvable? I.e., what are good points to track?

- $\mathbf{A}^T \mathbf{A}$ should be invertible
- $\mathbf{A}^T \mathbf{A}$ should not be too small due to noise
 - eigenvalues λ_1 and λ_2 of $\mathbf{A}^T \mathbf{A}$ should not be too small
- $\mathbf{A}^T \mathbf{A}$ should be well-conditioned
 - λ_1 / λ_2 should not be too large (λ_1 = larger eigenvalue)

Does this remind you of anything?

Criteria for Harris corner detector

Low-texture region



$$\sum \nabla I (\nabla I)^T$$

- gradients have small magnitude
- small λ_1 , small λ_2

Edge



$$\sum \nabla I (\nabla I)^T$$

- gradients very large or very small
- large λ_1 , small λ_2

High-texture region



$$\sum \nabla I (\nabla I)^T$$

- gradients are different, large magnitudes
- large λ_1 , large λ_2

Dealing with Larger Movements: Iterative Refinement

1. Initialize $(x', y') = (x, y)$

Original (x, y) position

2. Compute (u, v) by

$$\begin{bmatrix} \sum I_x I_x & \sum I_x I_y \\ \sum I_x I_y & \sum I_y I_y \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = - \begin{bmatrix} \sum I_x I_t \\ \sum I_y I_t \end{bmatrix}$$

2nd moment matrix for feature
patch in first image

displacement

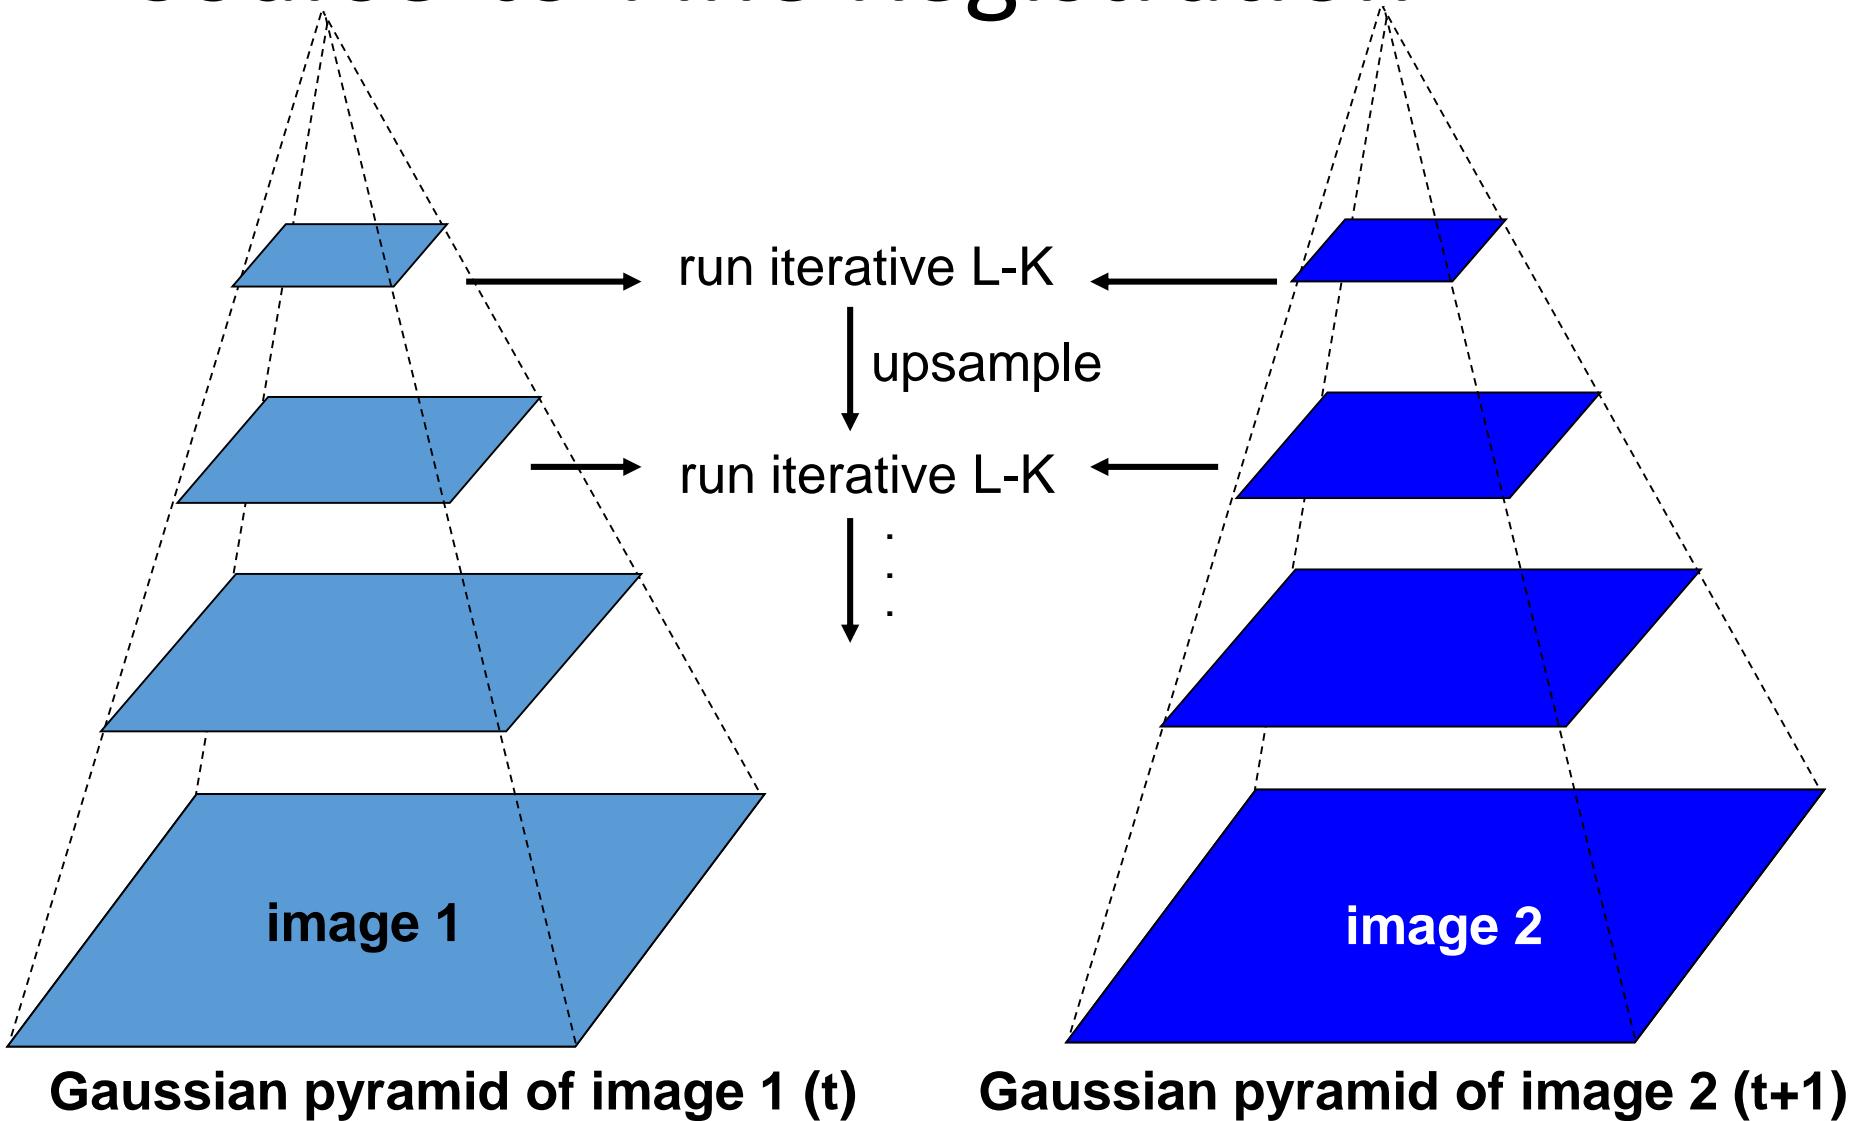
3. Shift window by (u, v) : $x' = x' + u$; $y' = y' + v$;

4. Recalculate I_t

5. Repeat steps 2-4 until small change

- Use interpolation for subpixel values

Dealing with Larger Movements: Coarse-to-Fine Registration



Shi-Tomasi Feature Tracker

- Find good features using eigenvalues of second-moment matrix (e.g., Harris detector or threshold on the smallest eigenvalue)
 - Key idea: “good” features to track are the ones whose motion can be estimated reliably
- Track from frame to frame with Lucas-Kanade
 - This amounts to assuming a translation model for frame-to-frame feature movement
- Check consistency of tracks by *affine* registration to the first observed instance of the feature
 - Affine model is more accurate for larger displacements
 - Comparing to the first frame helps to minimize drift

Tracking example

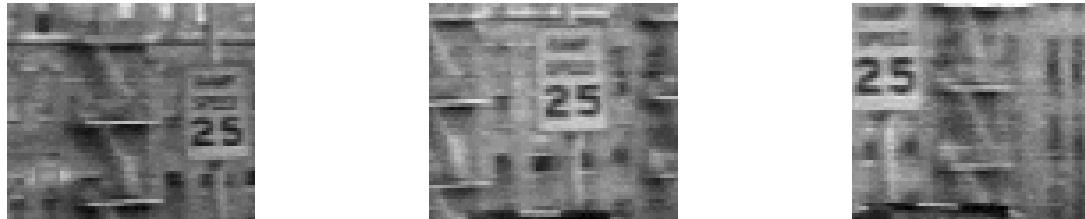


Figure 1: Three frame details from Woody Allen's *Manhattan*. The details are from the 1st, 11th, and 21st frames of a subsequence from the movie.

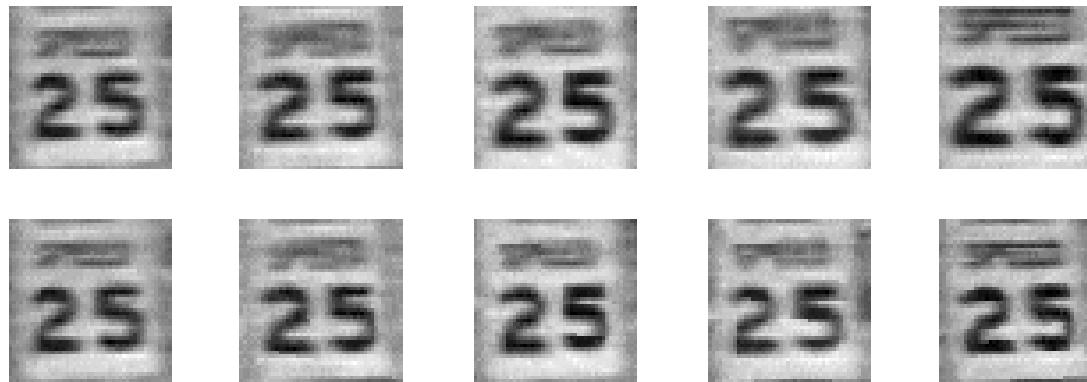
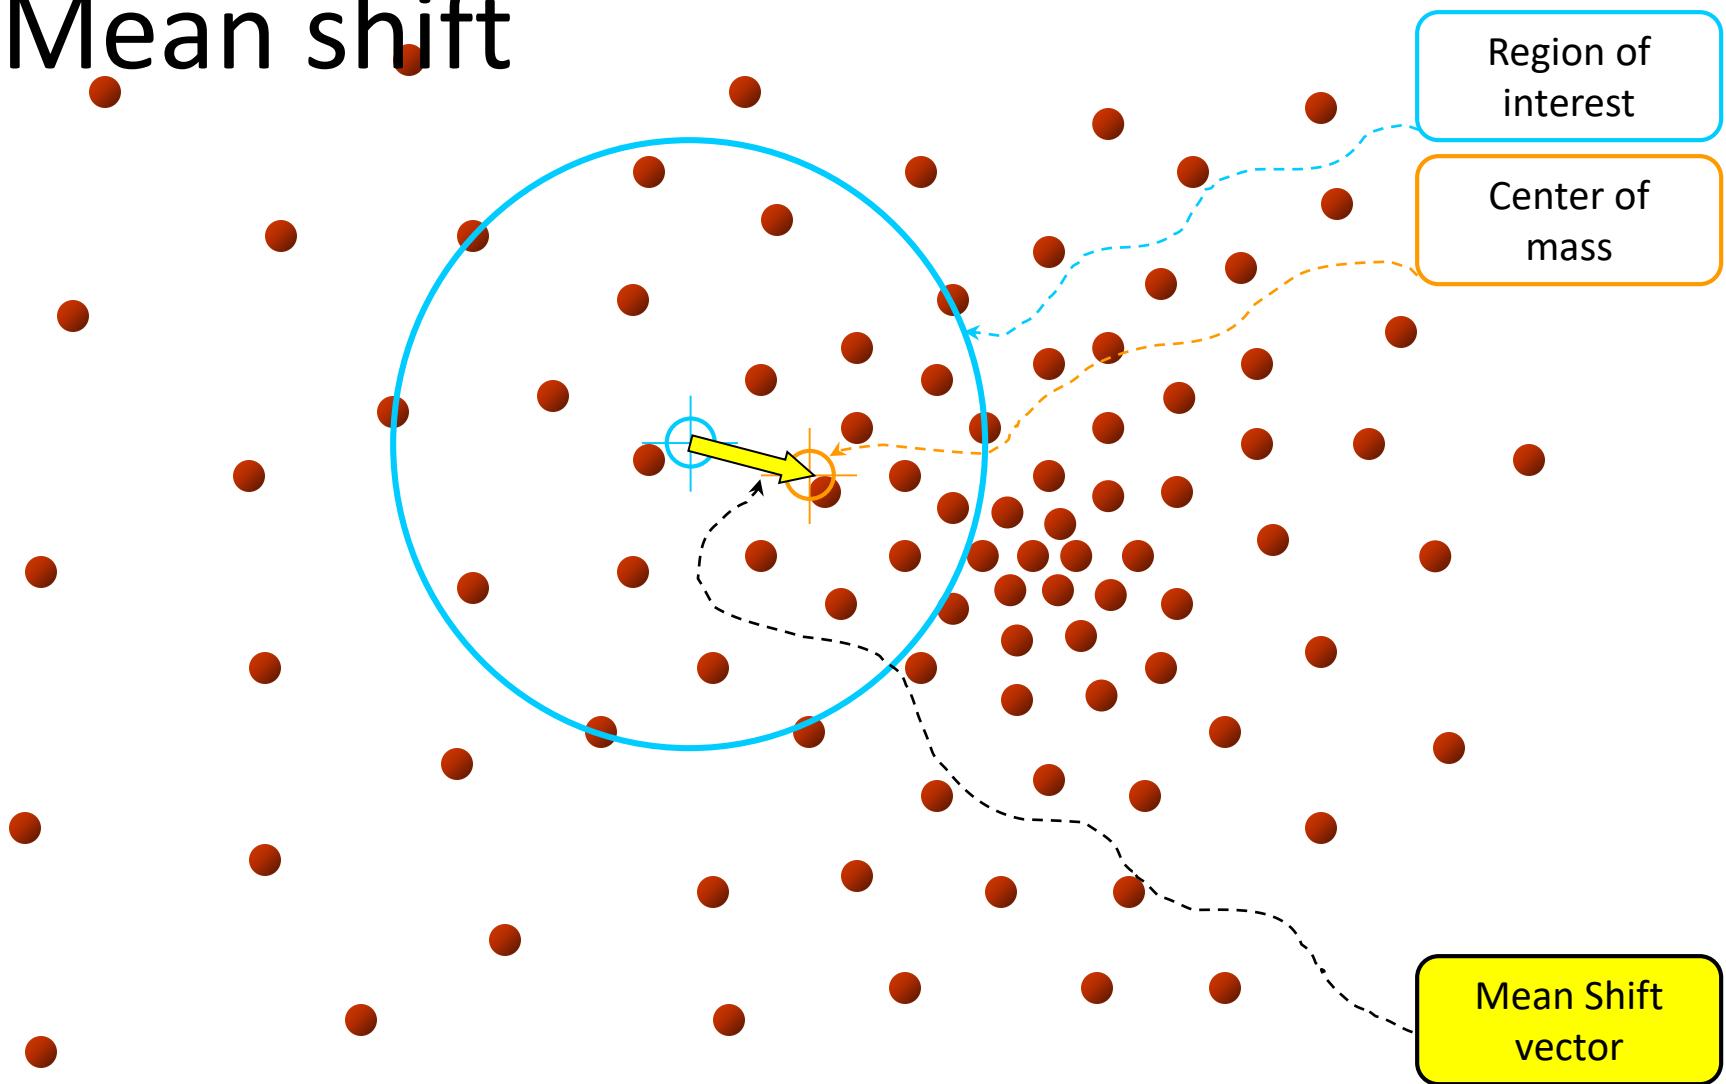


Figure 2: The traffic sign windows from frames 1,6,11,16,21 as tracked (top), and warped by the computed deformation matrices (bottom).

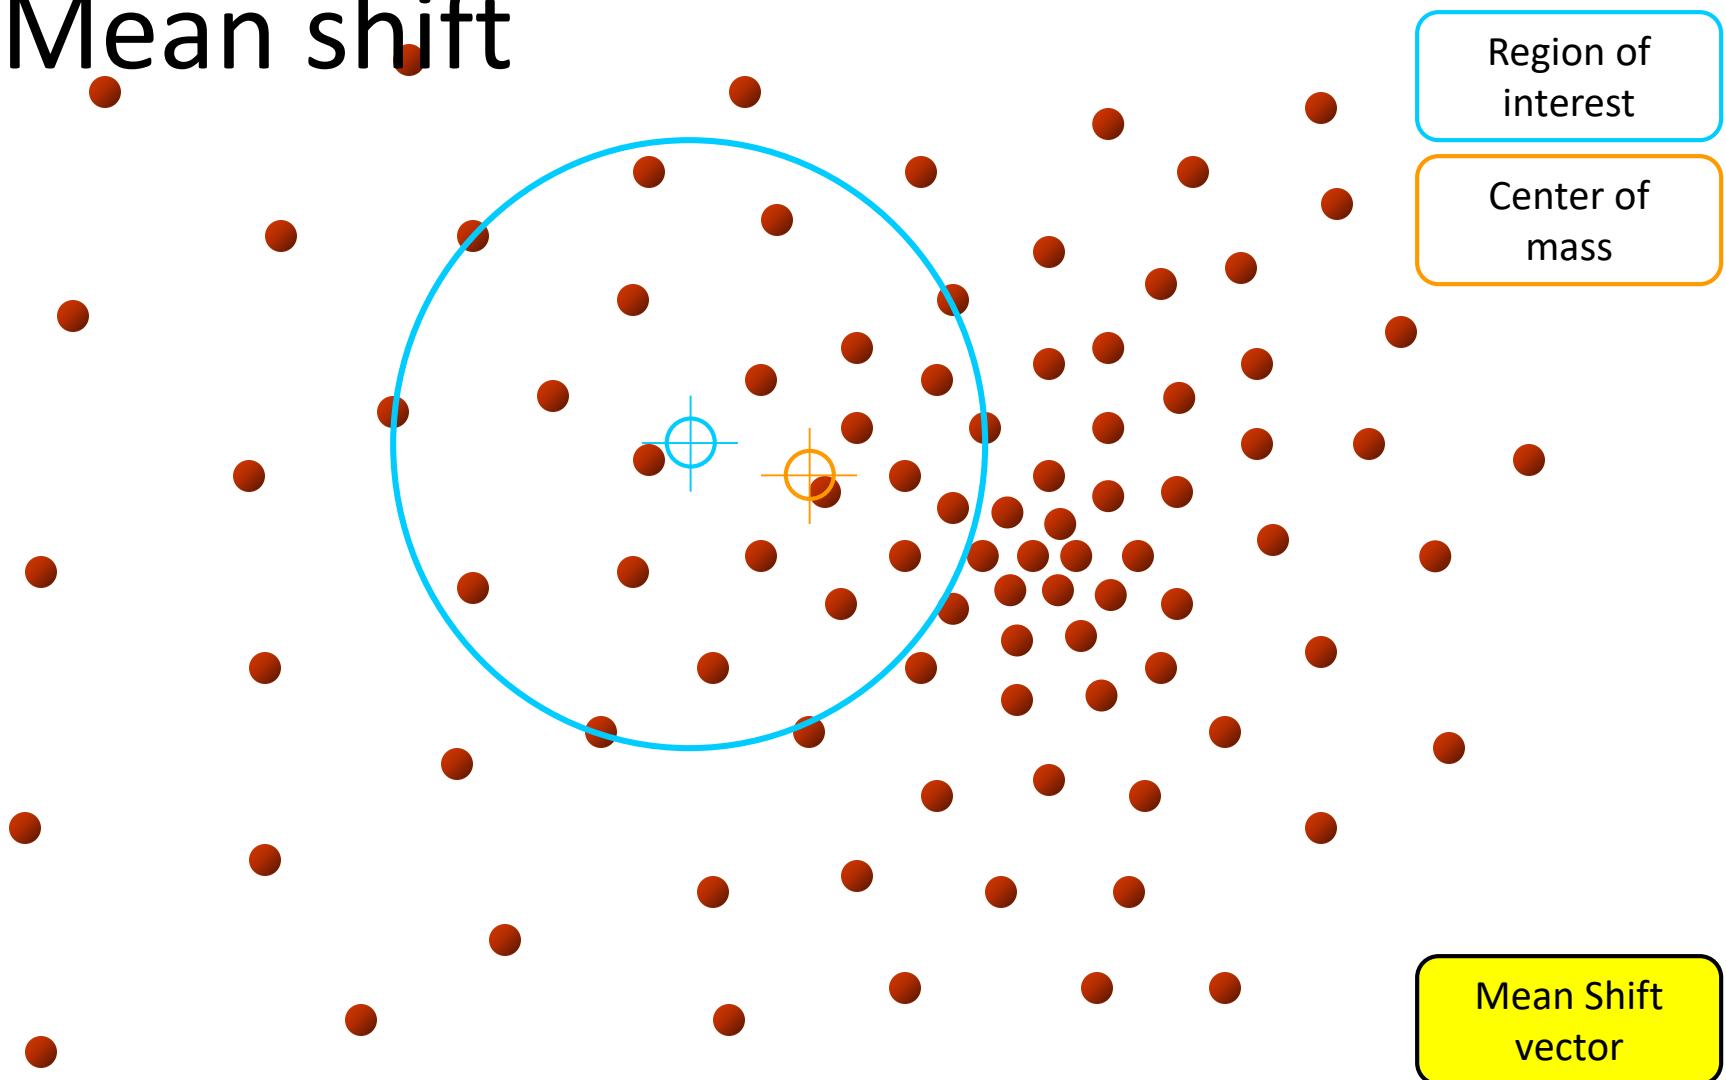
Summary of KLT Tracking (Kanade-Lucas-Tomasi Feature Tracker)

- Find a good point to track (harris corner)
- Use intensity second moment matrix and difference across frames to find displacement
- Iterate and use coarse-to-fine search to deal with larger movements
- When creating long tracks, check appearance of registered patch against appearance of initial patch to find points that have drifted

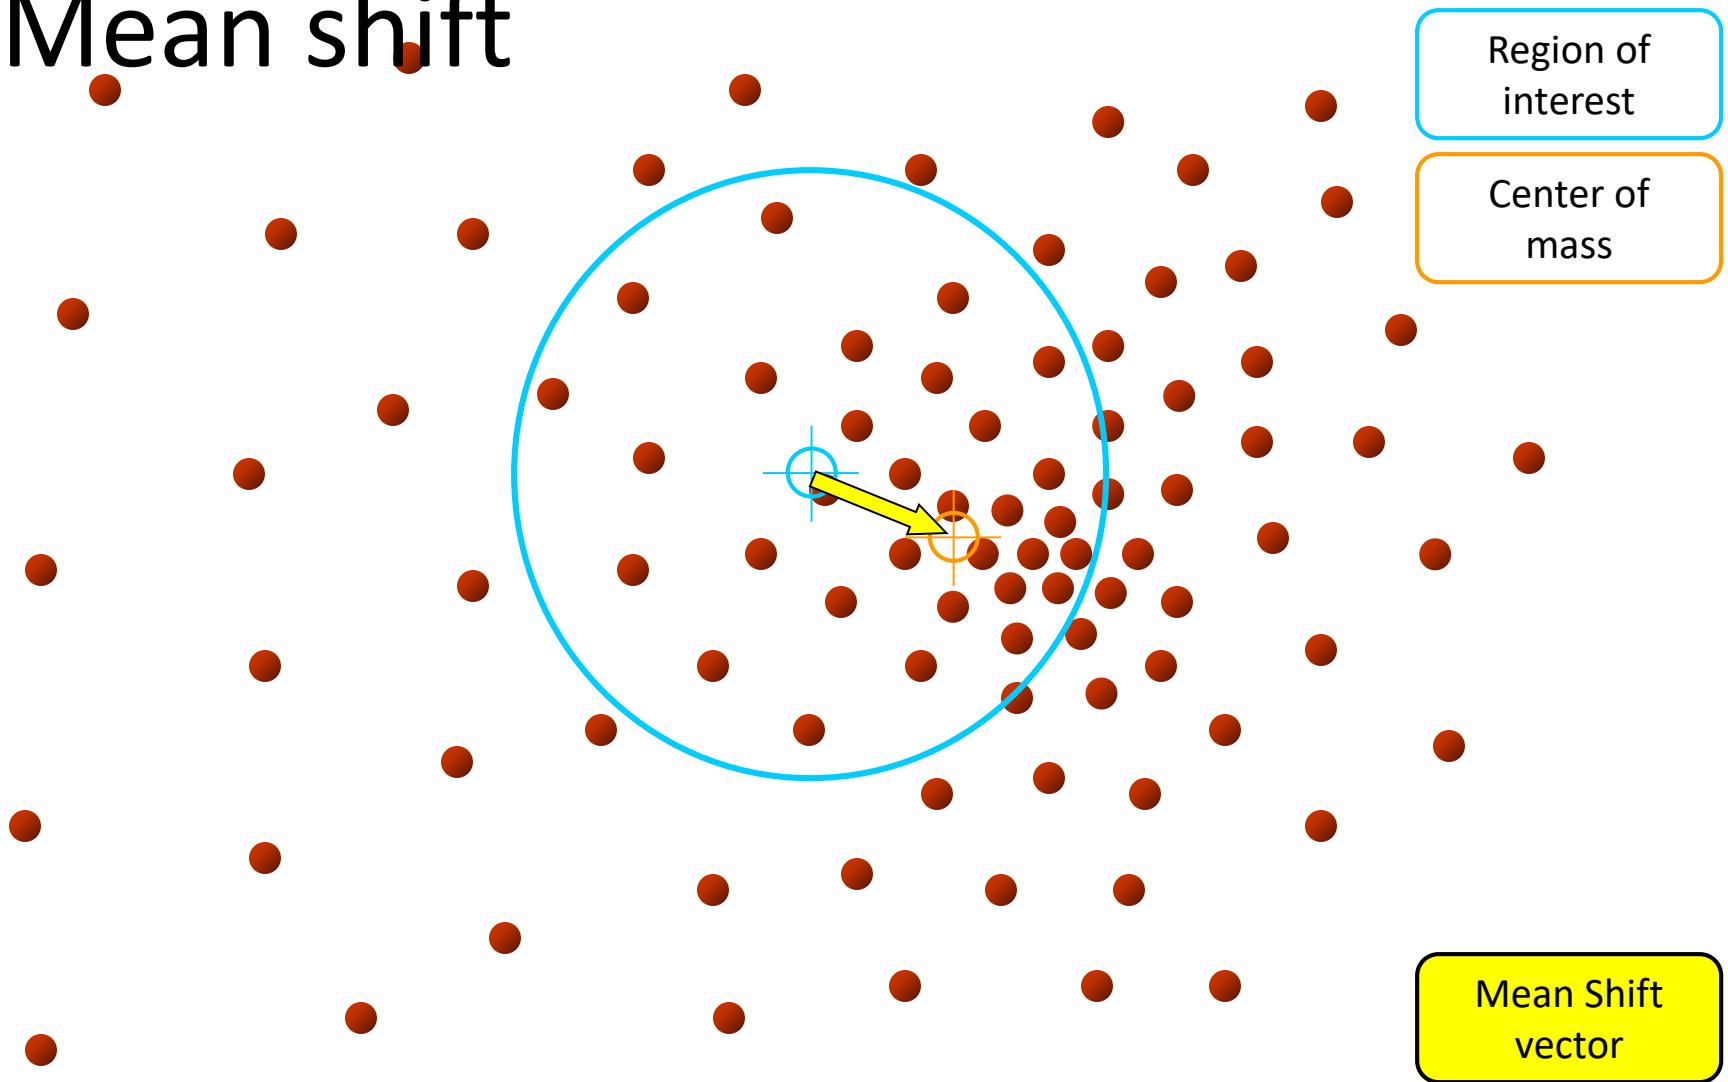
Mean shift



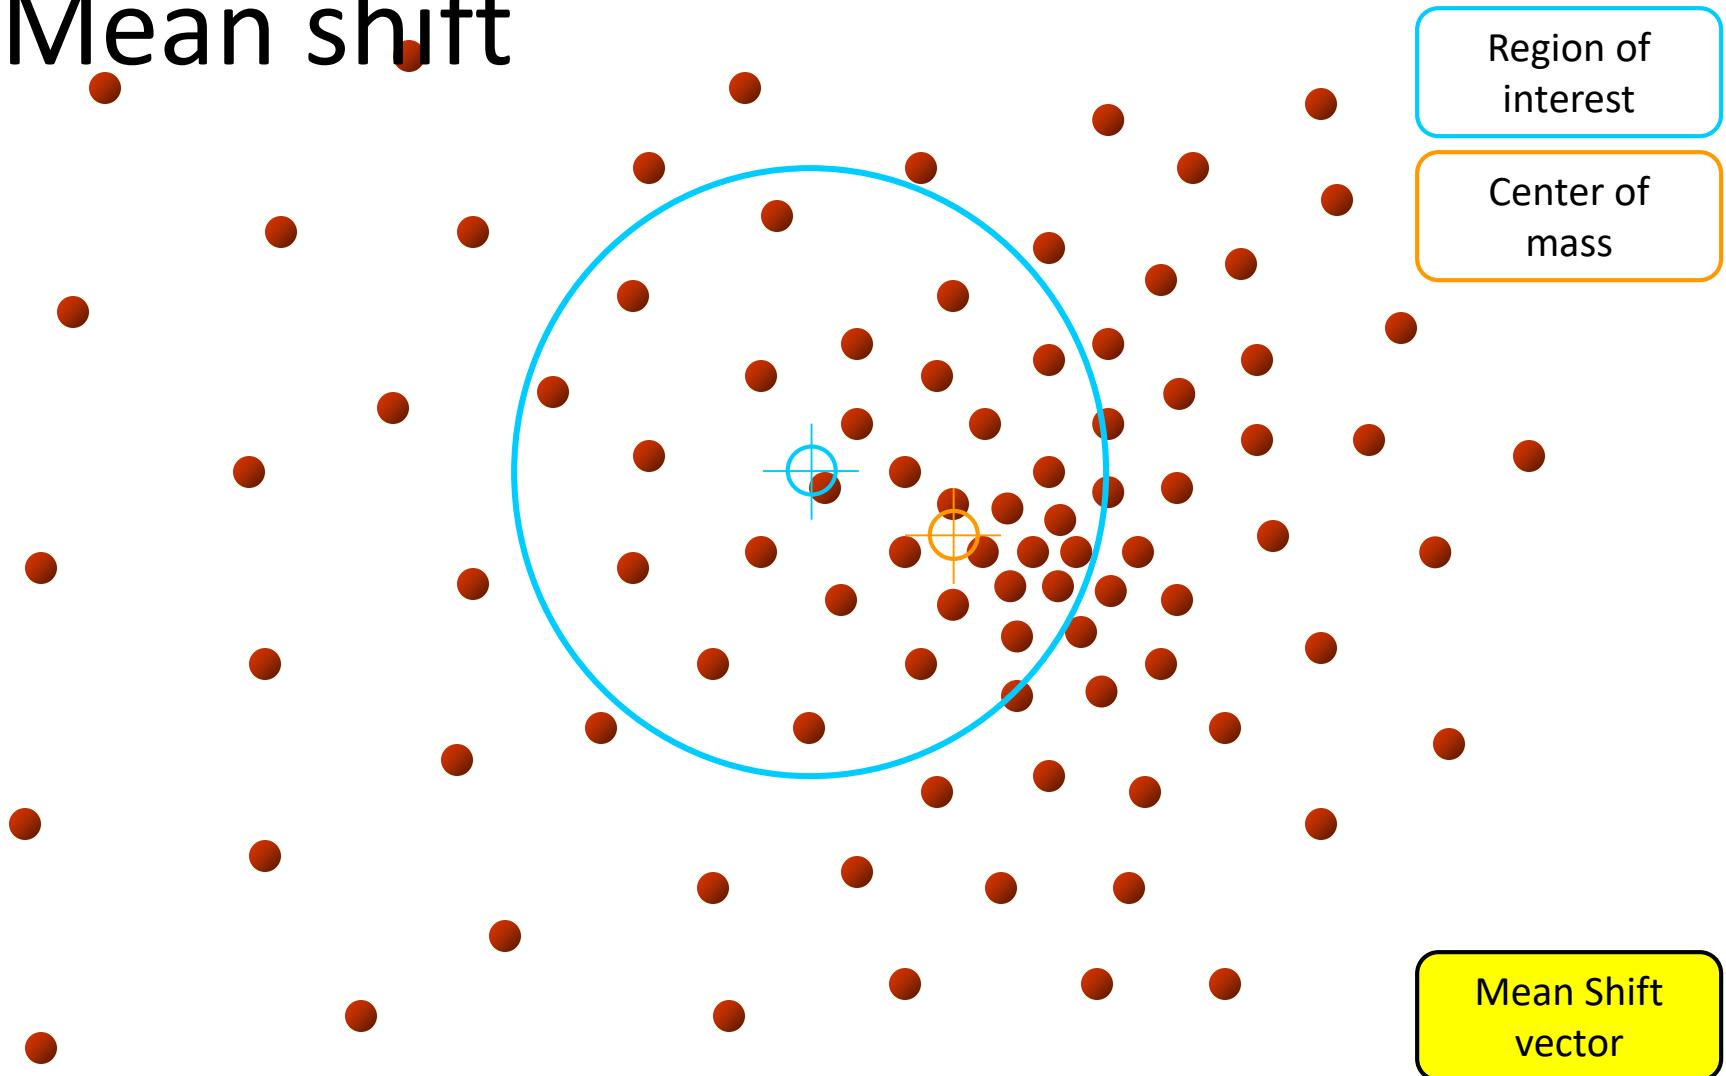
Mean shift



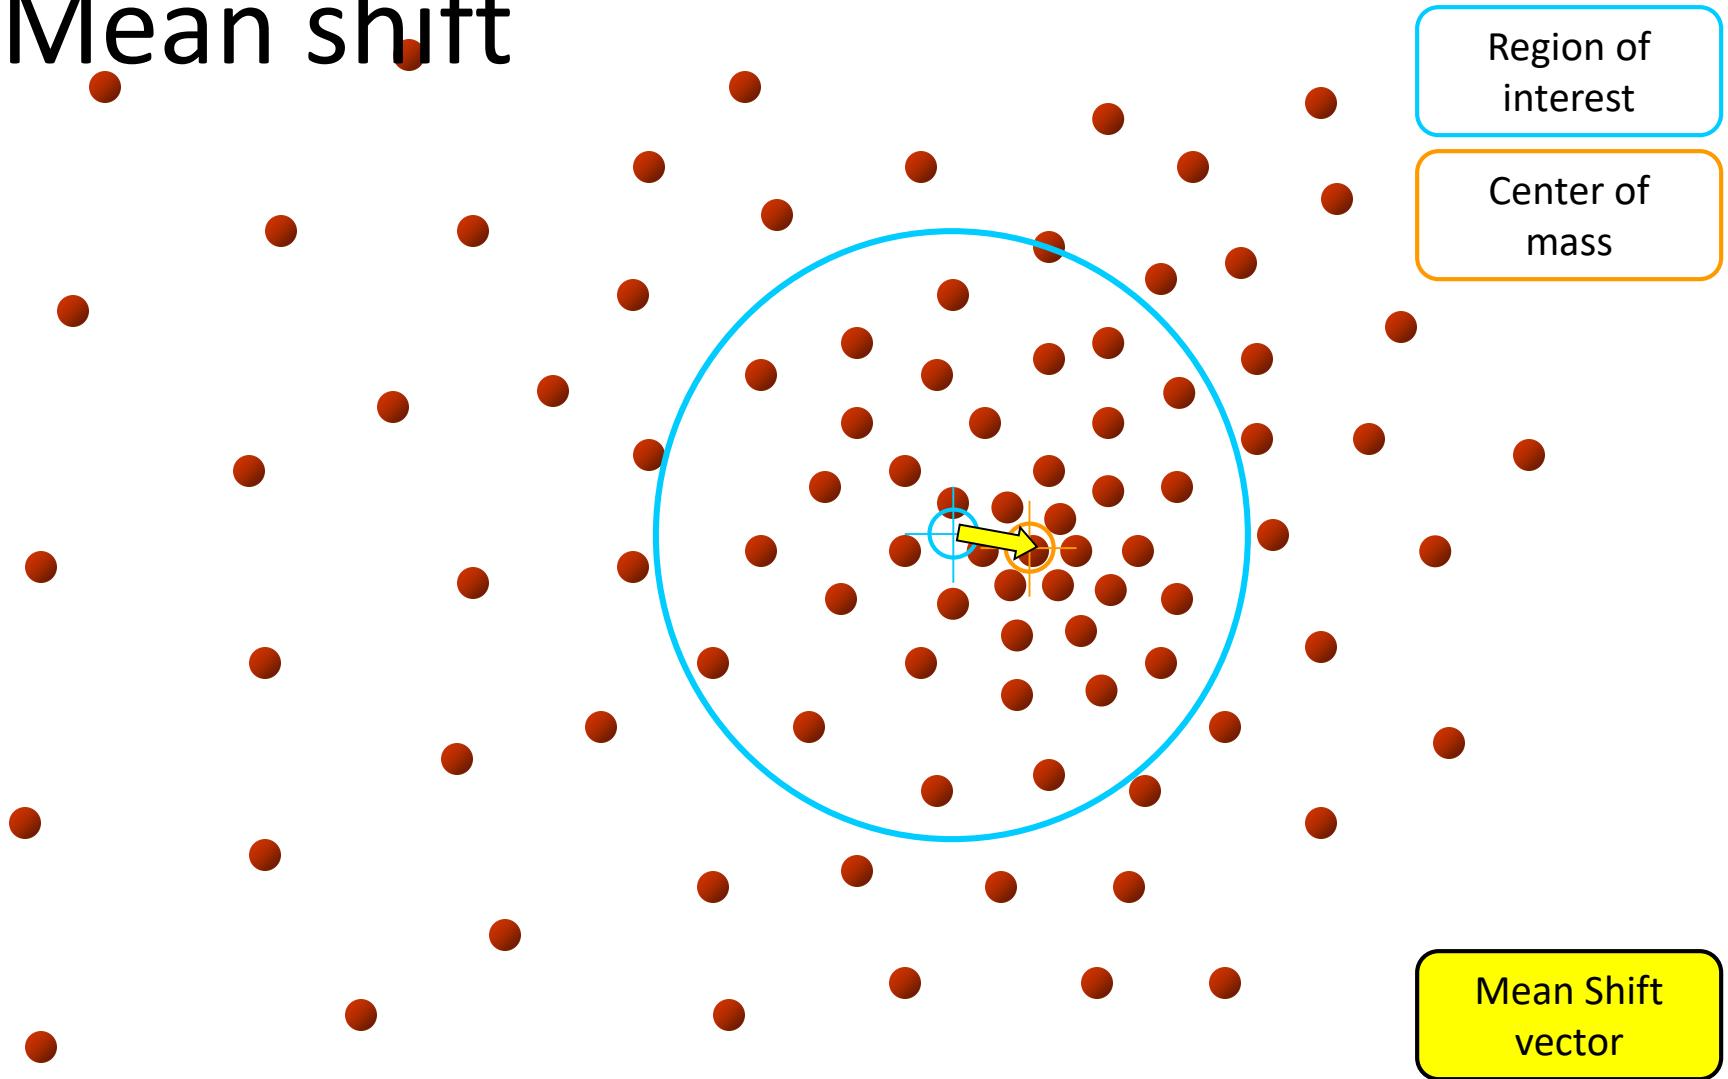
Mean shift



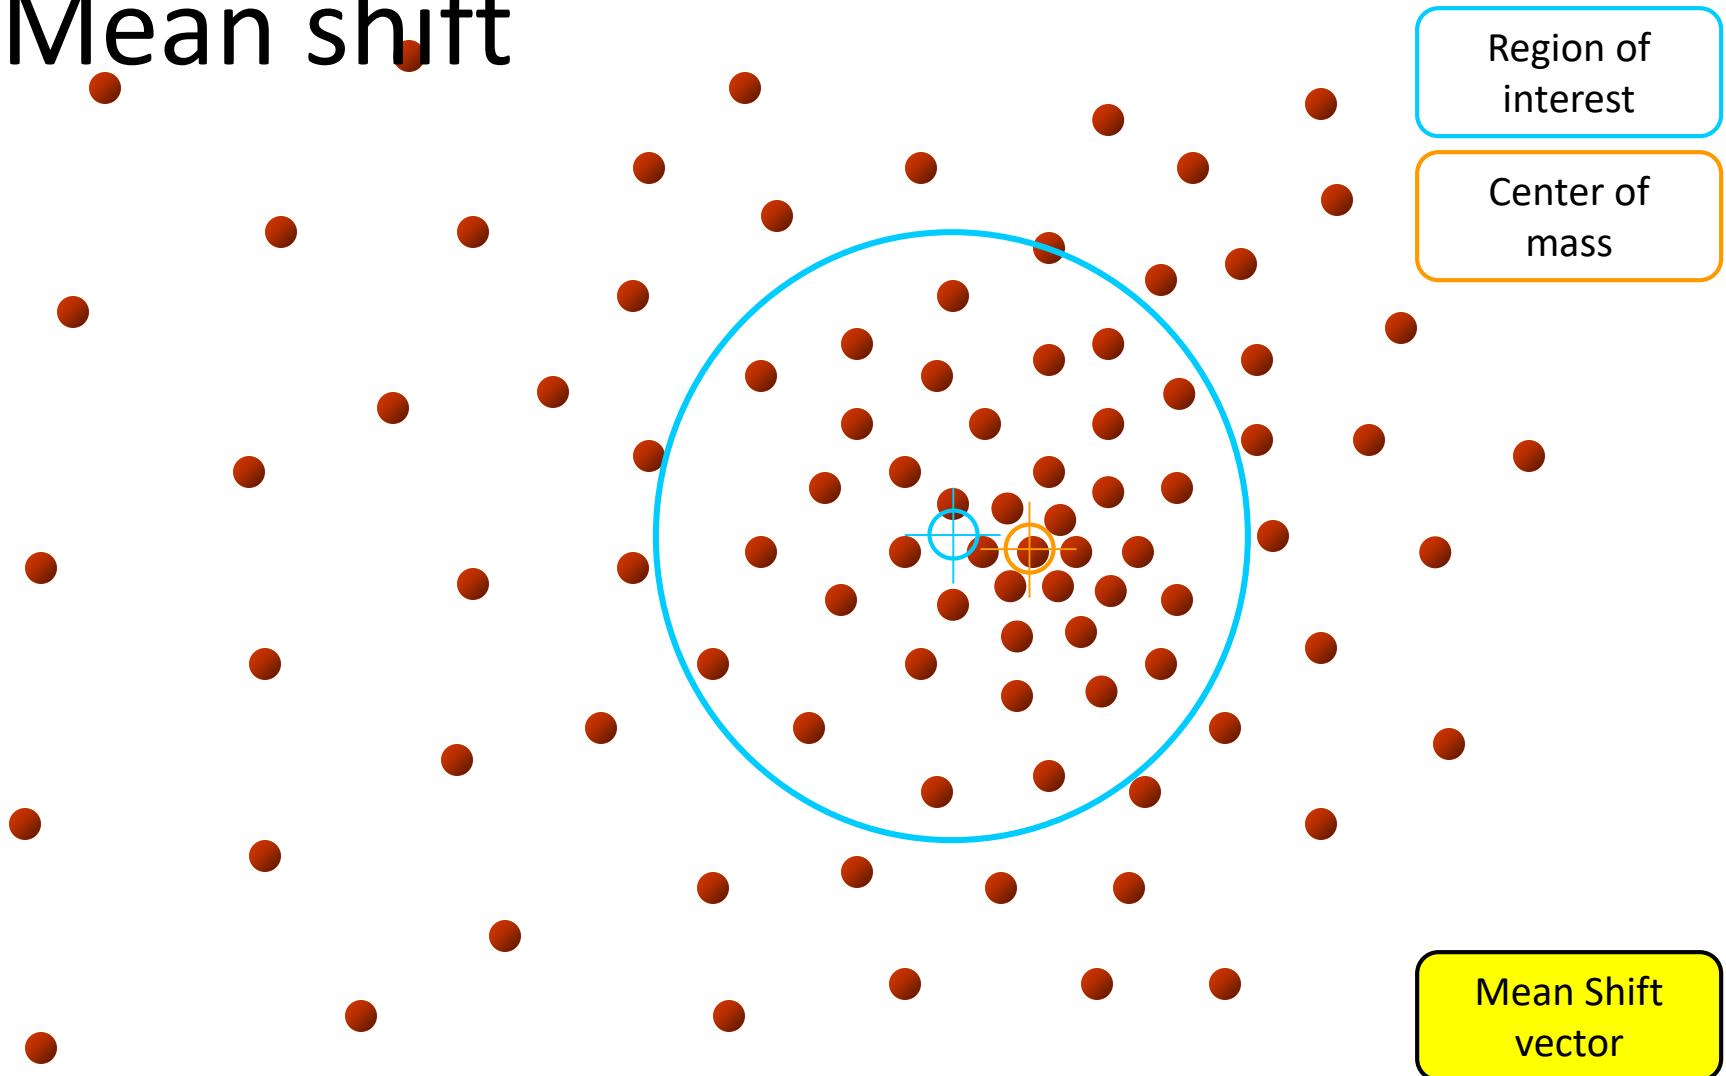
Mean shift



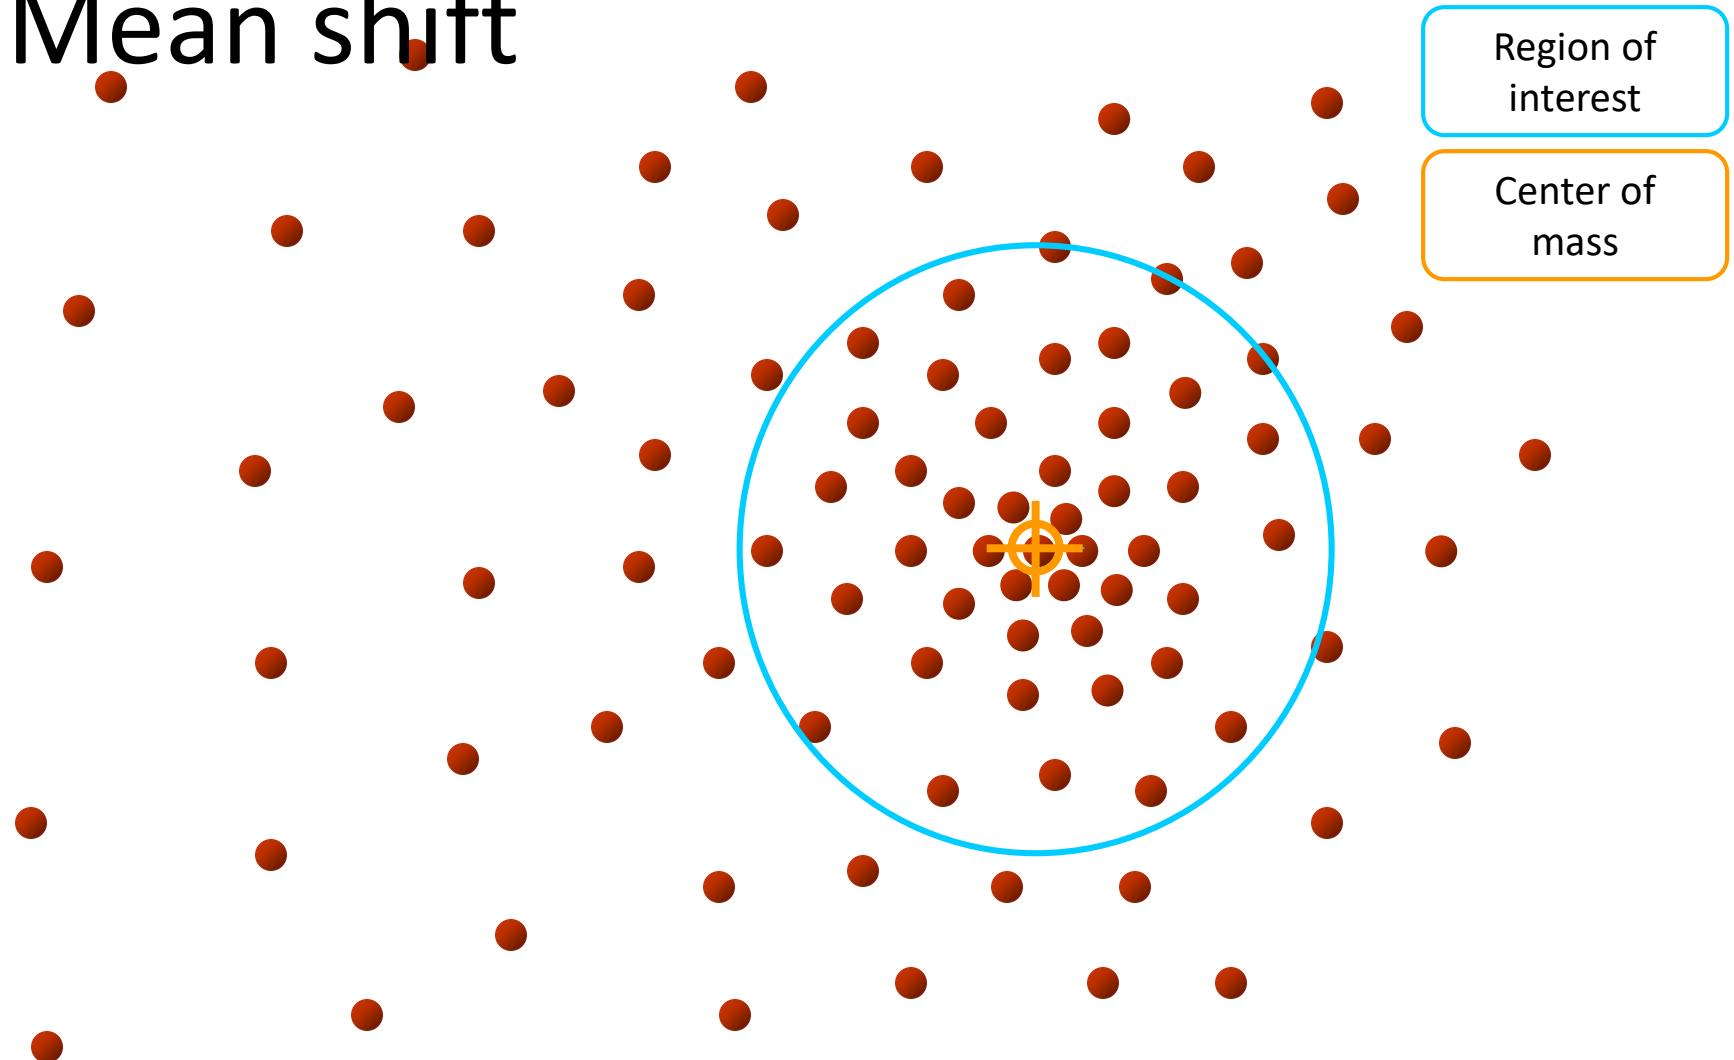
Mean shift



Mean shift



Mean shift



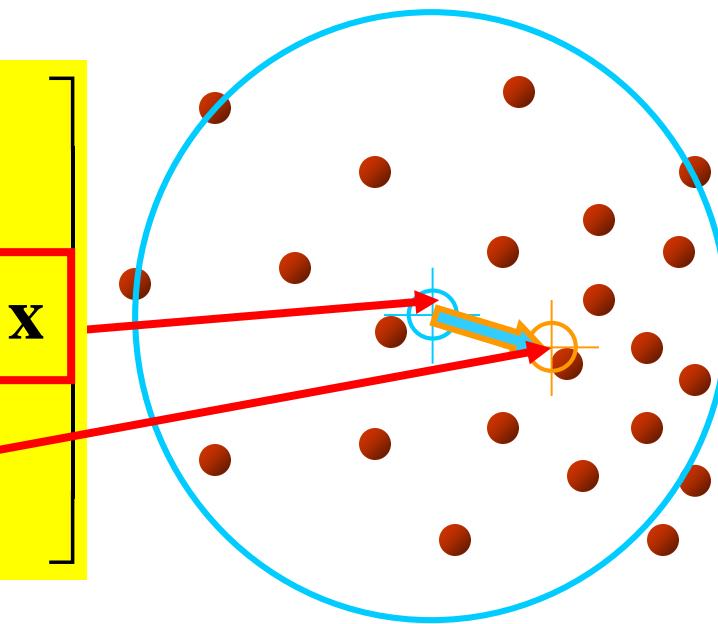
Region of
interest

Center of
mass

Computing the Mean Shift

Simple Mean Shift procedure:

- Compute mean shift vector
- Translate the Kernel window by $\mathbf{m}(\mathbf{x})$

$$\mathbf{m}(\mathbf{x}) = \left[\frac{\sum_{i=1}^n \mathbf{x}_i g\left(\frac{\|\mathbf{x} - \mathbf{x}_i\|^2}{h}\right)}{\sum_{i=1}^n g\left(\frac{\|\mathbf{x} - \mathbf{x}_i\|^2}{h}\right)} \right] \mathbf{x}$$


Kalman Filter

- Kalman filters are modeled on a Markov chain built on **linear operators** perturbed by **Gaussian noises**.
- At time k , each target has state x_k

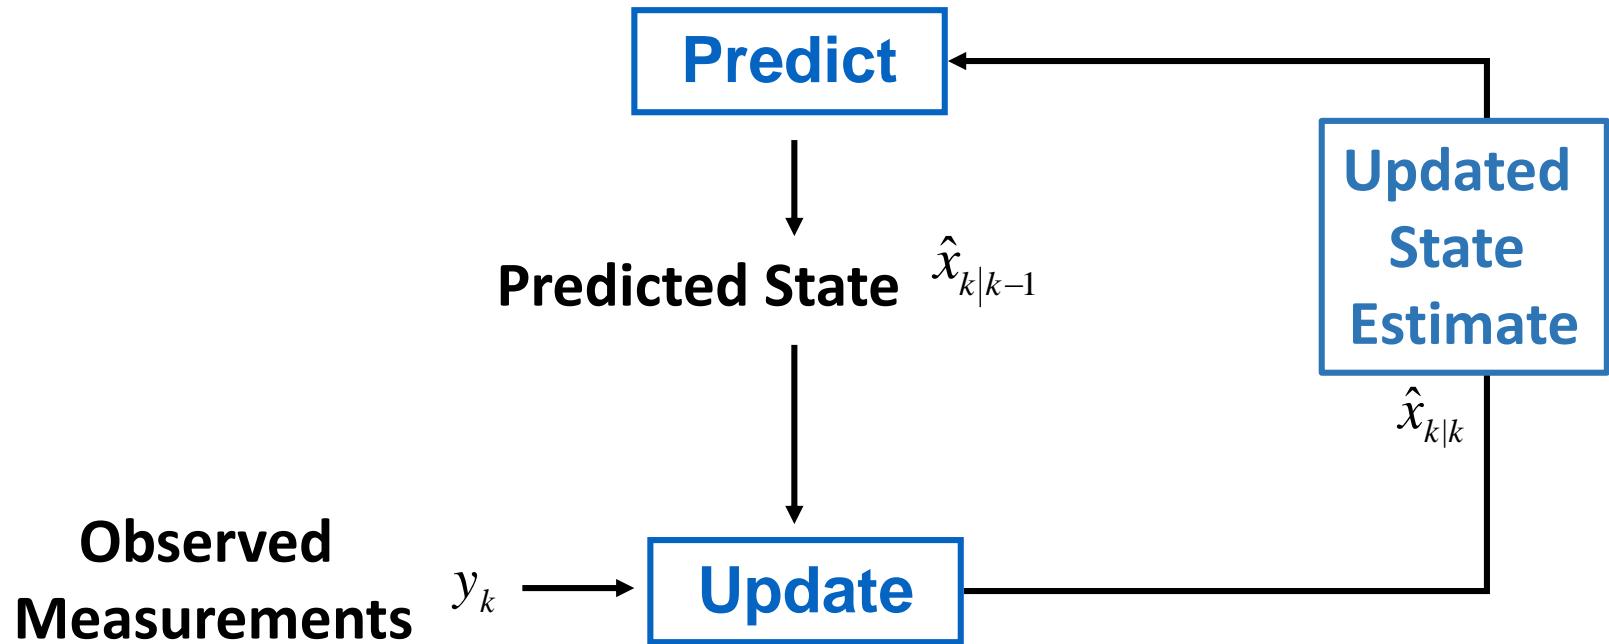
$$x_k = F_k x_{k-1} + w_k \text{ , where } w_k \sim N(0, Q_k)$$

and observation (measurement) y_k

$$y_k = H_k x_k + v_k \text{ , where } v_k \sim N(0, R_k)$$

Kalman, R. E. "A New Approach to Linear Filtering and Prediction Problems,"
Transactions of the ASME - Journal of Basic Engineering Vol. 82: pp. 35-45, 1960.

Kalman Filter Phases



$$x_k = y_k = [u_k \ v_k \ \dot{u}_k \ \dot{v}_k]^T$$

$$F_k = \begin{bmatrix} 1 & 0 & 1 & 0 \\ 0 & 1 & 0 & 1 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

$$H_k = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

Kalman Filter Phases

Predict Phase

- Predicted State

$$\hat{x}_{k|k-1} = F_k \hat{x}_{k-1|k-1}$$

- Predicted Estimate Covariance

$$P_{k|k-1} = F_k P_{k-1|k-1} F_k^T + Q_k$$

Update Phase

- Updated State Estimate

$$\hat{x}_{k|k} = \hat{x}_{k|k-1} + K_k \tilde{y}_k$$

- Updated Estimate Covariance

$$P_{k|k} = (I - K_k H_k) P_{k|k-1}$$

- Kalman Gain

$$K_k = P_{k|k-1} H_k^T S_k^{-1}$$

- Innovation Covariance

$$S_k = H_k P_{k|k-1} H_k^T + R_k$$

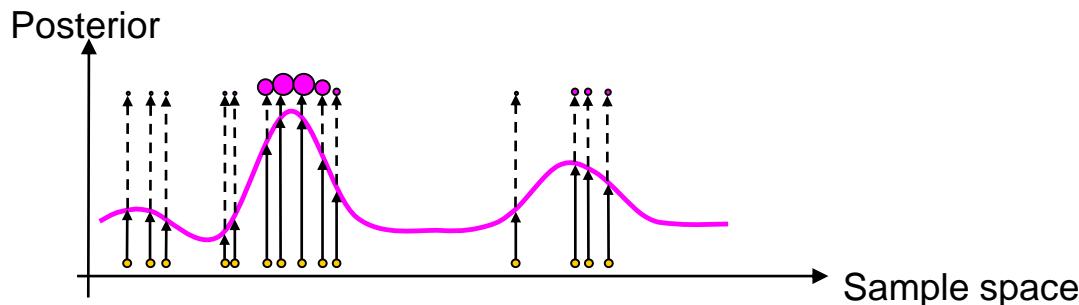
- Innovation (Measurement) Residual

$$\tilde{y}_k = y_k - H_k \hat{x}_{k|k-1}$$

Particle filtering

- Many variations, one general concept:

Represent the posterior pdf by a set of randomly chosen weighted samples (particles)



- Randomly Chosen = Monte Carlo (MC)
- As the number of samples become very large – the characterization becomes an equivalent representation of the true pdf

Generic PF

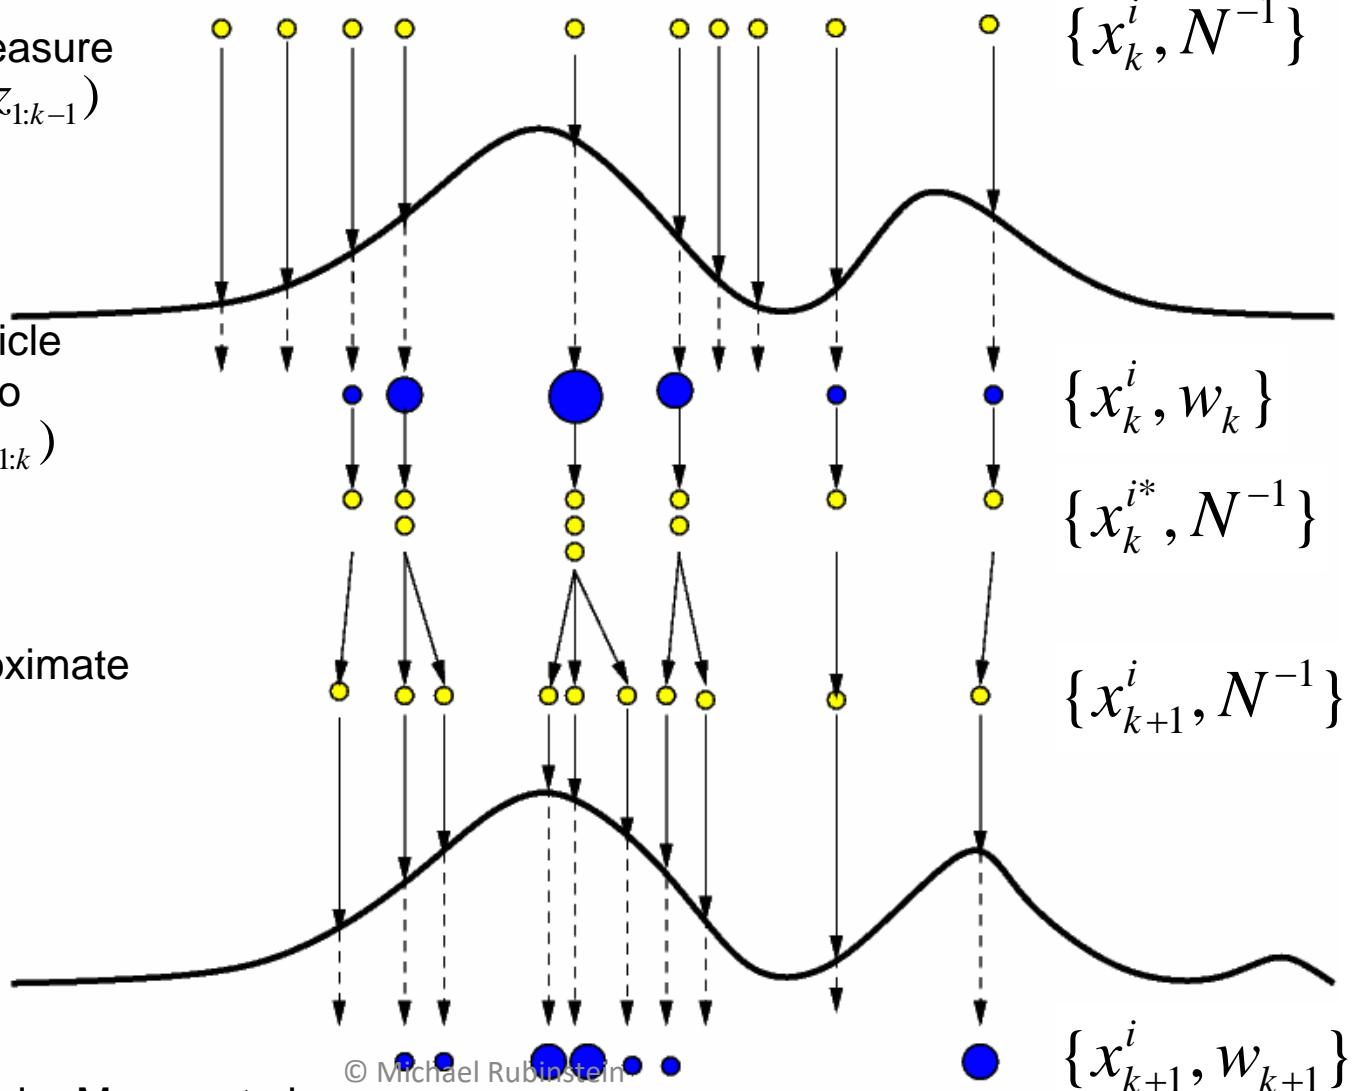
$i=1, \dots, N=10$ particles

Uniformly weighted measure
Approximates $p(x_k | z_{1:k-1})$

Compute for each particle
its importance weight to
Approximate $p(x_k | z_{1:k})$
(Resample if needed)

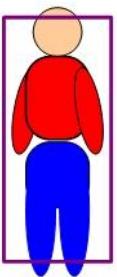
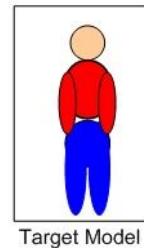
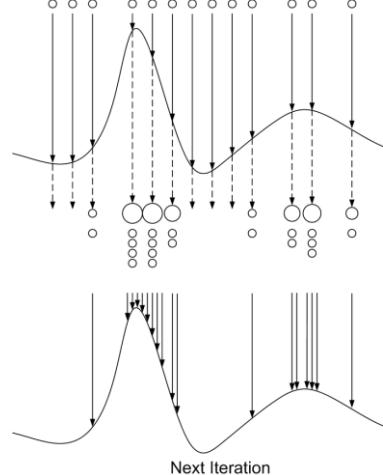
Project ahead to approximate
 $p(x_{k+1} | z_{1:k})$

$p(x_{k+1} | z_{1:k+1})$



Particle Filter

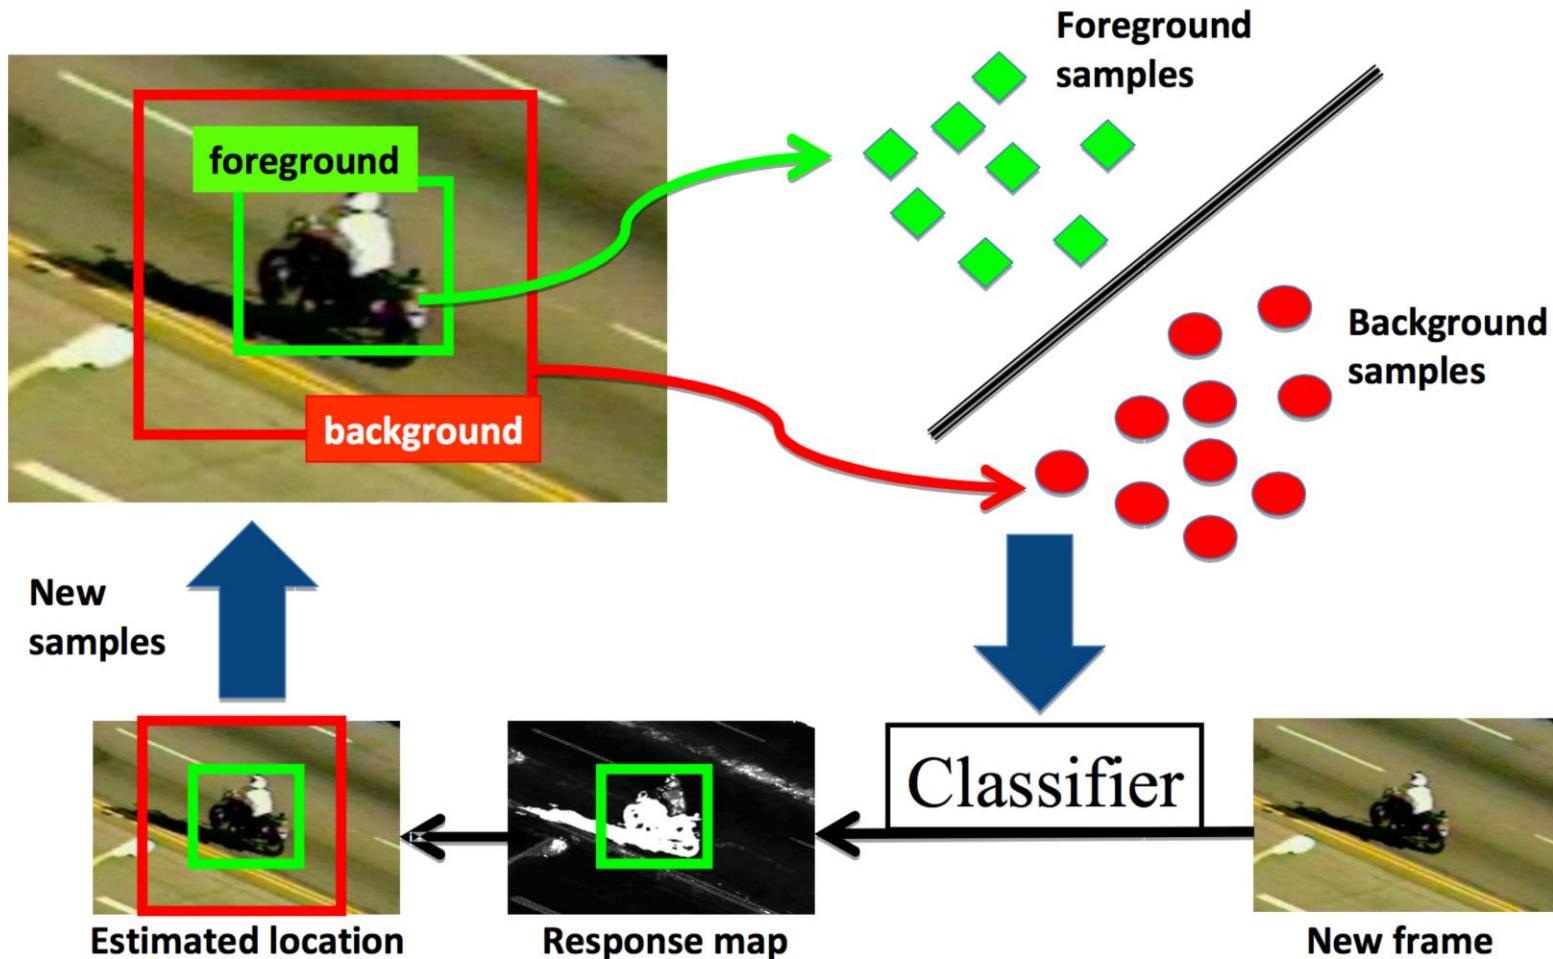
- Try to approximate $p(\underline{X} | \underline{Z}_t)$
- Sampling → evaluation → resampling



On-line Learning

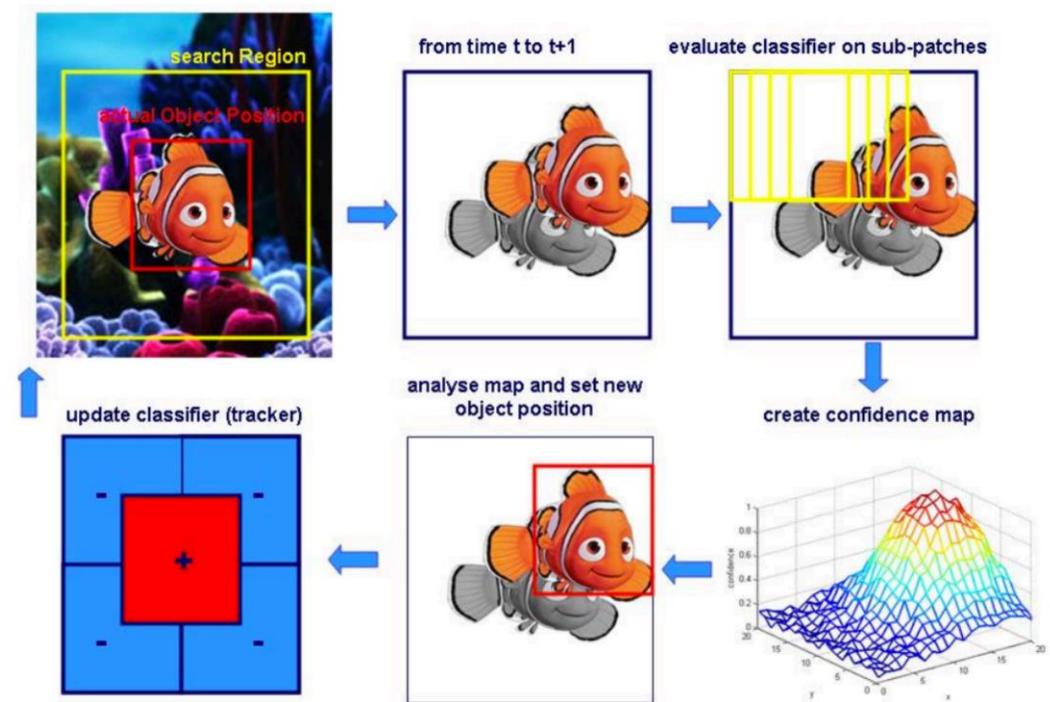
- Discriminative modeling (tracking-by-detection)
- Learn and apply a detector or predictor
- Challenges:
 - What are the training data? Label?
 - How to avoid drift? Handle occlusion?
 - How to control complexity?

On-line Learning



On-line Learning

- On-line discriminative learning
- One-shot learning
- On-line update of the classifier



Grabner and Bischof CVPR06

On-line Learning

- Examples of on-line discriminative learning
 - Multiple Instance Learning [1]
 - Kernelized Structured SVM [2]
 - Combine short track + detector [3]

[1] Babenko, Boris, Ming-Hsuan Yang, and Serge Belongie. "Visual tracking with online multiple instance learning." CVPR 2009

[2] Hare, Sam, Amir Saffari, and Philip HS Torr. "Struck: Structured output tracking with kernels." ICCV 2011

[3] Kalal, Zdenek, KrysEan Mikolajczyk, and Jiri Matas. "Tracking-learning-detection." PAMI 2012

On-line Learning

- Example: TLD



Multi-camera Tracking

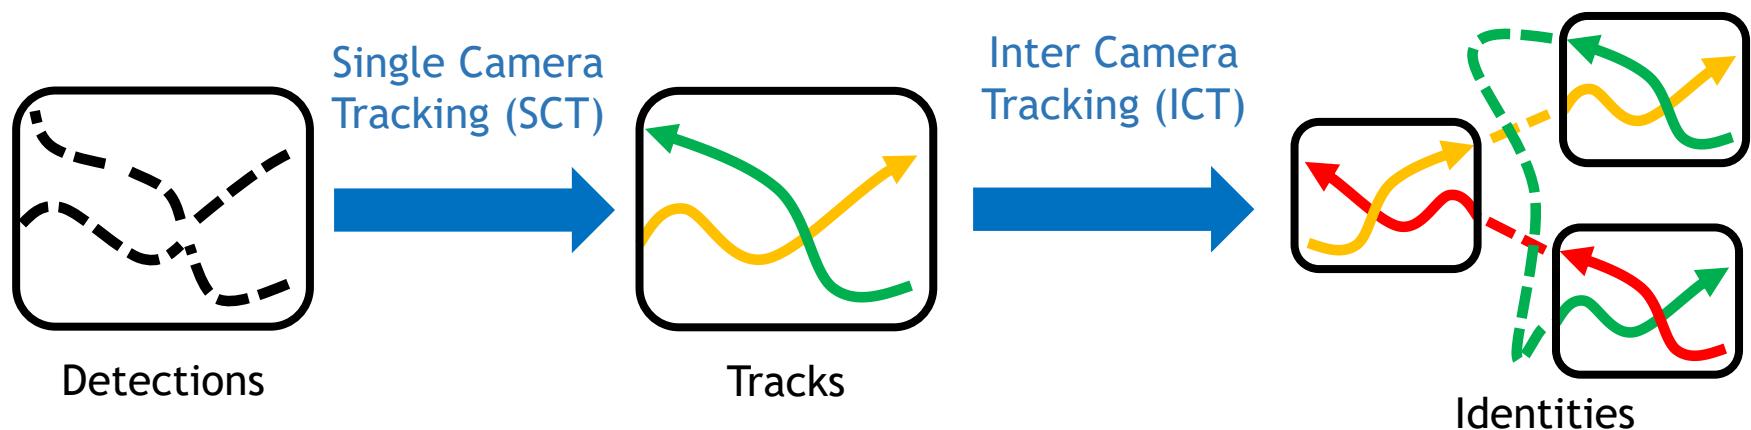
- Track objects in across multiple cameras



Ref: Y. Cai and G. Medioni, "Exploring context information for inter-camera multiple target tracking," WACV2014.

Multi-camera Tracking

- Multi-camera Tracking = single camera tracking (SCT)
+ inter camera tracking (ICT)

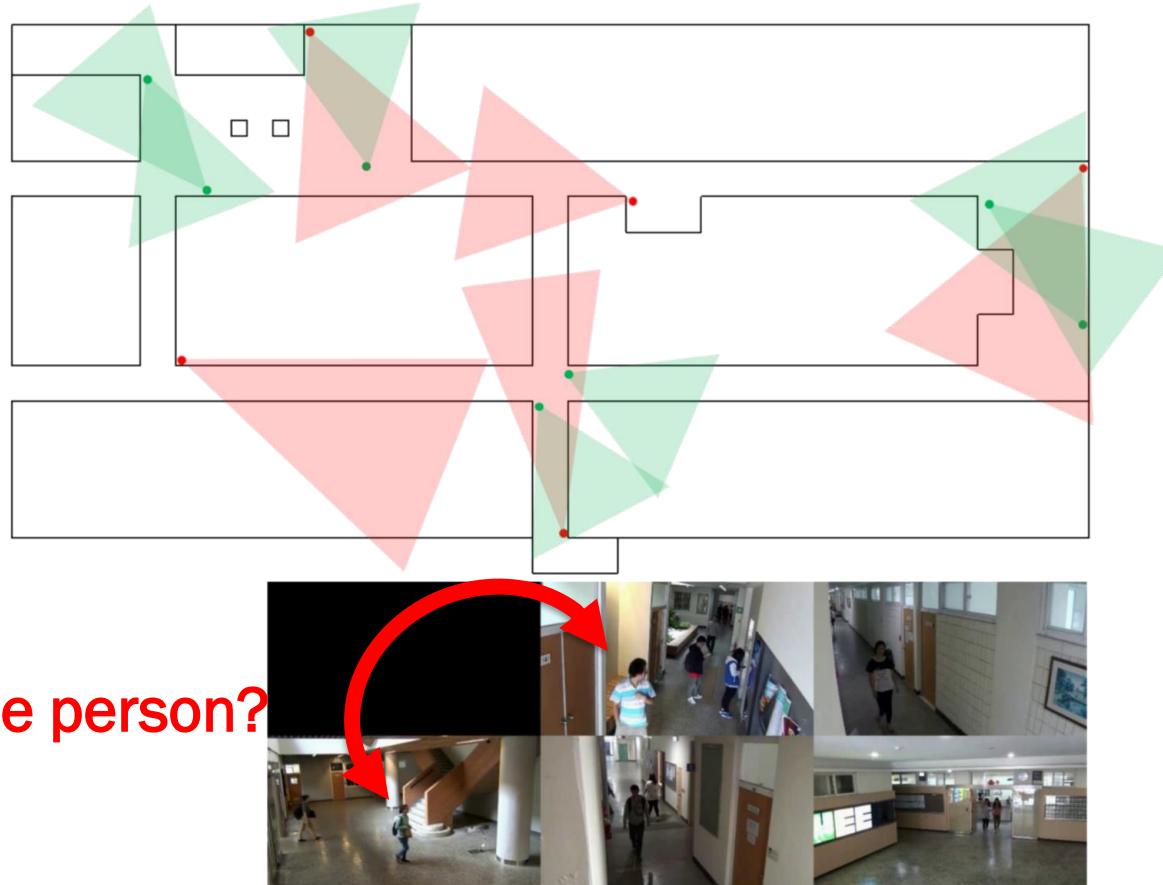


Ref: S. Zhang, E. Staudt, T. Faltemier, and A. K. Roy-Chowdhury, "A Camera Network Tracking (CamNeT) Dataset and Performance Baseline," WACV 2015.

Multi-camera Tracking

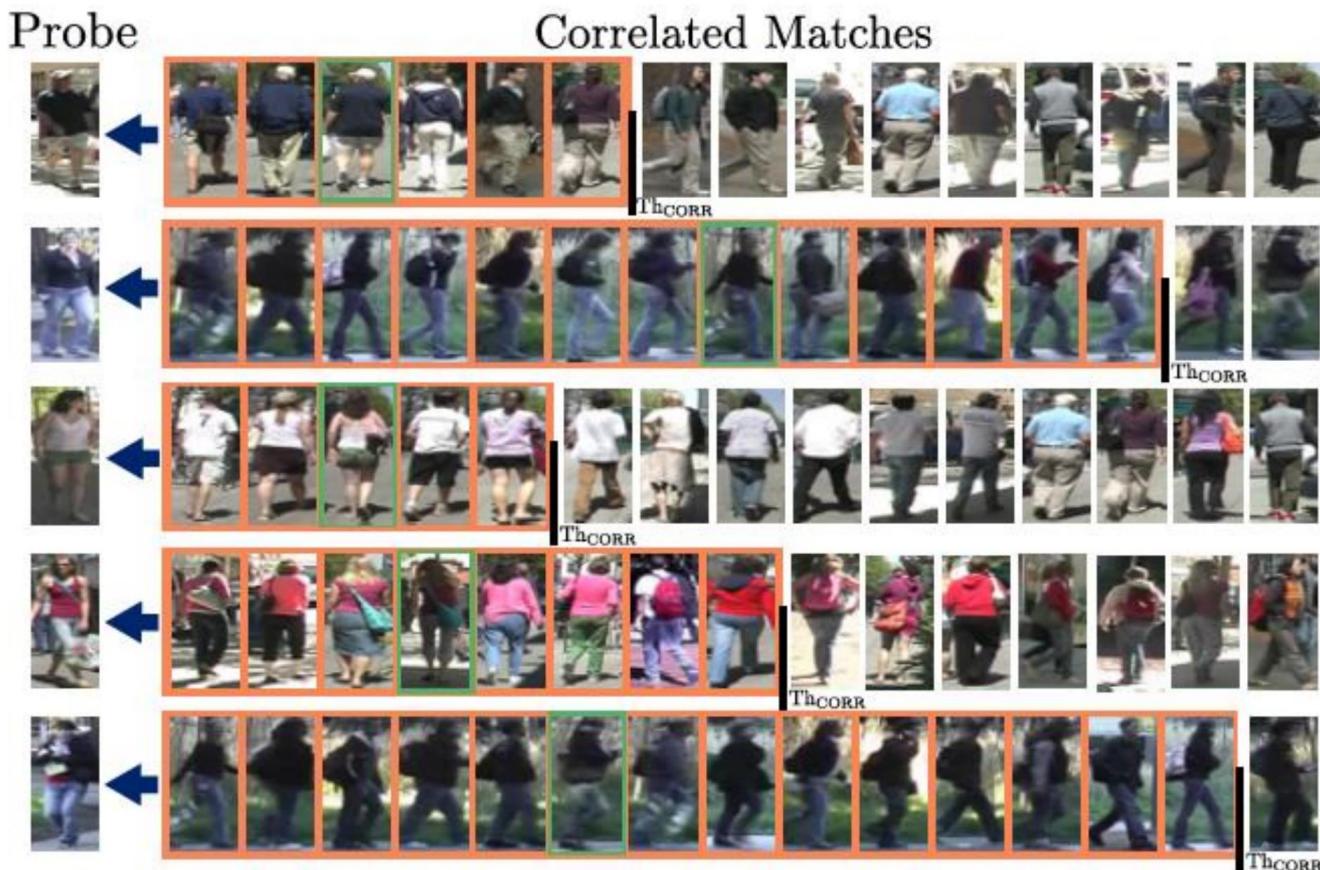
- Applications
 - Surveillance camera network
 - Customer flow analysis
 - Homeland security
 - Monitoring of elders, children, and patients
 - Crime prevention
 - Traffic control
 - ...

Re-identification



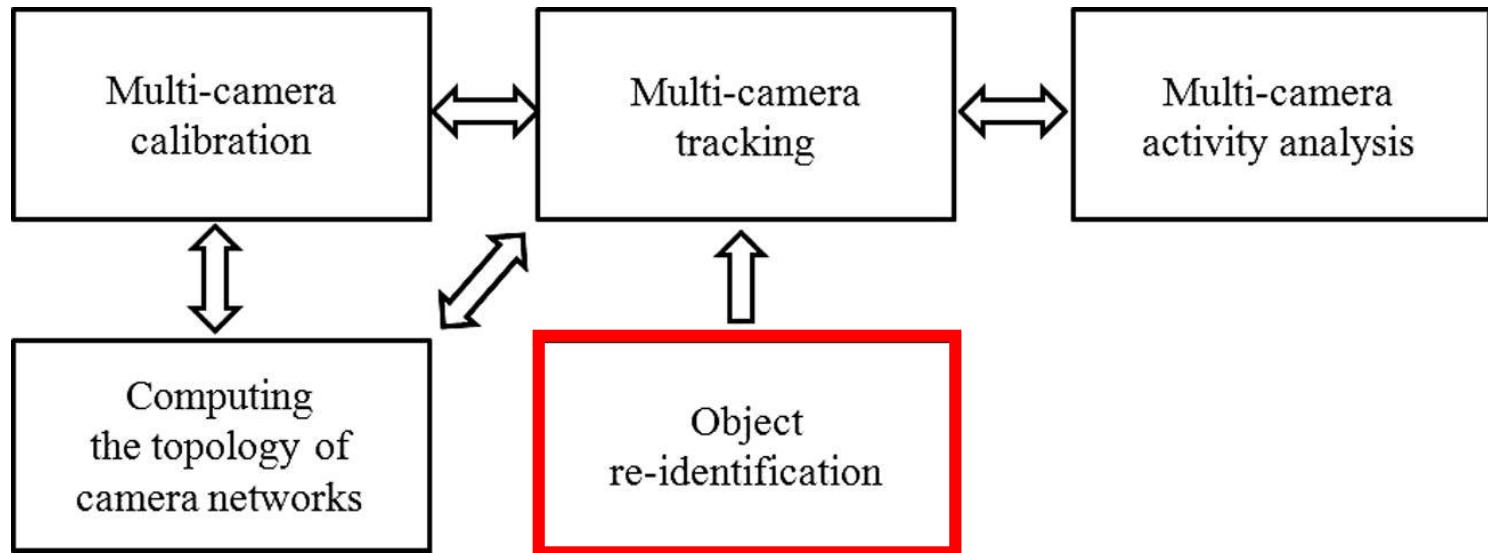
Re-identification

- is not an easy task...



The Relationship of Multi-camera Tracking and Re-identification

- Multi-camera tracking framework

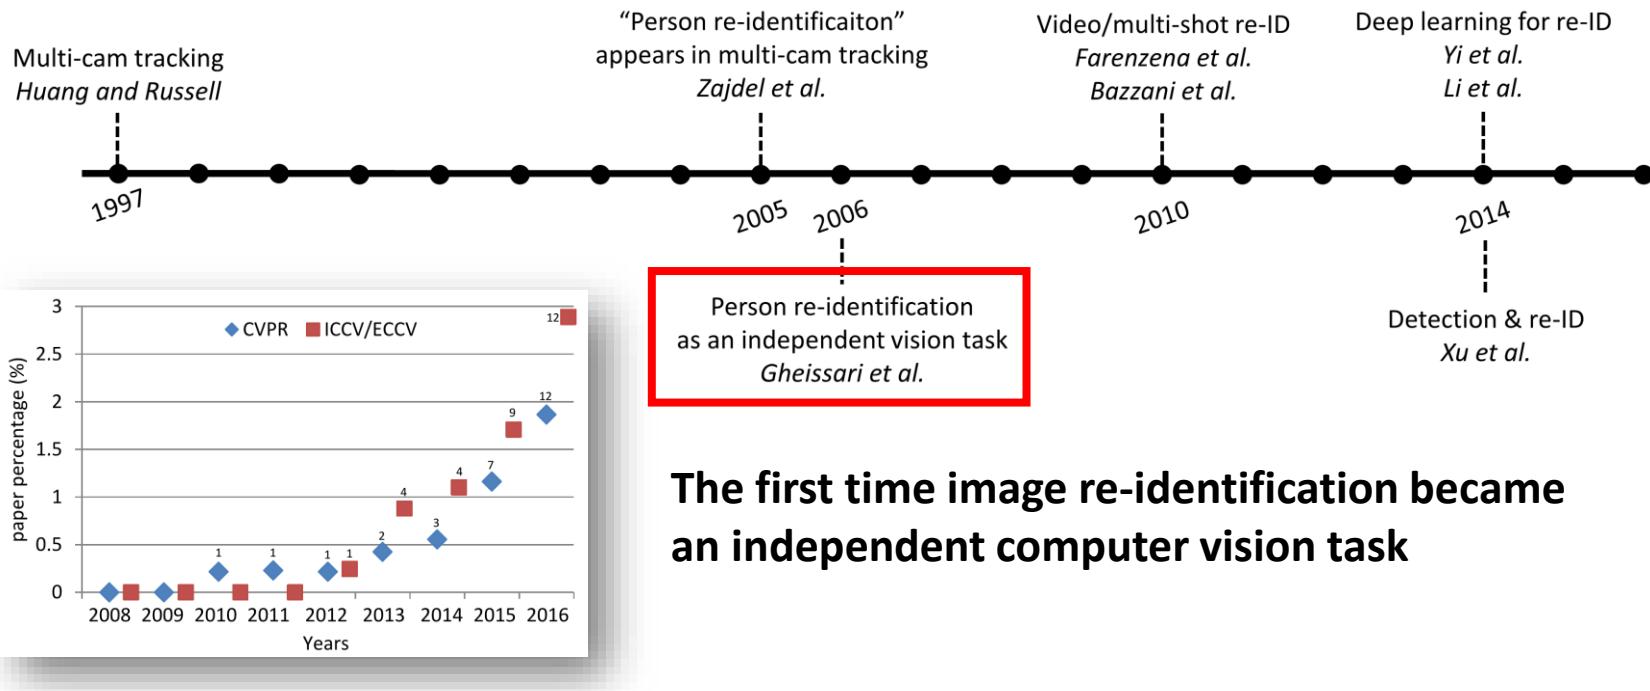


An important part of multi-camera tracking

Ref: X. Wang, "Intelligent multi-camera video surveillance: a review,"
Pattern recognition letters, vol. 34, no. 1, pp. 3–19, 2013.

The Relationship of Multi-camera Tracking and Re-identification

- History of re-identification

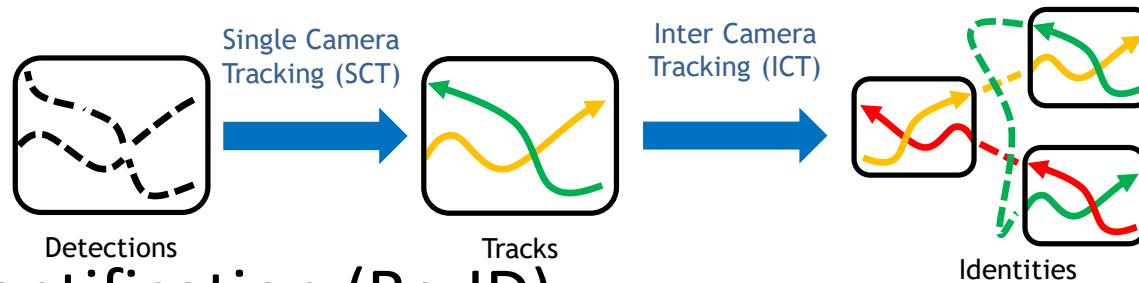


The first time image re-identification became an independent computer vision task

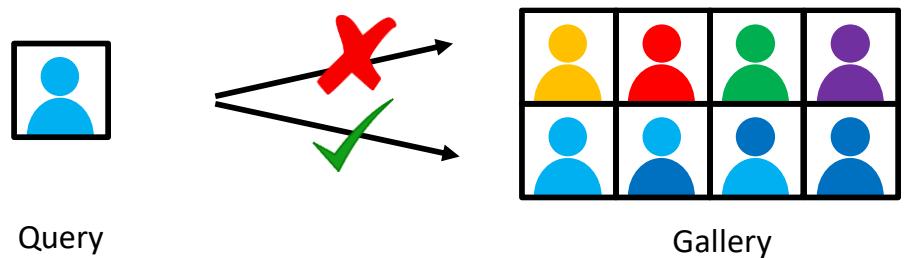
Ref: L. Zheng, Y. Yang, and A. G. Hauptmann, “Person reidentification: Past, present and future,” *arXiv preprint arXiv:1610.02984*, 2016.

The Relationship of Multi-camera Tracking and Re-identification

- Multi-camera tracking (MCT)
 - A **data association** problem



- Re-identification (Re-ID)
 - A **retrieval** problem

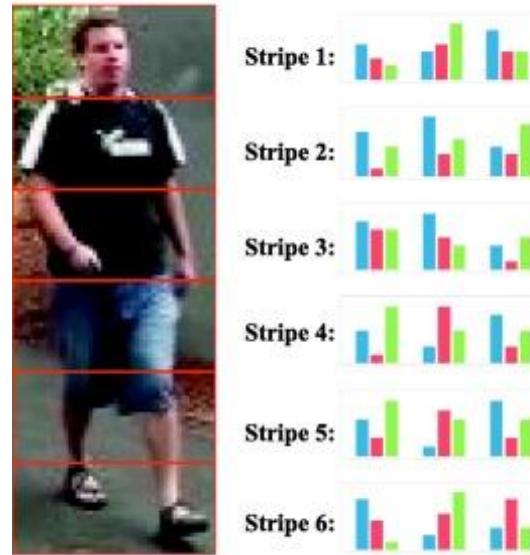


Person Re-identification

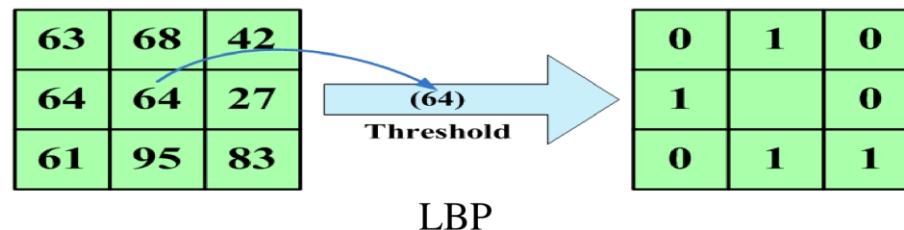
- Feature + distance function
 - Feature: LOMO, BoW, HistLBP, DGD-CNN, ...
 - Distance function: XQDA, KISSME, Mahalanobis distance, and Euclidean distance, ...
- Re-identification system with deep learning

Features

- Typical feature



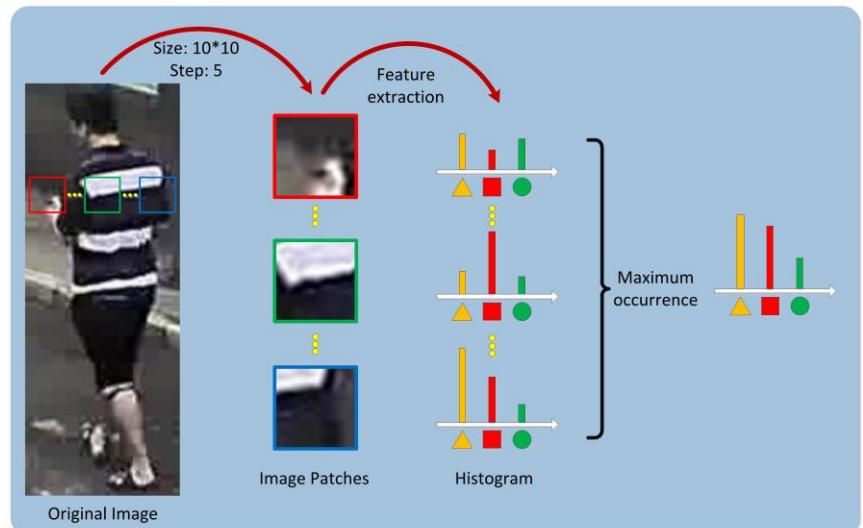
- HistLBP



Ref: B. Prosser, W.-S. Zheng, S. Gong, T. Xiang, and Q. Mary, “Person re-identification by support vector ranking,” *BMVC2010*.

Features

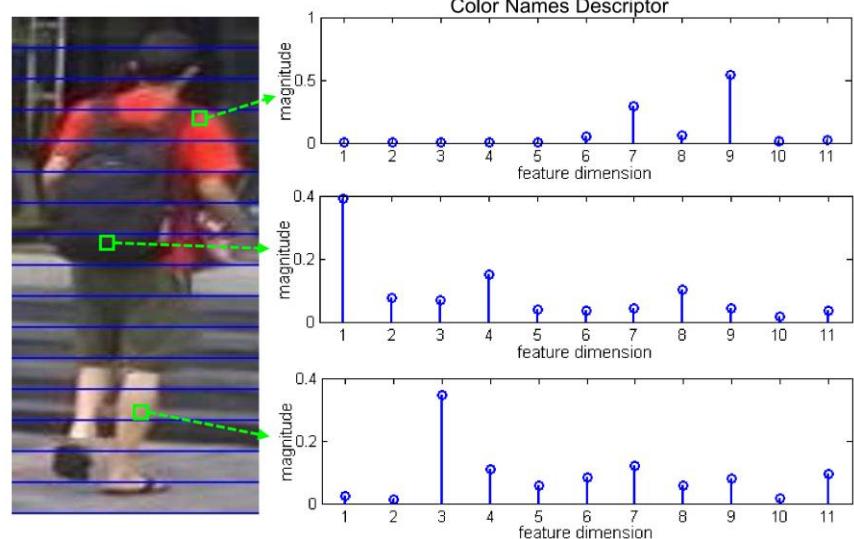
- LOMO
 - Retinex algorithm for illumination variation
 - HSV color histogram as a color descriptor, Scale Invariant Local Ternary Pattern (SILTP) as a texture descriptor
 - Sliding window + maximum occurrence for view invariance



Ref: S. Liao, Y. Hu, X. Zhu, and S. Z. Li, "Person re-identification by local maximal occurrence representation and metric learning," *CVPR 2015*.

Features

- BoW
 - Feature Extraction
 - Color name (CN) descriptor for each 4x4 non-overlapped patch
 - Local features are quantized, and pooled in a histogram for each horizontal stripe.
 - Code book is generated from a training set
 - Quantization: a feature is represented by 10 visual words



Ref: L. Zheng, L. Shen, L. Tian, S. Wang, J. Wang, and Q. Tian, “Scalable person re-identification: a benchmark,” *ICCV2015*.

Distance Function

- Mahalanobis distance

$$d(\mathbf{x}_i, \mathbf{x}_j) = (\mathbf{x}_i - \mathbf{x}_j)^T \Sigma^{-1} (\mathbf{x}_i - \mathbf{x}_j)$$

- KISSME



Ω_I Intra-person variation



Ω_E Extra-person variation

$$P(\Delta | \Omega_I) = \frac{1}{(2\pi)^{d/2} |\Sigma_I|^{1/2}} e^{-\frac{1}{2} \Delta^T \Sigma_I^{-1} \Delta},$$

$$P(\Delta | \Omega_E) = \frac{1}{(2\pi)^{d/2} |\Sigma_E|^{1/2}} e^{-\frac{1}{2} \Delta^T \Sigma_E^{-1} \Delta},$$

$$d(\mathbf{x}_i, \mathbf{x}_j) = (\mathbf{x}_i - \mathbf{x}_j)^T (\Sigma_I^{-1} - \Sigma_E^{-1}) (\mathbf{x}_i - \mathbf{x}_j)$$

Ref: M. Kostinger, M. Hirzer, P. Wohlhart, P. M. Roth, and H. Bischof, “Large scale metric learning from equivalence constraints,” *CVPR 2012*.

Distance Function

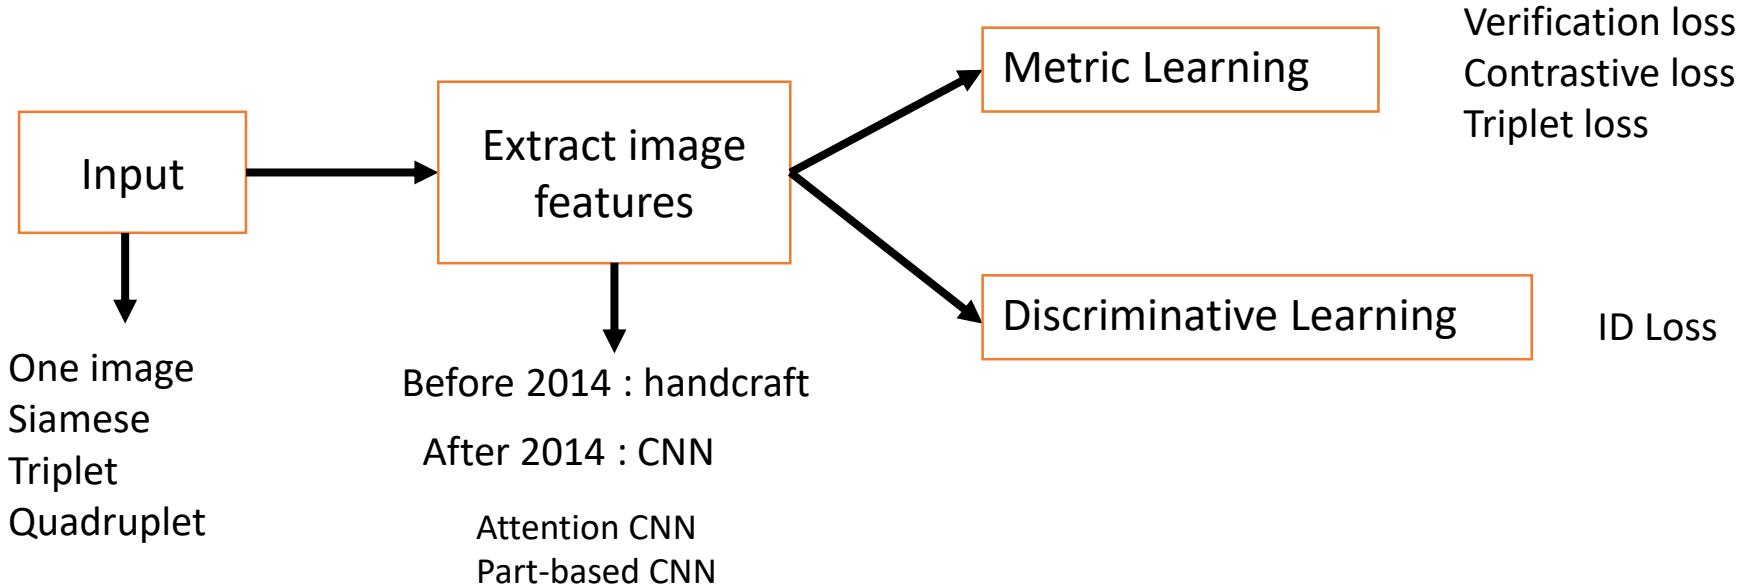
- XQDA: Cross-view quadratic discriminant analysis
 - Goal: to find a subspace mapping matrix \mathbf{w} to reduce the dimension for better classification
 - by maximizing
$$W = (\mathbf{w}_1, \mathbf{w}_2, \dots, \mathbf{w}_r) \in \mathbb{R}^{d \times r}$$

$$J(\mathbf{w}) = \frac{\mathbf{w}^T \Sigma_E \mathbf{w}}{\mathbf{w}^T \Sigma_I \mathbf{w}}$$

Ref: S. Liao, Y. Hu, X. Zhu, and S. Z. Li, "Person re-identification by local maximal occurrence representation and metric learning," *CVPR 2015*.

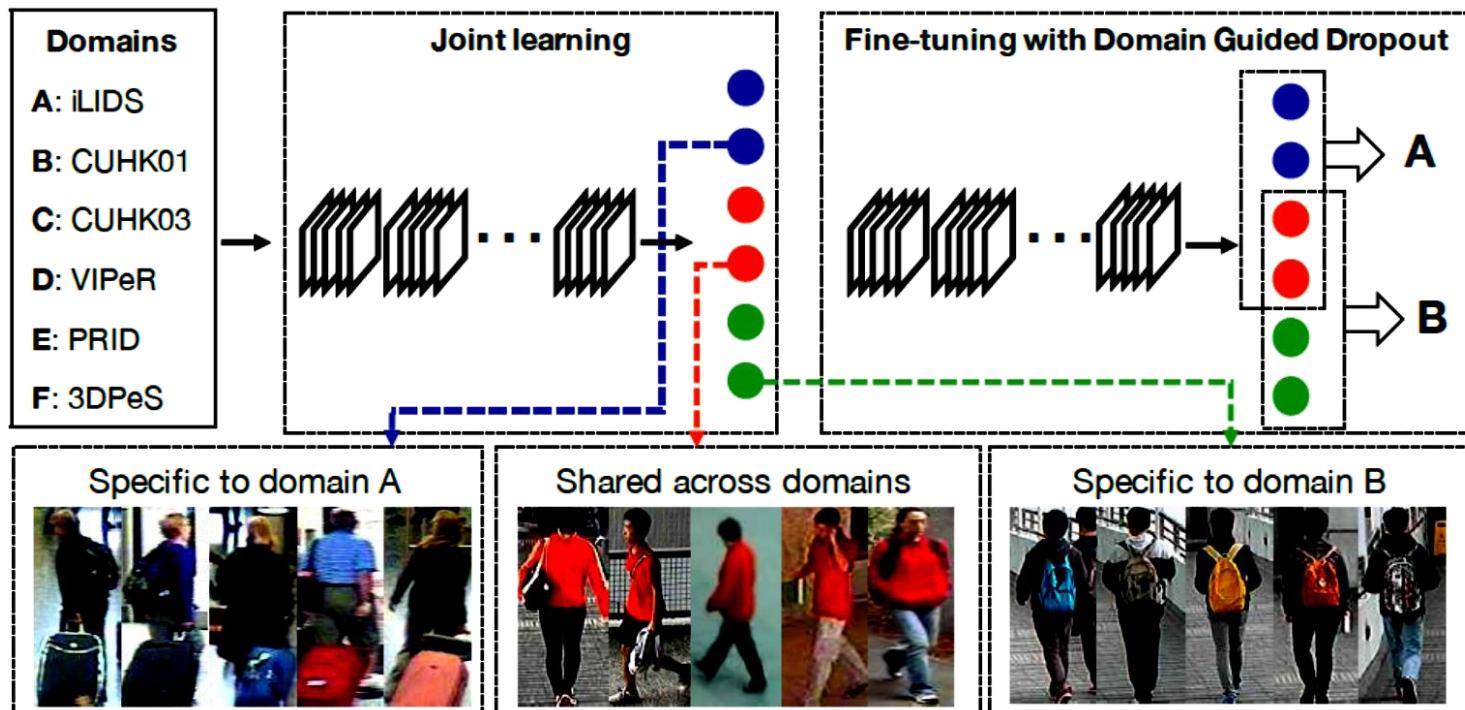
Re-identification System with Deep Learning

- Image-based Re-ID



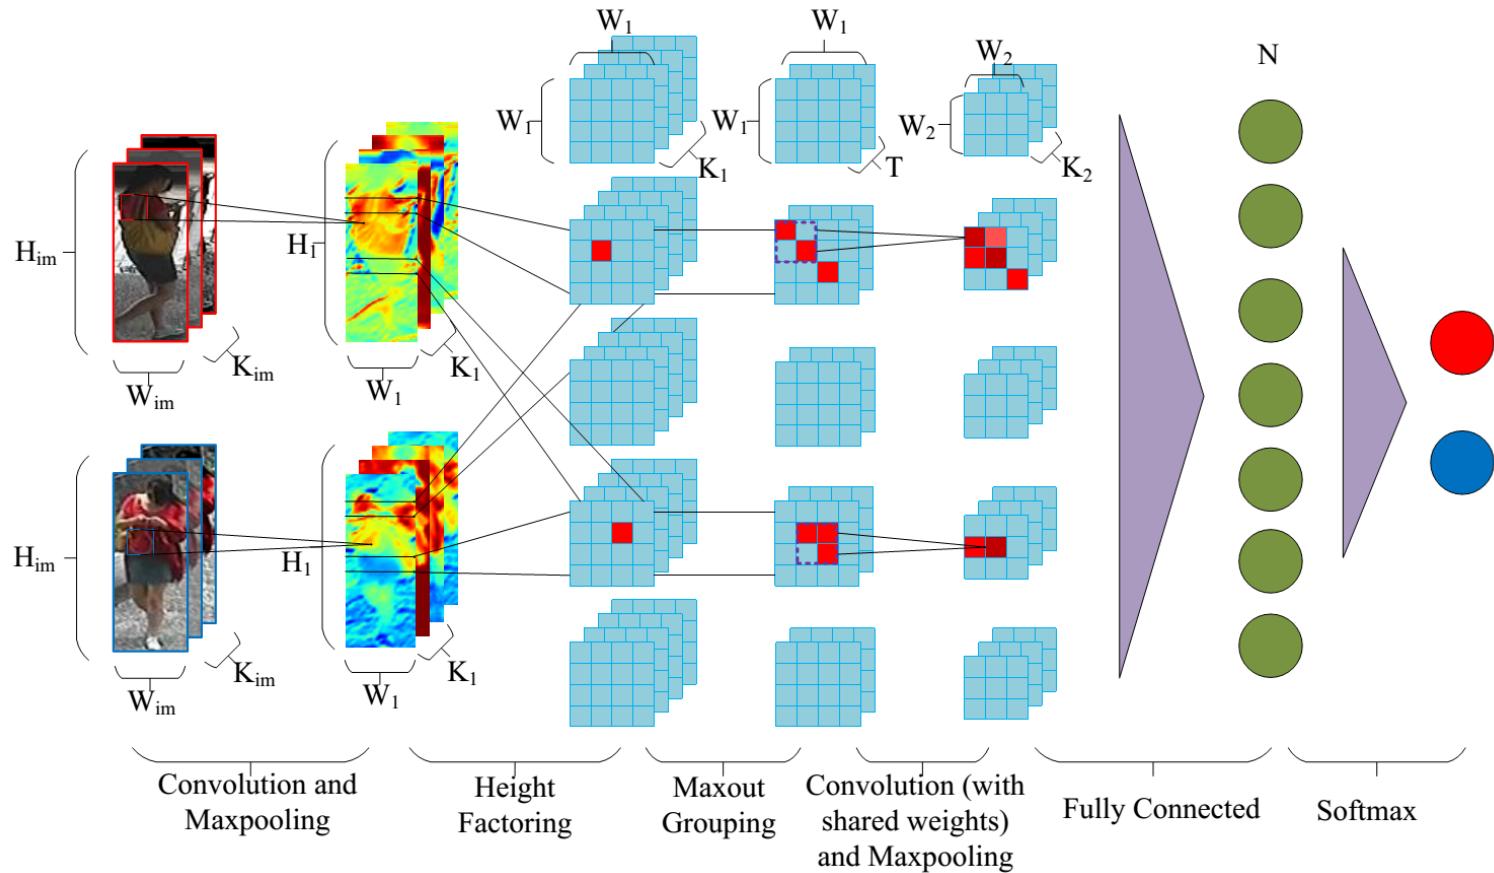
Re-identification System with Deep Learning

- Learning deep feature representation with domain guided dropout



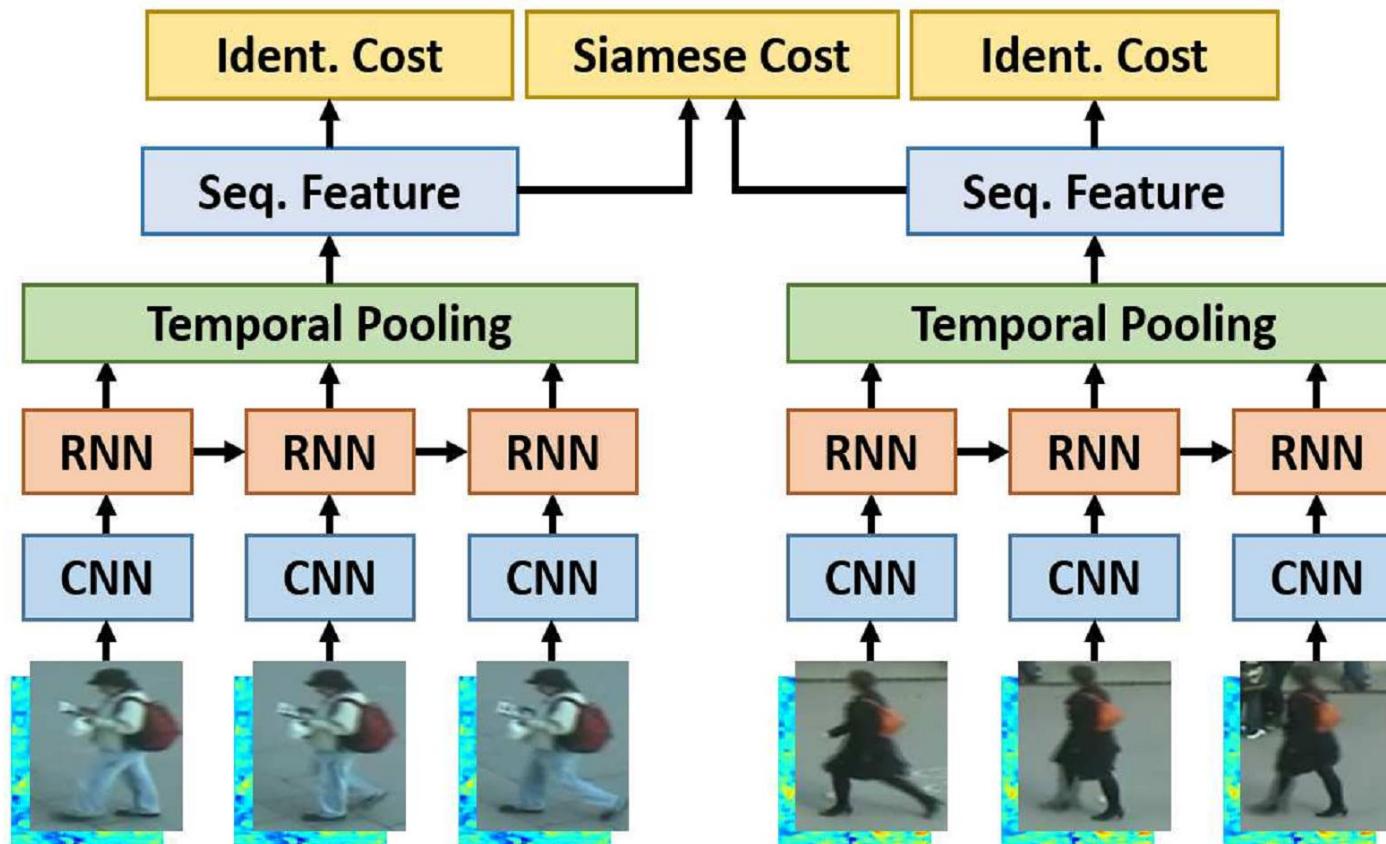
Ref: T. Xiao, H. Li, W. Ouyang, and X. Wang, "Learning deep feature representations with domain guided dropout for person re-identification," *CVPR 2016*.

Re-identification System with Deep Learning



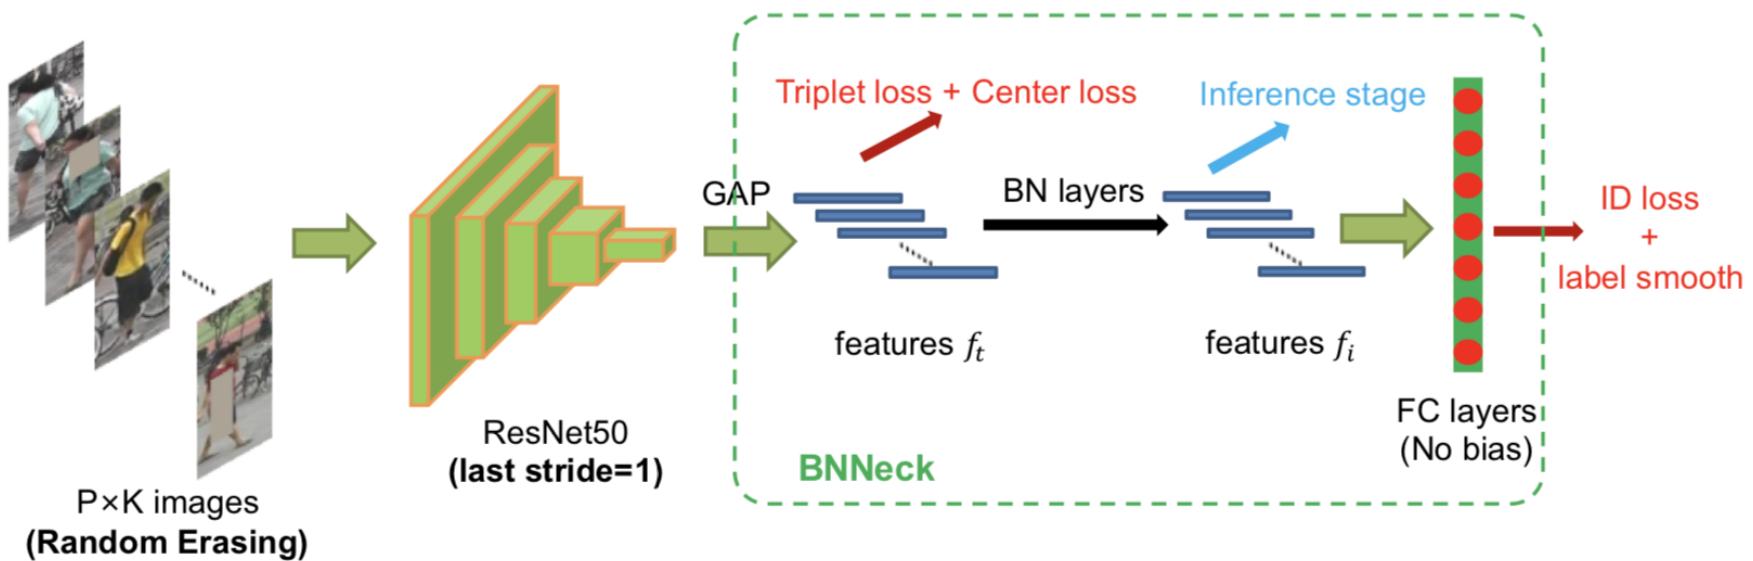
Ref: W. Li, R. Zhao, T. Xiao and X. Wang, "DeepReID: Deep Filter Pairing Neural Network for Person Re-identification," *CVPR2014*.

Re-identification System with Deep Learning



Ref: N. McLaughlin, J. M. del Rincon, and P. Miller, "Recurrent Convolutional Network for Video-Based Person Re-Identification," *CVPR 2016*.

Re-identification System with Deep Learning



Re: H. Luo, Y. Gu, X. Liao, S. Lai and W. Jiang, "Bag of Tricks and a Strong Baseline for Deep Person Re-Identification," *CVPR Workshops (CVPRW) 2019*.